

Thalia Project Report

Team Charlie

**Martti Aukia, Albert Boehm, Arthur-Louis Heath,
George Stoian, Daniel Joffe, Weronika Kakavou,
Marcell Veiner**



University of Aberdeen
Friday 17th April, 2020

Our thanks go out to Dr. Beacham and Dr. Compatangelo for their continued support and guidance, to our logo designer and video assistant Marton Kottmayer and the folks at Stackoverflow for assistance in any technical matters.

Contents

1	Introduction	5
1.1	Project Overview	5
1.2	Motivation	5
1.3	Project Management Strategy	5
1.4	Budget	6
2	Background and Competitors	7
2.1	User Types	7
2.2	Competitor Analysis	8
2.2.1	Portfolio Visualizer	8
2.2.2	Other Competing Software	9
2.2.3	Summary	10
3	Branding	11
3.1	Name and Logo	11
3.2	Motto	11
3.3	Official colours	11
4	Requirements	12
4.1	Functional Requirements	12
4.2	Non-functional Requirements	14
4.3	Use Case	16
4.3.1	Backtest an Investing Strategy	16
4.4	Feasibility Analysis	16
4.4.1	Technical	17
4.4.2	Economic	17
4.4.3	Legal	18
4.4.4	Operational	18
4.4.5	Scheduling	18
4.4.6	Conclusion	18
5	Design	18
5.1	General Design Decisions	18
5.2	Overview of System Architecture	18
5.3	GUI Structure	19
5.3.1	Thalia Web	19
5.3.2	The Dashboard	22
5.3.3	Changes to the Initial Design	25
5.4	Business Logic Structure	26
5.5	Data Harvester Structure	27
5.5.1	Choosing APIs	27
5.6	Database Structure	28
5.6.1	Data Segregation	28
5.6.2	The Data Layer Module - Finda	29
6	Coding and Integration	29
6.1	Overview	29
6.2	Planning	30
6.3	Key Implementation Decisions	32
6.3.1	Web Framework	32
6.3.2	GUI	33
6.3.3	Business Logic	34
6.3.4	Database and Finda	35

6.3.5	Data Harvester	36
6.3.6	Continuous Updating Mechanism	36
6.3.7	API Handler	37
6.3.8	Standardization of Data	40
6.3.9	Interpolation	40
6.3.10	Security	40
6.4	Integration and Deployment	41
6.4.1	Continuous Integration	42
6.4.2	Hosting and Continuous Deployment	44
7	Testing	47
7.1	Testing Strategy	47
7.1.1	Unit Testing	47
7.1.2	Integration Testing	48
7.1.3	End-to-end Testing	48
7.1.4	Data Harvester Logs	48
7.2	Testing Tools	49
7.3	Test Data	50
7.3.1	Mock Data	50
7.3.2	Real-world testing data	51
7.4	Testing Results	51
8	Evaluation	51
8.1	Approach	51
8.2	Functional Requirements	52
8.3	Non-Functional Requirements	53
8.3.1	Usability	53
8.3.2	Reliability	53
8.3.3	Performance	54
8.3.4	Supportability	54
8.3.5	Implementation	55
8.3.6	Interfacing	55
8.3.7	Operations	55
8.4	Packaging	55
8.4.1	Legal	55
8.4.2	Accessibility	56
8.5	Evaluation Results	57
9	Conclusion and Further Work	57
9.1	Conclusion	57
9.2	Further Work	57
9.2.1	Payment Service Integration	57
9.2.2	Implementation of a Scripting Language	57
9.2.3	Additional Risk Metrics	57
9.2.4	Expanding Asset Selection	57
9.2.5	Equity Dividends	58
10	Appendices	59
11	Appendix A - User Manual	59
11.1	Launching Thalia	59
11.2	Browsing Thalia	59
11.2.1	Homepage	60
11.2.2	About Page	61
11.2.3	Log In and Sign Up Pages	62

11.2.4 Contact Us Page	63
11.2.5 About Us Page	64
11.2.6 My Portfolios Page	64
11.3 Running a Backtest	65
11.3.1 Dashboard	65
11.3.2 Ticker Selector	65
11.4 Viewing the Results	67
11.4.1 Summary	67
11.4.2 Metrics	68
11.4.3 Returns	68
11.4.4 Drawdowns	69
11.4.5 Assets	70
11.4.6 Overfitting	70
12 Appendix B - Maintenance Manual	72
12.0.1 Maintenance Design	72
13 Appendix C - Problem Domain Glossary	74
14 Appendix D - Project Terminology	76
15 Appendix E - UML Diagrams	77
15.1 Anda	77
15.2 Data Harvester	78
15.3 Finda	79
15.4 Thalia Web	80
16 Appendix F - User Types Diagrams	81
References	89

1 Introduction

1.1 Project Overview

The aim of this project was to create a portfolio backtesting software which enables creating custom portfolios and measuring their performance with different backtesting functions. The user can pick from a variety of assets, some of which are Equities, Fixed Income, Currencies, Commodities, and Cryptocurrencies. Risk metrics and performance are then visualized for the given asset allocation.

1.2 Motivation

Since retail investing is a growing market, our target audience consists of retail investors, who instead of seeking the full package that comes with financial advising would like to take over the wheel and assess the viability of their investment strategies themselves. As retail investors are non-professionals and invest comparatively small amounts (with financial advice services being non affordable for those individual clients) our goal was to create a product that would not only be more affordable for small retail investors, but would also include a variety of international assets, which most existing backtesting software fails to provide [1].

1.3 Project Management Strategy

The creators of Thalia are:

- Martti Aukia, *Team Leader*
- Albert Boehm, *Chief Editor*
- Arthur-Louis Heath, *Deputy Team Leader*
- Daniel Joffe
- Weronika Kakavou
- George Stoian
- Marcell Veiner

Our goal was the creation of a high quality industrial prototype of the Thalia backtesting tool, based on the identified project requirements and optional features [2]. The team held regular meetings, both with the project guide, Dr. Nigel Beacham, and the course coordinator Dr. Ernesto Compatangelo. During the analysis stage, meetings took place weekly and involved discussion of ideas and requirements. Later, during the implementation stage, the team held meetings at least once per week, some of which consisted of discussing the design for the tool while others involving joint programming sessions.

As in the past [2], our workflow was centered around the GitHub platform and the tools it provides. Furthermore, we continued to follow the Egalitarian Team structure, as this worked very well during the first term of the course. In our Technical Report, we also discussed the use of effort-oriented metrics, i.e. story points, which were assigned to tickets based on the time needed and the functionality of the task. However, this term we concluded that in many occasions these metrics failed to successfully measure the size of a task, as they were seemingly arbitrary. Based on this observation, we decided not to make use of them for the rest of the development process.

As we were fortunate enough to gain an additional member this term, some of our initial effort was focused on introducing her to the project and team dynamics. Given our access to a large team, it was common to see coding tickets assigned to pairs and groups rather than individuals. We believe this not only resulted in writing better-quality code, but also in faster team communication when making design decisions.

1.4 Budget

As previously discussed, we did not have any budget restrictions other than time. All of us agreed to allocating a minimum of 10 hours per week to development efforts and discussed personal expectations well in advance. We estimate the value of our products goodwill to be approximately £13,000. We base this estimate on the developer time allocated to this project, the total duration of 20 weeks, and average developer salaries in the UK [3] (taking into account that we have no industry experience, i.e. that the value of our work is towards the lower end of the range provided).

2 Background and Competitors

The following sections provide for a discussion of how we defined our target audience through creating User Types and an in-depth analysis of competing backtesting tools.

2.1 User Types

In the creation process of our user types, we have mainly used the methodology described in [4]. In addition, the advice given by [5] and [6] have been very helpful. The process started with us gathering as much data as possible about the categories of people that would be interested in backtesting and, therefore, in investing. The next step was that of creating several skeleton personas [7] and validating their legitimacy with our team members and the information we had about the general types of investors that exist. After that, we decided that the creation of four personas was required to cover all of our target demographic. The four personas created are described here with the full charts attached to this document in the appendix 16.

- **Addriana:** Addriana is a 56-year old flower shop owner who is trying to invest in order to gain a comfortable retirement. She has used the services of a financial advisor but was disappointed with them. As a result, she wants to have more control and more hands-on knowledge of the investments she is making.
- **Allison:** Allison is a finance student in her 20s who wants to use the knowledge she acquired from university in a way that would allow her to both make a profit and gain insights that would help her in her career.
- **Sam:** Sam is an accomplished middle-aged man who wishes to know more about investing and what benefits and disadvantages come with it.
- **Timmy:** Timmy is a high-school student who is interested in short term investments. He would described himself as being part of the new wave of cryptocurrency enthusiasts.

With the personas created, we conceived of a number of scenarios where backtesting provided value. We then put the personas through the created scenarios. This process helped us extract the empathy maps [8] of our personas. We then validated our empathy maps with our team and outsiders. The validation with people from outside our team was crucial in eliminating some of our internal biases. Based on the user personas and their empathy maps, we have started looking at what types of metrics are best at describing our personas. The ones of critical importance we have seen are the knowledge on the topic of finance and the size of funds they have available for investing. Another metric that was found to be important was that of the types of investing timeframes. Based on those metrics, we have identified four user types.

1. **Financial Proficiency and Large amounts of Capital:** This person already knows the benefits of backtesting and only needs a tool that can speed up the process of constructing an investment portfolio. The resources they have means that the price of using our product would not be a problem for them. This type of user would do engage in both long and short term investments.
2. **High Financial Knowledge and Little Capital:** This is most likely someone young who has a keen interest in finance or is studying it. The price could be a deterrent for using our product if they are doing short term investing. If they would be doing long term investing, however, the price did may not pose as much of a problem since it is a part of their initial investment.
3. **Little Financial Knowledge and Large amounts of Capital:** This is a person who wants to switch from having a savings account or a financial advisor. They are aware of the learning curve and need additional information before they will commit to anything. They will be happy to pay for our product in order to gain access to our data, analysis and general advice.
4. **Little Financial Knowledge and Little Capital:** This last category is represented by people who are not confident in their ability to invest successfully. As a result, they do not want to risk large sums of money until they will be sure that they can make the right decisions. There is a good chance that they will acquire our product to form an idea of what investing is like. If they get more confident, they can increase the funds they'd be comfortable to invest.

2.2 Competitor Analysis

As a result of our initial competitor analysis, we can group the competing software into two categories. For a comprehensive list of available backtesters, please see [9]. The first consists of full-suited trading software, such as Fidelity Active Trader Pro [10], MetaTrader4 [11] and NinjaTrader [12] which, among others, offer backtesting features. Services in this category include features such as buying and selling of financial assets, prediction of future prices and storing and managing of a user's capital. Although it may be beneficial to consider some features related to real time trading, these would only be part of a future development process and are not within scope of our prototype.

The second category consists of software that does not offer trading as a service and solely focuses on backtesting. Thus, for now, we shall only consider the feature set provided by the second category. In the following paragraphs we will consider some of the design decisions made by competing software, together with the brief analysis and conclusion on each feature.

2.2.1 Portfolio Visualizer

Our main competitor of the second category is an online backtesting tool named Portfolio Visualizer [13]. During the inception phase of development, we relied heavily on this website for writing the requirement analysis. Let us now briefly dissect what it has to offer.

Upon opening the website we are greeted with a brief description of the domain, together with the following input form:

The screenshot shows a web-based input form for a portfolio analysis. It includes the following fields:

- Time Period**: A dropdown menu set to "Year-to-Year".
- Start Year**: A dropdown menu set to "1985".
- End Year**: A dropdown menu set to "2019".
- Initial Amount**: An input field containing "\$ 10000 .00".
- Cashflows**: A dropdown menu set to "None".
- Rebalancing**: A dropdown menu set to "Rebalance annually".
- Display Income**: A dropdown menu set to "No".
- Benchmark**: A dropdown menu set to "Vanguard 500 Index Investor".

Figure 1: Portfolio Visualizer - Input (source: <https://www.portfoliovisualizer.com/>)

As we can already see, the key elements of portfolio analysis are provided. Considering the functionality of our prototype as a baseline (i.e. testing an allocation of assets with a fixed initial investment over a fixed time period), we see that in addition, users are able to perform the following set of actions:

- Set the endpoints of the investment period, up to months.
- Set the initial investment.
- Specify a regular cash flow and its frequency.
- Select a rebalancing strategy.
- Select a benchmark strategy for comparison.
- Compare multiple strategies at the same time.
- Adjust for inflation.

- Select from a set of lazy portfolios (benchmarks).
- Calculate additional metrics.
- Export the results to PDF, Excel, or save link.

At first, it may seem as if Portfolio Visualizer meets all the requirements needed for a backtesting tool. Indeed, our main criticism is with respect to the UI, responsiveness and user experience, which are all findings from initial user testing.

The UI design is simplistic and has a non-commercial, bare-bones look. Although this was appreciated while testing the system, it certainly does not improve the user experience. The input, mostly using dropdown menus, is straightforward to use, except for the selection of assets, which we will discuss briefly at the end of this section.

The website also has user accounts. However, little to no data is associated with a user's profile, which further decreases the overall user experience. Finally, as a last remark, which holds for most of the backtesting tools we have tried, the website is heavily US biased, hence the selection of asset classes is limited.

2.2.2 Other Competing Software

The rest of the alternatives are of a significantly lower quality. In the remaining parts of this section, we will briefly consider a few design decisions made by these websites.

Tree-like Asset Selection

One challenging aspect of designing such a system is the following:

What is the most intuitive way for selecting a portfolio item?

Where a portfolio item could be: an equity, ETF index, a commodity, bond, or the like. Within these classes, we have more options to choose from. Many backtesters opted for a search bar, which is a sensible approach. Nevertheless, many portfolio items are named similarly, and for a new user it is a high entry barrier as they might not know what is available.

Another option is what ETFReplay [14] has implemented, which is a tree-like selection form, shown below.

The screenshot shows a web browser window for 'etfreplay.com/symbols.aspx?id=ContentPlaceHolder1...'. At the top, there is a search bar labeled 'Looking for:' with radio buttons for 'ETF' (selected) and 'Mutual Fund'. Below this, a section titled 'Select a category:' contains several expandable categories. The 'U.S. Equity Index ETFs' category is expanded, showing sub-options: Primary U.S. Equity Index ETFs (with sub-options for DIA, IWM, QQQ, SPY, VTI), Smart Beta Sample, U.S. Dividend ETFs, U.S. Growth vs Value, U.S. Large Cap, U.S. Small & Midcap, and Other U.S. Equity ETFs. The 'U.S. Sector ETFs' and 'International Equity' categories are collapsed. The entire interface is styled with dark blue and orange highlights on the buttons.

Figure 2: ETFReplay - Tree-like Structure (source: <https://www.etfreplay.com/>)

However, pursuing this form of gathering input from the user involves multiple repetitive actions of expanding the tree structure, which becomes tedious once one is more familiar with the tool. Thus, we have chosen to pursue the search bar approach.

Tiles

It may be worth briefly discussing one example of a backtester with a good design. We found the simplistic and tiled design of Backtest Curvo [15] a good choice, as it gives the website an overall modern and fresh look, whereas most backtesters looked old and out of date. Our goal is to achieve a layout similar to this. For a further analysis on design decisions, please see 5.

Simple Scripting Language

In our whitepaper, we briefly discussed implementing a simple scripting language which would allow users to simulate dynamic trading strategies [1]. As mentioned in our white paper, this would not be in the scope of the course, but for inspiration we can have a look at the approach of Stockbacktest [16].

Signals (Need at least one buy or short signal)

Buy When:

Signal 1:	Upper Bollinger Band (days devs) ▾	Crosses Below ▾	Percentage % ▾	AND ▾
20 2	none	70		
Signal 2:	Close Price ▾	Is At Least % Above (%) ▾	Percentage % ▾	
none	15	10		

Figure 3: Stockbacktest - Simple Logic (source: <http://stockbacktest.com/>)

Although not intuitive at all, it demonstrates that this feature is possible to implement. We believe that after nailing down precisely how the user would create the rules, this would be a straightforward task, as the calculations involved are not more complicated than in the static case. As mentioned previously, this is something we would like to implement in the future.

2.2.3 Summary

After testing competitor software during the inception phase of development, we were left with three key observations. These were the following:

- Transparency: Many of the testers' web pages did not give the impression of being operated by a reputable businesses. In the worst case, they gave the impression of being nontransparent or untrustworthy (misspelled words, advertisements with questionable subject matter), which is something we should avoid.
- Learning Curve: After getting to know our target audience we concluded that most seem willing to learn how to use a complex system, if it is worth it.
- Design of the UI: It was easy to tell which backtesters are still getting updated by looking at their design. We should aim at looking fresh. In addition, some backtesters offered so many features that every corner of their UI was packed with information. This is something we should also avoid doing.

The analysis was conducted informally with the help of colleagues studying or otherwise involved in finance. This allowed us to develop a more neutral and accurate impression, despite the subjective nature of the observations. Having analysed the competing software, we are now better suited to determine the requirements of our system.

3 Branding

We have started the process of creating a brand for our product by looking at books and articles on branding. The two most useful resources we have seen were [17] and [18]. Based on the knowledge acquired, we have decided that our brand needs a name, a logo and a motto.

3.1 Name and Logo

In order to find a name and a logo that would represent us, we have conducted several brainstorming sessions where we have laid out the keywords that represent our product. Out of all the words we listed, two recurring categories appeared, which were ‘flourishing’ and ‘data’. We then went on to find ways of representing those. At the end, we have decided on choosing ‘Thalia’ as the name of the product, which is the name of a Greek muse often described as ‘the growing one’. We saw the association between our product and the Greek muse of growth as beneficial.

To create a logo to our liking, we have asked a young designer willing to create it Pro Bono for us. After a few iterations, we had two logos made from copyright free images, the smaller of which can be seen on Figure4.



Figure 4: Thalia Logo - Small

3.2 Motto

Based on our previous findings in 2.2.3, we have decided that our motto should be something that would reassure our customers that they have made the right choice. As a result, we have come up with a simple motto: ‘Thalia, the reliable Backtester’.

3.3 Official colours

The idea of defining a set of official colours came from homogeneity. We strove for a uniform design both for our product and this report, so we believed it was a good way of ensuring this. Given our team’s limited experience with design, we have decided to start by choosing the base colour of our palette. In order to be up to date with the latest design trends, we have sought to use something close to the colour of the year 2020 [19]. Based on this colour, we have created three colour palettes by using a triadic colour scheme [20].

In the following figure, the final three colour schemes are presented. After some deliberation within the team and some feedback from a few outside people, we have decided to use the option labelled ‘Thalia’.



Figure 5: Color Schemes

4 Requirements

Earlier in the inception phase of the development, we classified our requirements using the established FURPS+ model [21]. Ever since identifying requirements for our initial project report [2], we have also identified some additional requirements that Thalia should fulfill, which are based on feedback from our project guide. Let us briefly revisit our main functional and non-functional requirements. The items outlined below focus on some of the key requirements we have so far identified as features necessary for providing a compelling product for paying customers.

4.1 Functional Requirements

Portfolio Configuration	
Allocate fixed amount/proportions of the portfolio to given assets	Choose how much each asset contributes to the portfolio's total value using either percentages or raw monetary amounts
Find assets quickly by category or name	When adding an asset the user can search a category for assets or search for a specific asset by its name
Share portfolio	Portfolios can be shared between people using a URL
Edit portfolio	Change asset allocation and their distributions in a portfolio

Portfolio analysis	
Compare portfolios	Use multiple portfolios in a single analysis to see differences in their performance
Use a selection of lazy portfolios	Select an existing common portfolio to compare against, such as common index funds (e.g. Vanguard 500 Index Investor or SPY)
Plot portfolio as a time-series	View portfolio performance as a line graph for a quick overview
Specify a time frame for the analysis	Select start and end dates for portfolio analysis
Choose rebalancing strategy	Optionally choose a strategy for buying and selling assets to meet your strategy e.g. buying and selling stocks each year to ensure the value of portfolio stays at 60% stocks and 40% bonds (i.e. maintain the initial allocation)
Change the distribution of assets in a portfolio using a slider	A slider for each asset to quickly increase or decrease its proportion of the total value
Edit portfolio analysis	Change parameters for portfolio's analysis after running it (e.g. date range or rebalancing strategy)

View results	
See key numerical figures	Show important numerical metrics for a portfolio's performance such as Initial Balance, Standard Deviation, Worst Year, Sharpe Ratio, and Sortino Ratio
See both real and nominal values	See portfolio's value as both adjusted and not adjusted for inflation
A breakdown of portfolio value at specific points of time	See what the value of the portfolio is at some point in time (e.g. January 3rd 1997)

User accounts	
Combine portfolios	Combine two portfolios' assets into one single portfolio
Save portfolio analysis for later	Save portfolio analysis parameters to the account so you can rerun it with a single click
Delete saved portfolio analysis	Remove a stored portfolio analysis from your account
Manage portfolio analyses	Edit saved portfolio analysis with different assets, distributions or other parameters
Sign-up, log in and log out	Basic authentication

Assets	
Choose assets from European market	Data for European assets were found to be lacking in competing products
Choose assets from Equities, Fixed Income, Currencies, Commodities, and Cryptocurrencies	Coverage of some of the largest asset classes

4.2 Non-functional Requirements

Non-functional requirements were elicited using the FURPS+ model developed by IBM to address the issues of the FURPS model, namely failure to take developer requirements into consideration [22]. The decision to use the FURPS+ model was made based on its wide adoption in industry as well as the high volume of literature and guides for identifying key requirements available. Using this model, we identified various categories of non-functional requirements based on either user or developer needs. Elicitation of user-focused non-functional requirements was performed based on feedback from our focus group from the first semester and the capabilities and properties of competing software (e.g. the legal requirements of providing a backtesting service). Although most requirements were identified during requirement elicitation in the first semester, several requirements and categories that we initially missed were added later during development, namely ‘Accessibility’ and ‘Supportability’ requirements.

1. Usability

- The product must be easily usable for users who already have some financial investment experience.
- The basic backtesting interface needs to look familiar to people already experienced with the process.
- The product must have detailed instructions on how to use its advertised functions.
- All major functions must be visible from the initial landing page.
- Must work in both desktop and mobile browsers.
- The results page should scale with mobile.

2. Reliability

- The product must have a greater than 99% up time.
- All our assets need to have up to date daily data whenever the asset is still publicly traded.
- All assets supported by the system must provide all publicly available historical data.

3. Performance

- The website should load within 3 seconds on mobile [2].
- Large portfolios must be supported with up to 300 different assets in one portfolio.

4. Supportability

- Software should be well documented and easy to maintain.
- The system should be fault-tolerant and be designed to support graceful degradation [23].
- The system should be easy to install and migrate between hosting providers.
- The system should be internationalized and be easy to extend to support new languages and currencies.

5. Implementation

- The system needs to work on a cloud hosting provider.

6. Interfacing

- The Data Gathering Module must never use APIs stated to-be-deprecated within a month.
- The Data Gathering Module must not exceed its contractual usage limits.

7. Operations

- An administrator on-call will be necessary for unexpected issues.

8. Packaging

- The product needs to work inside a Linux container (e.g. Docker).
- All dependencies need to be installable with a single command.

9. Legal

- All user testing must be done with ethical approval from the University.
- The UI must display a clear legal disclaimer about the service not providing financial advice.
- All third-party code should allow for commercial use without requiring source disclosure (e.g. no GPL-3).
- User data handling should comply with GDPR.
- Provided services should not constitute financial advice under UK law to avoid being subject to financial advice legislation and potential liability.

10. Accessibility

- Display items should be clearly labelled.
- UI should scale to accommodate different screen sizes and aspect ratios.
- UI elements and text superimposed over one another should have high contrast in their colors.
- UI should allow for the use of assistive technologies to accommodate for individuals with accessibility issues.

4.3 Use Case

4.3.1 Backtest an Investing Strategy

Actors: User

Purpose: Show a user how their strategy would have performed in the past.

Overview: The user selects the assets they want in their portfolio, along with how much to invest in each asset. They may also select contribution, withdrawal and rebalancing options. On completion, the user will see a graph plotting the value of their portfolio over time and a table of a few key metrics summarising performance.

Preconditions: The user must be logged in.

Flow of Events:

1. The user navigates to the *dashboard* page of the website.
2. The system sends a list of available assets.
3. The user enters the period of time over which they wish to backtest.
4. The user selects a financial asset.
5. The user enters the proportion of their portfolio they want to invest in that asset.
6. The user submits their strategy for evaluation.
7. The system analyses the user's strategy over the specified period.
8. The user receives their strategy's performance.

Alternate Flow of Events

- Steps 4 and 5 may repeat any number of times, allowing the user to have as large a portfolio as they like. However, they must happen at least once.
- In step 6, the system may have insufficient data for the user's desired time range. If this is the case, the system should give results for the time it does have data for.
- Between steps 3 and 6, the user may optionally enter a rebalancing frequency and a contribution amount and frequency. The system should take these into account when calculating.

4.4 Feasibility Analysis

The feasibility of this project hinges not only on technical considerations, but must be evaluated across multiple dimensions. To conduct our analysis, we have used the TELOS methodology [24]. TELOS is an acronym for:

- Technical
- Economic
- Legal
- Operational
- Scheduling

The following sections will analyse each of these areas in detail to allow for drawing conclusions about the project feasibility.

4.4.1 Technical

This area concerns the question of whether there exist technologies that enable us to realise our project idea. We consider the prototype developed over the last semester sufficient evidence for the technical feasibility of our project as we have successfully developed implementations for each of our system components, albeit in stripped-down form. As a result of familiarising ourselves with the technologies we used in developing this prototype, we are confident that scaling it up to a fully functional product will pose little technical challenge to our team.

4.4.2 Economic

To reduced reliance on external funding for operating our business over time, license sales must outweigh costs. As part of our effort to determine economic feasibility, we have forecasted our cash flow for the year of 2020:

Figure 6: Thalia Cash Flow Forecast

Thalia Cash Flow Forecast												
Starting cash on hand	£ 10,000.00											
Starting date	Jan 2020											
Cash minimum balance alert	£ 2,000.00											
	January 2020	February 2020	March 2020	April 2020	May 2020	June 2020	July 2020	August 2020	September 2020	October 2020	November 2020	December 2020
Cash on hand (beginning of month)	£ 10,000.00	£ 37,690.00	£ 34,530.00	£ 32,120.00	£ 29,960.00	£ 19,050.00	£ 8,890.00	£ 48,980.00	£ 37,820.00	£ 29,910.00	£ 21,750.00	£ 14,340.00
Total												
Cash Receipts												
License sales	£ -	£ 500.00	£ 1,000.00	£ 1,500.00	£ 3,000.00	£ 4,000.00	£ 4,000.00	£ 3,000.00	£ 6,000.00	£ 6,000.00	£ 6,500.00	£ 7,500.00
Business Loan	£ 20,000.00	£ -	£ -	£ -	£ -	£ -	£ -	£ -	£ -	£ -	£ -	£ 20,000.00
Venture Capital	£ 25,000.00	£ -	£ -	£ -	£ -	£ -	£ 50,000.00	£ -	£ -	£ -	£ -	£ 75,000.00
Total Cash Receipts	£ 45,000.00	£ 500.00	£ 1,000.00	£ 1,500.00	£ 3,000.00	£ 4,000.00	£ 54,000.00	£ 3,000.00	£ 6,000.00	£ 6,000.00	£ 6,500.00	£ 7,500.00
Total Cash Available	£ 55,000.00	£ 38,190.00	£ 35,530.00	£ 33,620.00	£ 32,960.00	£ 23,050.00	£ 62,890.00	£ 51,980.00	£ 43,820.00	£ 35,910.00	£ 28,250.00	£ 21,840.00
Cash Paid Out												
Advertising	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 2,400.00
Contract work	£ -	£ 250.00	£ -	£ 250.00	£ -	£ 250.00	£ -	£ 250.00	£ -	£ 250.00	£ -	£ 250.00
Materials and supplies (in COGS)	£ 14,000.00	£ 100.00	£ 100.00	£ 100.00	£ 100.00	£ 100.00	£ 100.00	£ 100.00	£ 100.00	£ 100.00	£ 100.00	£ 15,100.00
Office expenses	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 150.00	£ 1,800.00
Rent or lease	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 1,000.00	£ 12,000.00
Loan Repayments	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 1,300.00	£ 15,600.00
Repairs and maintenance	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 250.00	£ 3,000.00
Licenses	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 210.00	£ 2,520.00
Utilities	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 200.00	£ 2,400.00
Wages (less emp. credits)	£ -	£ -	£ -	£ -	£ 10,500.00	£ 10,500.00	£ 10,500.00	£ 10,500.00	£ 10,500.00	£ 10,500.00	£ 10,500.00	£ 84,000.00
Total Cash Paid Out	£ 17,310.00	£ 3,660.00	£ 3,410.00	£ 3,660.00	£ 13,910.00	£ 14,160.00	£ 13,910.00	£ 14,160.00	£ 13,910.00	£ 14,160.00	£ 13,910.00	£ 14,160.00
Cash on hand (end of month)	£ 37,690.00	£ 34,530.00	£ 32,120.00	£ 29,960.00	£ 19,050.00	£ 8,890.00	£ 48,980.00	£ 37,820.00	£ 29,910.00	£ 21,750.00	£ 14,340.00	£ 7,680.00

This forecast makes several assumptions that we deem justifiable:

- Seed and Series A venture capital investments of a total of £75,000.
- Approximately linear growth of license sales throughout the year.
- Team members being able to sustain themselves without a salary for the four starting months.
- Being able to acquire a £20,000 business loan with less than 5% interest p.a.

These assumptions lead to our company being profitable by mid 2021. Time to profitability of around 18 months is very typical for startups, as shown e.g. by Olsen and Kolvereid [25]. While these findings are indicators of economic feasibility, we will have to constantly revise our cash flow forecast to account for changing market conditions and advances in product development.

4.4.3 Legal

Fundamentally, backtesting is not considered to encourage purchases or sales of assets. Our company is merely providing information to our customers and thus can not be considered to provide financial advice. Decisions drawn from the information presented by Thalia solely consists of the opinions of our customers and do not reflect the opinions of the Thalia team or any associates. A disclaimer of this form will be placed on our website and within our terms and conditions which have to be accepted by users upon sign up. Within the framework of the Financial Conduct Authority, our product is understood to provide guidance on investment decisions. Entities offering guidance “are responsible for the accuracy and quality of the information they provide” [26]. Thus, it is paramount that the information provided by our tool is accurate.

Our research has not revealed any other potential legal issues with our service. Given correctness of our calculations, we thus consider this project to be legally feasible.

4.4.4 Operational

After deployment of our service, we will need to continually update price data for all assets while integrating new assets to stay competitive. For a discussion of how live data is integrated into our service, please refer to the design section concerning the Data Collection module 6.3.5. Given that these requirements are the only way in which we deviate from other web applications, we are confident that this project is feasible from an operations standpoint.

4.4.5 Scheduling

The Thalia Technical Report established a schedule for this sub-session. The aforementioned strict requirement of correctness of our service makes it difficult to estimate the time required for developing a system component. We are using the Agile development methodology as described in our white paper [1] as a means of adjusting for this uncertainty. Our previous schedule must then be understood to be less rigid and used instead for orientation of how to prioritise work efforts. Delivery of a correct product is thus ensured at any point. Scheduling issues are therefore less of an issue, as reducing the feature specification in favour of accuracy of results is the preferred course of action.

4.4.6 Conclusion

The areas in the TELOS analysis yielding questions about the feasibility of this project are Economic and Legal. Accordingly, we have incorporated actions that reduce any potential negative impact of these risks into our strategy. In particular, we will be including legal disclaimers on our website and as part of our terms of service and try to acquire external funding for the initial development stage of our business as soon as possible. We believe that these discussions solidify the case for our project being feasible.

5 Design

5.1 General Design Decisions

We have chosen pandas dataframes [27] as the common data format for both data exchange and any calculations. Pandas provides us with data structures “that cohere with the rest of the scientific Python stack” [28], such as NumPy, which we are using to calculate historical returns and risk metrics (5.4). Additionally, it is supported natively by many third party APIs and can even be used to read data from SQL databases. For additional information, please consult the official pandas documentation [27].

5.2 Overview of System Architecture

Before discussing any specific component and architectural layer in depth, it is worth revisiting our original architecture [2].

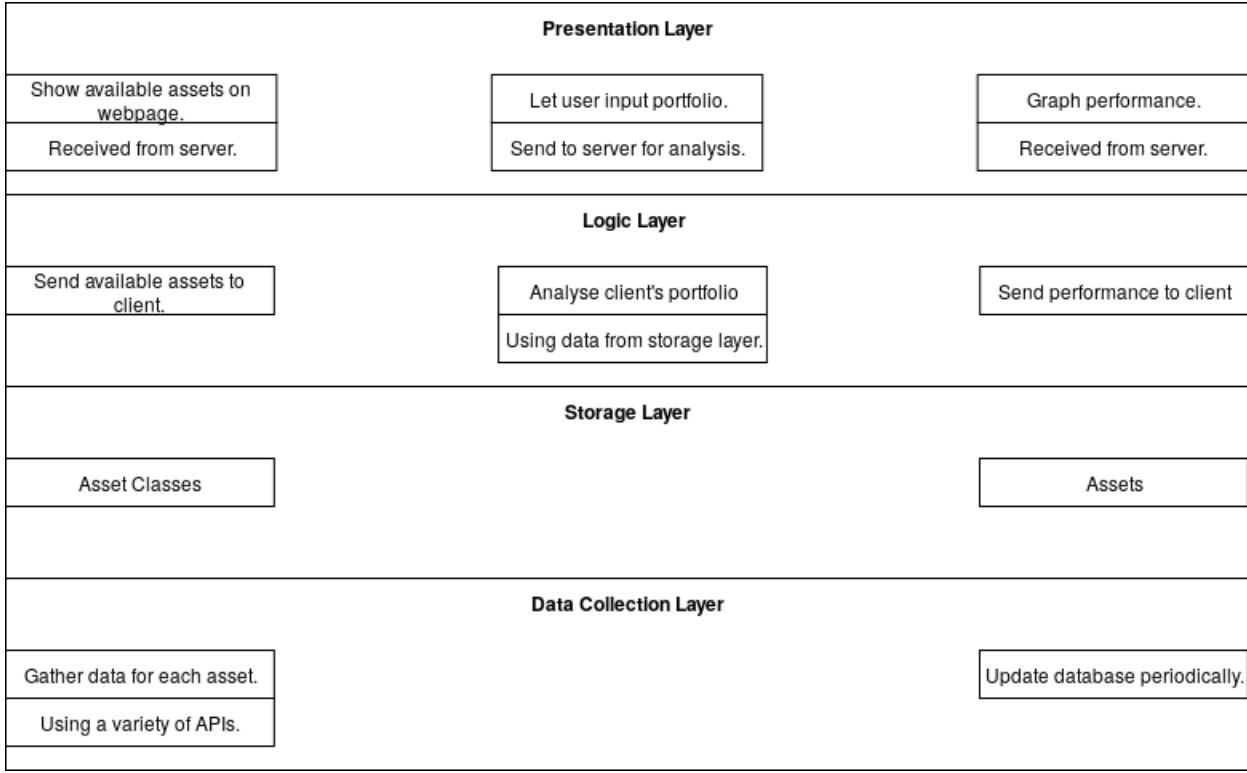


Figure 7: Original Thalia System Architecture [2]

The above schematic illustrates how we modified a typical Three Layer architecture to include a Data Collection Layer. Subsequent discussion will refer to this layer as the ‘Data Harvester’. It consists of adapters for third party APIs which offer price data for financial assets. The associated API will be queried according to a configurable interval to allow for live updates to our price data. Additionally, initial seeding of the database with historical price data can also be achieved via the Data Harvester. The decision to isolate this system component has been made in order to increase security by limiting the access of the Thalia web service to the financial data database to read-only and to allow for independent scaling of the Data Harvester and our web service [2].

While the architecture of our system has stayed the same, there have been some modifications to individual components. Most notable are the changes to the Database Structure examined in 5.6. The following sections will provide for a discussion of design decisions made on a layer-by-layer basis.

5.3 GUI Structure

The Graphical User Interface of Thalia consists of two main parts. One is a *static* website created with HTML5 [29] and stylised with the help of Bulma [30]. The second is a *dynamic* Dash [31] page, which is basically the bulk of our application. In the following paragraphs we will discuss these separately.

5.3.1 Thalia Web

Thalia Web consists of 5 pages, which all serve as an introduction to our application. We aimed at creating an authentic and modern looking website that attracts users. This was done by following a consistent style for the whole website, i.e. using the official colours, motto and Thalia logo, as discussed in 3.

While Thalia was not meant to be used on smaller devices due to the nature of our application, with the help of Bulma, every component of Thalia Web is responsive. Consequently, the layout of the website changes so that it fits the browser size perfectly.

Even though it is possible to navigate between pages by URLs directly, we have designed a navigation bar for this purpose. The look of this navigation bar depends on the device of the current user and whether the user has already been logged in or not. By default, an unauthenticated user will see the navigation bar visible on Figure8.

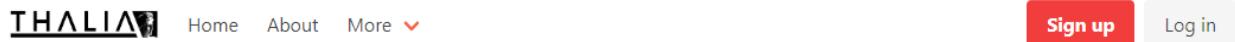


Figure 8: Thalia Navigation Bar

In case Thalia is launched on a mobile phone, the navigation bar becomes a so-called ‘hamburger button’ and dropdown menu visible on Figure9.

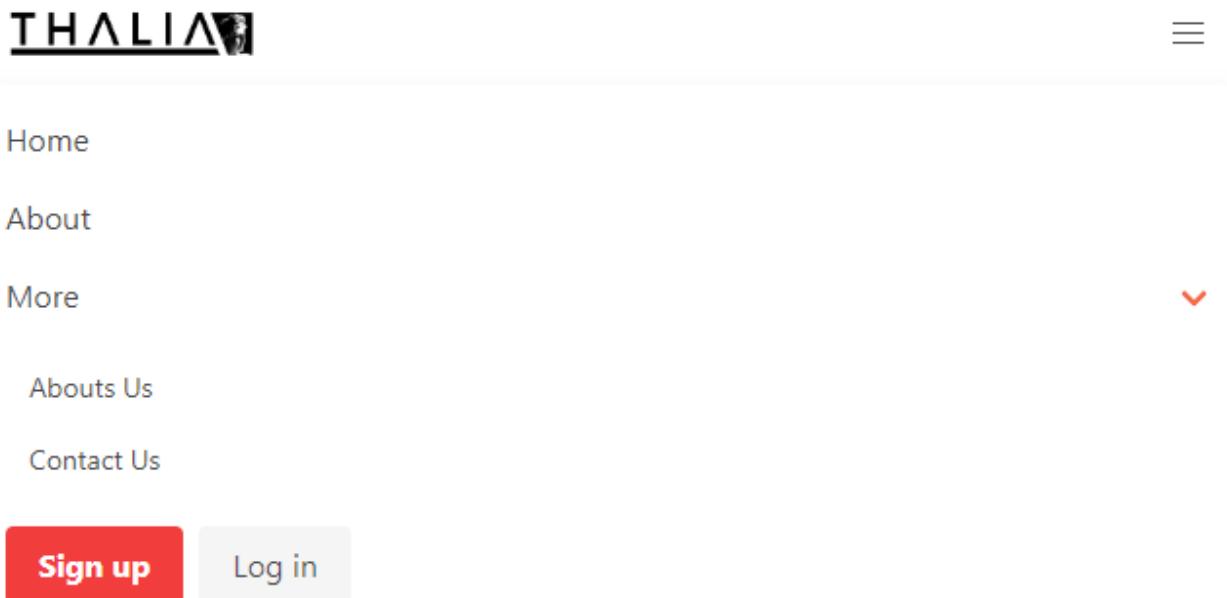


Figure 9: Thalia Navigation Bar - Hamburger

For the rest of this section we shall assume that the application was launched on a desktop, although the layout is identical and intuitive in the mobile case.

Homepage

We shall not discuss every page separately, as they are identical. For a full walkthrough, please refer to 11. As an example of our design choices, let us look at the homepage where users should arrive by default.

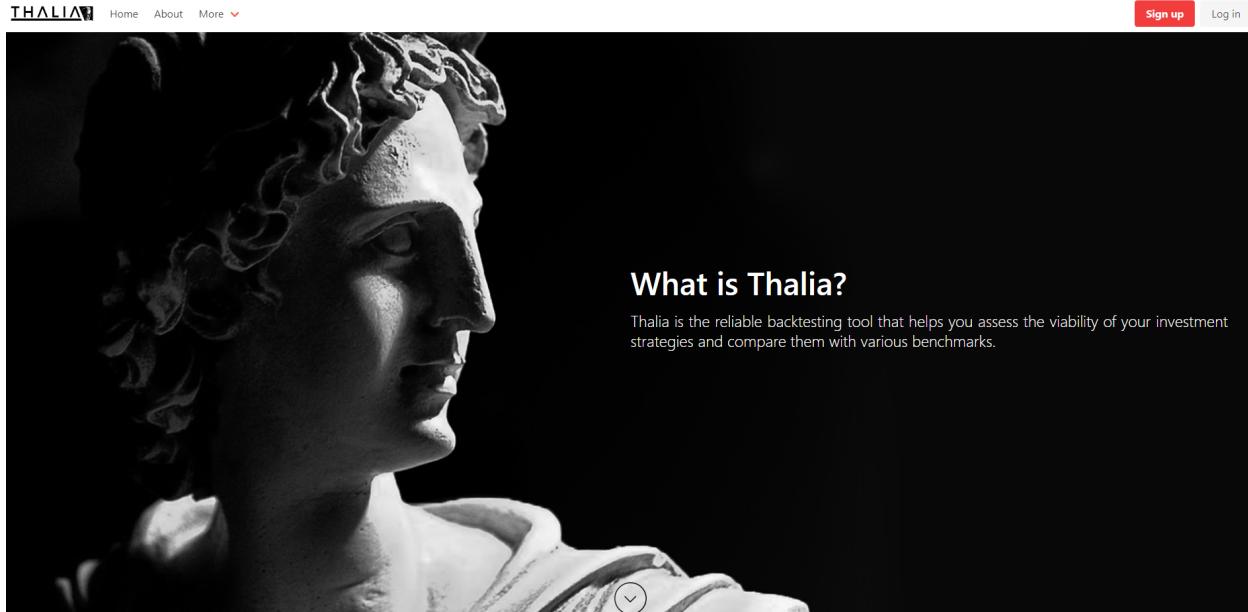


Figure 10: Thalia Homepage (source: <http://thaliabacktest.xyz/>)

As users arrive on the homepage, they are greeted with a short description of Thalia and the process of backtesting. We have designed the homepage such that the navigation bar and the Thalia Logo fills the screen completely. To indicate that the page does not end there, we have added a scroll down button, which with the help of a JavaScript [32] animation takes the user down to the bottom of the page. Here, the user may encounter a small register form, or in case the user is logged in, a link to the dashboard, i.e. the main application.

 A screenshot of the Thalia registration page. The title "Register" is at the top. It contains three input fields: "Username", "Password", and "Confirm Password". Below these is a red "Register" button. To the right of the "Register" button are links for "Already have an account?" and "Login". A horizontal line separates this from the bottom section.

 The bottom section is titled "Tweets" and shows a tweet from the user @Thalia99627941. The tweet reads: "Help pass the time during lockdown by perfecting your investment strategy with the #Thalia backtesting tool!". It includes a timestamp of "Apr 5, 2020" and a "View on Twitter" link. There are also "Load more Tweets" and "Embed" buttons.

 On the left side of the bottom section, there is a "Still with us?" section with a "Sign up" button and a "Log in" button. Below this is a "Already signed in?" section with a "Logout" button.

Figure 11: Top: User not logged in; Bottom: User recognised

Log In and Sign Up Pages

In case the user is not yet logged in, links for the ‘Log In’ and ‘Sign Up’ pages are visible on the navigation bar as seen on Figure8 or Figure9. Both of these forms are quite common, with the login requiring:

- Username
- Password
- (Optional) Remember me

And for signing up, the fields are:

- Username
- Password
- Confirm Password

As standard, the registration fails when the user enters different values to the ‘Password’ and ‘Confirm Password’ fields. In this case, the user is prompted to try again.

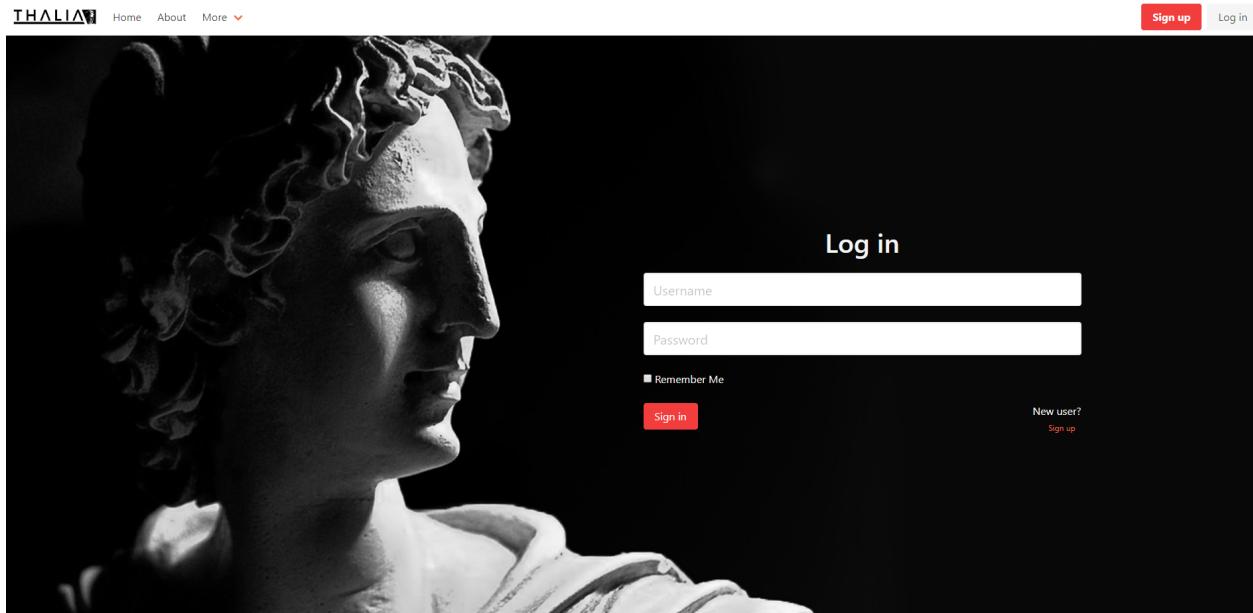


Figure 12: Thalia Log In Page (source: <http://thaliabacktest.xyz>)

In case the user is already logged in and attempts to access these pages, they will be redirected to the homepage. Additionally, the navigation bar changes, allowing us to log out.

5.3.2 The Dashboard

“Dash apps are composed of two parts. The first part is the ‘layout’ of the app and it describes what the application looks like. The second part describes the interactivity of the application [...]” [33]. Creating the Dashboard with Dash meant using Python classes for all of the visual components of the application.

Writing code for the layout was intuitive, but cumbersome, as on top of the HTML-like structure, it also needed to be correctly indented. Consequently, we needed to rely on refactoring a lot, chopping the layout code into smaller, more manageable components. In our final design, we have decided to divide the Dashboard into tabs, which helped us achieve a more intuitive look and organised codebase. The layout of the dashboard can now be initialised as follows:

```

1 import dash_core_components as dcc
2 import dash_html_components as html
3 from . import tabs
4
5 layout = html.Div(
6     html.Main(
7         [
8             html.Div(
9                 [
10                dcc.Tabs(
11                    [
12                        tabs.tickers(),
13                        tabs.summary(),
14                        tabs.metrics(),
15                        tabs.returns(),
16                        tabs.drawdowns(),
17                        tabs.assets(),
18                    ],
19                    id="tabs",
20                    value="tickers",
21                ),
22                ],
23                className="column",
24            ),
25            ],
26            className="columns",
27        ),
28        className="section",
29    )

```

Listing 1: layout.py - Example of Dash Code

Upon arriving at the Dashboard, only the ‘Ticker Selection’ tab is available to the user. This is where the user is expected to input their investment strategy. Although this process is fairly intuitive and done by using dropdown menus, input fields and standard date selectors, a more thorough walkthrough can be found in 11.

For the selection of assets we have decided to take an alternative approach visible on Figure13. This was done mostly because of the limitations of Dash discussed in 6.3.2, but came with its own benefits.

Ticker:



	AssetTicker	Name	Allocation
x	SPY		1
x	VTSAX		2
x	JPY		1

Figure 13: Thalia Dashboard - Assets Table

The official documentation describes the dash table as “an interactive table component designed for viewing, editing, and exploring large datasets” [34]. Working with the dash table is also quite simple, as its data can be accessed by addressing the id of the table and then the data component. For a further explanation on registering Dash Components, read 6.3.2.

Accessing the data like this also allowed us to fulfill one of our requirements, which is to support portfolios

with a large number of assets. This would have been significantly harder, if not impossible, with alternative approaches, see 6.3.2.

To compare investment strategies, the user may add portfolios via the ‘Add Portfolio’ button. Due to the arguments presented in 6.3.2, we have decided to set the maximum number of portfolios to 5, compared to letting it be dynamic. In addition, it is possible to select from a set of benchmarks, also known as lazy portfolios. Having selected one, the user will find the table populated with the desired assets and proportions.

When the user is content with the input, they can start backtesting by clicking the ‘Submit’ button. If all required fields are populated, the user is taken to the summary tab. This, as well as all other tabs, is now unlocked.

Showing the Results

Many studies have shown that one most important factor when designing a dashboard application is to show the right data using the right visualization tools [35] [36]. In our case, we have already established which metrics and plots are most important to the users. When the user lands on the summary tab, it is these key metrics that are presented, leaving a more detailed analysis for the relevant tabs. This way the user is not overloaded with information as soon as the results are presented.

Having decided on ‘What’ the user should see we can now focus on the ‘How’. In our final design, we combine data visualization techniques to make the dashboard look interesting. Key metrics are shown either in a ‘box’, which is a core building block of a dashboard or in tables. Proportions are visualised using pie charts, whereas relative differences are shown with the help of a bar chart. Some examples can be seen on Figure14.



Figure 14: The Dashboard

Among these charts, the user can find one of the key components of our application, the Portfolio Growth Plot. This graph is crucial for our application, as it serves as a visualisation for two of our main functional requirements, i.e. showing the total returns of a portfolio over time and the comparison of strategies. Thanks to Dash, all graphs are fully interactive. The user may zoom in on selected areas, hover over desired data points, save the plot as an image, etc.



Figure 15: Dash Functionalities

The result of performing one such action on the Portfolio Growth Plot can be seen on Figure16. In this case, the user is about to enlarge the selected region to inspect the plot closer. Resetting the graph can then be done by one of the buttons visible on Figure15.

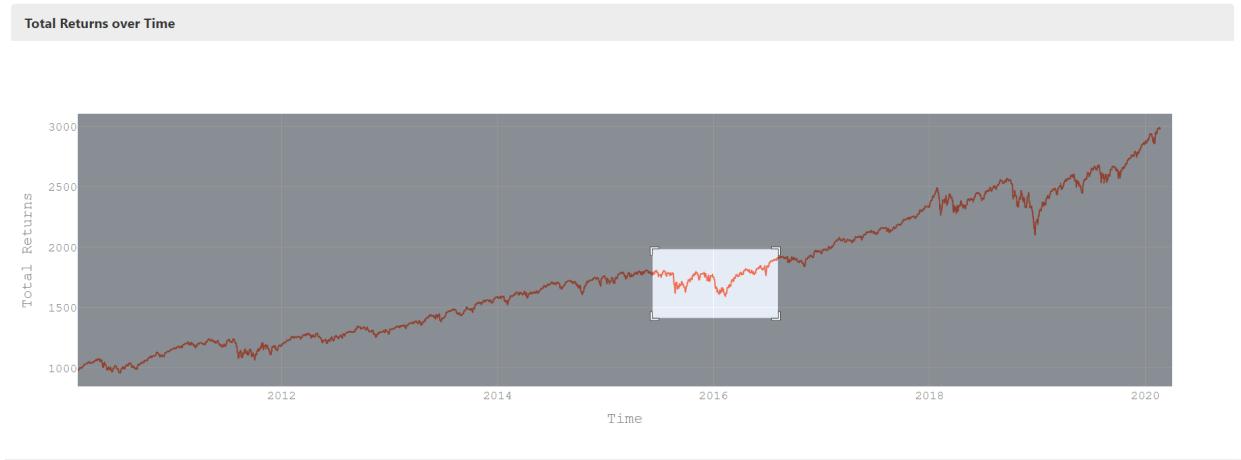


Figure 16: Dash Functionalities - Portfolio Growth

5.3.3 Changes to the Initial Design

The design of the Dashboard was done according to the wireframes we have provided last term [2], also visible on Figure17.

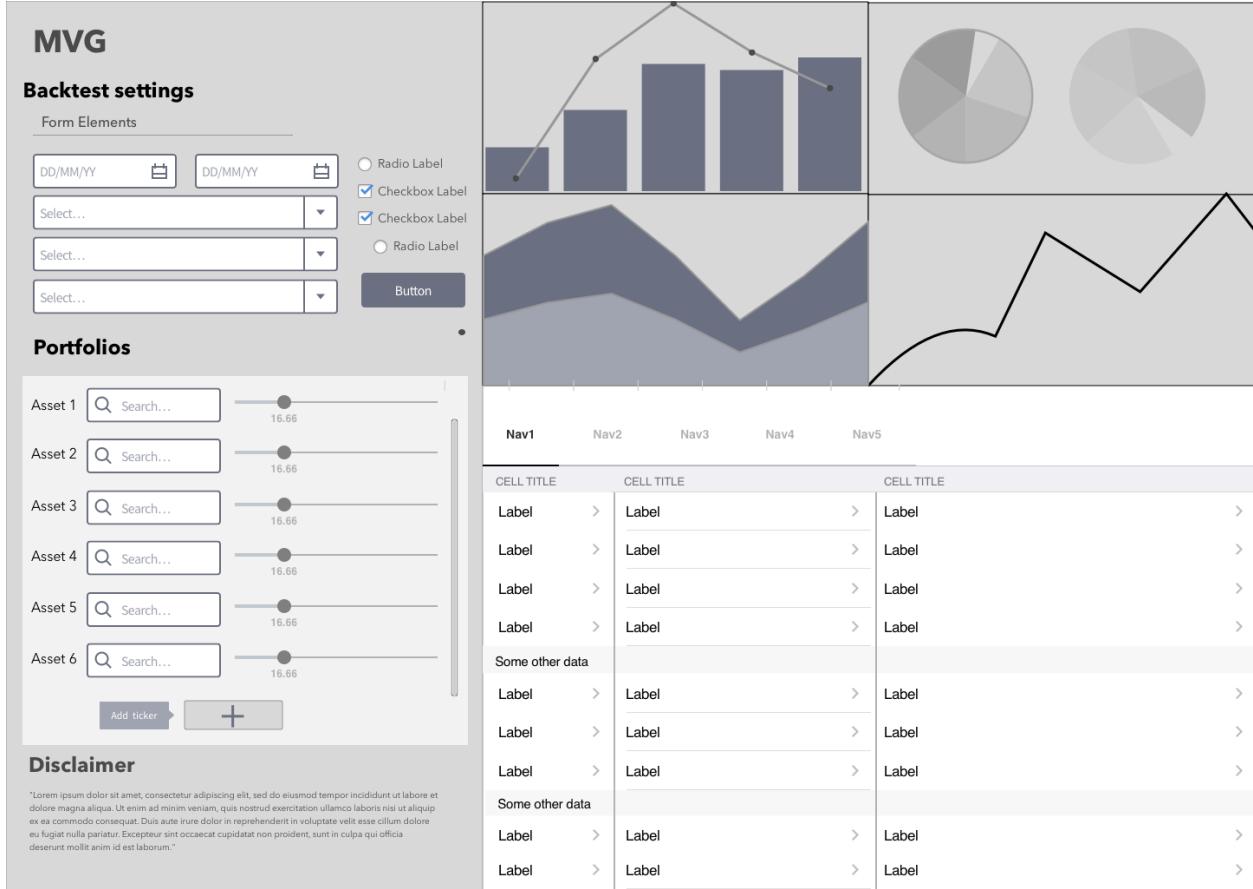


Figure 17: UI Wireframes

As we can see, we managed to encapture the overall look as we envisioned with only minor changes. The data is significantly better distributed by the use of tabs, giving a less cluttered UI. Additionally, we have decided not to implement the aside menu, partly because it was unnecessary, but mostly because it suggests that the backtest results can be changed on the go. This, as we have realised, was not feasible, due to expensive calculations in the Business Logic Layer.

One change we had to make was the use of the Dash table instead of the approach visible on Figure17. The benefits of this change have been illustrated further above. In addition, we have also decided to discard the sliders as a way of changing asset allocations.

For portfolio comparison, we also had to enable inputting multiple strategies. As seen on Figure17, this is something we have not accounted for when planning the design. The implementation of this feature led to a full refactoring of the UI code, as it came with an excessive amount of code due to the arguments presented in 6.3.2. Fortunately, this issue was addressed relatively early and resulted in the current design.

5.4 Business Logic Structure

The responsibility of our business logic can be summarised as follows: Given an investment strategy specification input from a user via the GUI, retrieve relevant financial data from the database to perform calculations for the historical performance and associated risk metrics.

To achieve this, we have developed a library (Anda 14) that performs the necessary calculations. Anda is decoupled from both the presentation and database layer by relying on external providers for any price data and the specification of a strategy. This decision has been made to allow for alternative sources of price data. One of our features is allowing users to input their own price data for assets not supported by Thalia. This data is uploaded as a CSV file. Without our current design, i.e. by coupling calculation of performance

and metrics to database access, we would have to modify the business logic to support multiple data sources. Given our current implementation, however, we can simply parse the user data into a pandas dataframe in a wrapper around Anda and then call functions within the library as required. Currently supported metrics include ‘Total Return’, ‘Max Drawdown’, ‘Best / Worst Year’, and the ‘Sharpe’ and ‘Sortino’ Ratios (please refer to 13 for definitions). However, the library is open for extension, hence additional metrics may be added at any point.

For a closer look at how an investment strategy is specified, consider the following class that serves as input to Anda library functionality (e.g. for calculating the Sharpe Ratio):

```

1 import pandas as pd
2
3 class Strategy:
4     def __init__(self,
5                  start_date: date,
6                  end_date: date,
7                  starting_balance: Decimal,
8                  assets: [Asset],
9                  contribution_dates, # implements __contains__ for date
10                 contribution_amount: Decimal,
11                 rebalancing_dates, # implements __contains__ for date
12                 ):
13         self.dates = pd.date_range(start_date, end_date, freq="D")
14         self.starting_balance = starting_balance
15         self.assets = assets
16         self.contribution_dates = contribution_dates
17         self.contribution_amount = contribution_amount
18         self.rebalancing_dates = rebalancing_dates
19

```

Listing 2: Strategy Object Specification

Here, ‘Asset’ is a simple dataclass consisting of a ticker string (e.g. ‘MSFT’ for Microsoft), a weight as a share of the portfolio overall (e.g. 0.25), a pandas dataframe holding historical price data ordered by date, and a pandas dataframe for dividends data (if any). An overview over the class relationships can be found in 15.

As alluded to earlier, functions within the library depend on a Strategy object for their calculations. For example:

```
1 def total_return(strat) -> pd.Series:
```

Listing 3: Example Function Signature

will calculate a series of total return values ordered by date within the date range specified in the passed Strategy instance.

Another important design decision has been the choice of data type to represent money, for example for price data. For this, we have chosen the Decimal type from the Python Standard Library ‘decimal’ module, since it “provides support for fast correctly-rounded decimal floating point arithmetic” [37]. As rounding errors and imprecision are unacceptable for our application, using the Decimal type will allow us to reliably compute figures for prices, risk metrics, etc.

Finally, we have chosen NumPy [38] for performing numerical calculations as this allows for highly optimized computation through the use of vectorized operations.

5.5 Data Harvester Structure

The Data Harvester’s role is to provide the application with the data it requires in order to allow the users to construct their investment portfolios. The choices we have made had to fulfil all our requirements while keeping in mind our financial and time-related limitations.

5.5.1 Choosing APIs

We started our quest for data by trying to find APIs that could we call to retrieve data the main financial assets used for constructing an investment portfolio. We have seen that there were several APIs that could

be used. After a quick look at them, we observed that there were two types of APIs. There were those that had FOREX data and those designed for Stock Market Data. The two APIs that we considered for FOREX data were Fixer and Nomics. For stock market data, we considered using yfinance, Alpha Vantage and Quandl.

FOREX API

Between nomics [39] and Fixer [40], we have chosen nomics due to the fact that nomics had unlimited free calls for free. Additionally, nomics also serves cryptocurrency data. However, Fixer is a more robust API with a larger user base and more data that we would have used if we would have been able to use their priced plans. To compare the two options, we have used the documentation provided by each API.

Stock Market Data

Out of yfinance [41], Alpha Vantage [42] and Quandl [43] we have chosen yfinance because it offered the most data and API calls for free, compared to the other two. If we had money to spend we would have used Alpha Vantage because of the larger number of available assets and the high number of requests per minute. As for Quandl we found it to be expensive and not centred on the type of data we need. We have made our choice based on the API specification given in the documentation of each API.

Standard Data Format

While looking for solutions that would help with solving the data format and the maintenance problem we have found pandas_datareader [44]. It is a wrapper for the biggest financial APIs. We saw that it included all of the Stock Market Data APIs named above and more. In addition, it standardizes the calling procedure for all APIs it includes and it returns the same format no matter the API you are calling. This is the only product that does this on the market so no comparison with other product can be done. In the case where we would have not used it we would have produced code that would have archived the same result. In conclusion, using pandas_datareader has saved us some valuable man-hours.

Update the database & Redundancy & Data Interpolation

For these requirements, we have seen that no already existing systems can do what we want. As a result, we had to write our system. In the building of the system, we have used python as a programming language. This decision has been made because the APIs chosen have been built in python. Writing the system in a different language while the APIs are in python would have created additional problems without any benefits. To manipulate the data we have used the pandas library. The persistent data required by the update mechanism is stored in CSV [45] format because of the readability and the ease of use when using pandas.

5.6 Database Structure

5.6.1 Data Segregation

The decision was made early on to horizontally partition the data stored by Thalia into two parts - one consisting of data related to users and user accounts and the other of financial data related to asset classes, assets and their historical prices. The following is the list of reasons the team documented for this decision:

- One alternative revenue stream we identified early on was the sale of our financial data as a separate product. This process would be trivial if it was stored in a separate database.

- Although the security of both types of data is important to our business model, protecting user's private information is the highest priority. The financial data is accessed by the Data Harvester, a separate program gathering data from many sources on the web and introducing additional security risks. Data segregation helps limit the scope of a potential data breach [46].
- The two types of data serve two separate purposes. The modules responsible for managing each are also decoupled. Thus, separation helps to enforce the principle of least concern.
- A large corpus of guides and examples on how to manage user accounts is available online. Extending any of these to include financial data might be difficult and risks leading to bad design.

The separation of dissimilar collections of data is a practice widely adopted in industry. Criteria for assessing when this approach is appropriate have also been documented [47]. Based on the decision to use SQLite as our DBMS and to maximize the portability and security of the financial data, we decided to use two separate databases.

5.6.2 The Data Layer Module - Finda

The Finda (for a definition see 14) module was designed to implement the data layer, acting as an intermediary between the Data Harvester / business logic and the financial data. It allows users to manage a number of databases implementing a common schema and give them access to a suite of tools for reading, writing and removing the data stored in each. In addition to this, the Finda module implements the following features:

- A system for managing user permissions to help reinforce separation of responsibilities among Thalia's other modules.
- Integrity checks to ensure the integrity of the data provided to the end user.
- A suite of administrative features to aid with managing the application on the back end.

Users can access these features through an outwards facing virtual interface, designed based on the facade design pattern [48].

Finda's design was modeled after object relational mappers (ORMs), which are libraries offered by most popular web frameworks. Although the implementation of what is essentially our own ORM proved to be costly in terms of developer time, it allowed us to create a more focused module tailored to our requirements. This helped to streamline the development of other modules.

UML diagrams for the Finda module can be found in 15.

6 Coding and Integration

After a brief overview of the project and the project plan, this section will focus on the main technologies used in the project and the rationale behind choosing them. Moving on, we will discuss how these components were integrated and eventually deployed.

6.1 Overview

One of the first decision we have made this term was to completely recreate our prototype of Thalia. Firstly, this radical move was the consequence of a new implementation decision (for more detail see 6.3.1). Secondly, as prototypes are meant to be disposable and are designed only to answer key questions about the system [49], the proof of concept served its purpose, giving us a chance to refine the structure and the quality of the code.

Despite the risks posed by using a Software Version Control (SVC) Host such as GitHub, we have decided to continue using it as our software development platform. The reasoning behind this builds on the argument developed in our Technical Report [2], which highlights that our Data Processing Module is a completely separate component in our system which none of the other components is able to access. Additionally, we are

storing API keys as environment variables in a file that is not tracked by our SVC system, which minimizes the probability of us exposing sensitive data.

We have also decided to develop the application with Python as our main choice of programming language. Even though this choice seemed obvious from the beginning, we did consider its main benefits, which are as follows:

- Python is a high level programming language allowing us to better focus on the application.
- A standard choice for prototyping.
- Provides superb third party libraries and frameworks for free.
- Easy to integrate if we were to choose other languages at some point in our development.
- The whole team was already familiar with the language, saving us the precious time needed to learn another programming language.
- Our application does not require an unreasonable amount of computation, so there is no need for a more efficient programming language such as C 6.3.3 [50].

We will discuss other technologies used in more detail after the discussion on project planning.

6.2 Planning

Early in the inception phase of development, we decided that our goal was not to have fixed responsibilities in our team, allowing everybody to work on each component of the system. This decision has also eased the code review process, as we had no status differences to distort the error-correcting mechanism [51]. Furthermore, since no team member was the sole developer of a system component, this allowed us to direct comments at the code and not the author [52]. For these reasons we have decided to follow the Egalitarian Team structure and make use of the flexibility offered by it.

Our workflow was centred around the tools provided by GitHub. We used a ticketing system to divide and distribute tasks among team members. These tickets were sometimes given by the team on the weekly meetings, but occasionally they were chosen by the member proposing the feature or change. Our goal with this approach was to divide larger jobs into smaller tasks.

A typical ticket in our project was an encapsulation of a user story, which consisted of a title, a one liner, value, acceptance criteria, and sub tasks. In the first half of development, we also used effort-oriented metrics (story points) to measure the amount of work incurred by taking on a ticket, but we decided to abandon this aspect. An example of a ticket can be seen on Figure18.

mara42 commented on Nov 17 2019 • edited

+ ...

As a user, I want to see the results of a backtest, so I can use the tool.

Value

- Working product

Acceptance criteria

- Does pressing submit on the portfolio page return results of the backtest?
- Do the results contain a table of key metrics?
- Do the results contain a graph?

SP:?

Subtasks:

- Send user input to backend
- Send backtest data back to frontend
- Create graph
- Create table

Figure 18: Ticket Example (source: <https://github.com/>)

Throughout the whole project, we had a total of 110 tickets, and 71 pull requests. The following graph also shows the number of contributions to master, excluding merge commits.

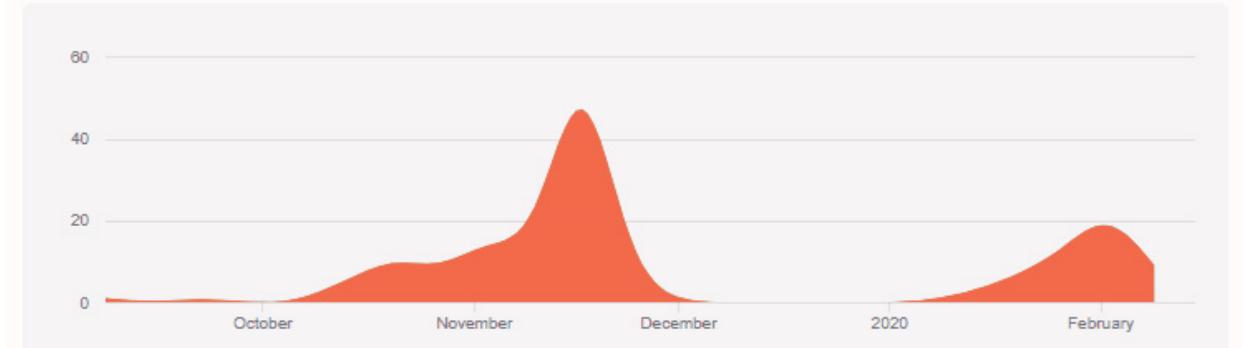


Figure 19: Contributions to the Master branch (source: <https://github.com/>)

Nevertheless, this rigour in creating tickets was not always achievable. In particular, it was difficult to create small enough tickets in the beginning of the development process when more fundamental components of the systems were developed. In these cases, we assigned the ticket to a pair or group of people. This approach achieved the following:

- Improved the overall code quality and fastened production [53].
- Minimised review time over the long run.
- Distributed the knowledge of large system components across a few people instead of one person.
- Eased introducing the new team member to the project.

In addition, we also had a scrum board as an overview for the ongoing tickets. Although the Scrum community engages in ongoing discussion about the benefits of a physical scrum board over an online one

[54], given no actual workplace, the former was not possible to achieve. The board allowed us to see which tickets needed to be reviewed and which were ready to be merged. The tickets/cards were distributed into columns, such as ‘To do’, ‘In progress’, ‘Review in progress’, ‘Review complete’ and ‘Finished’. A truncated picture of this scrum board can be seen on figure Figure20.

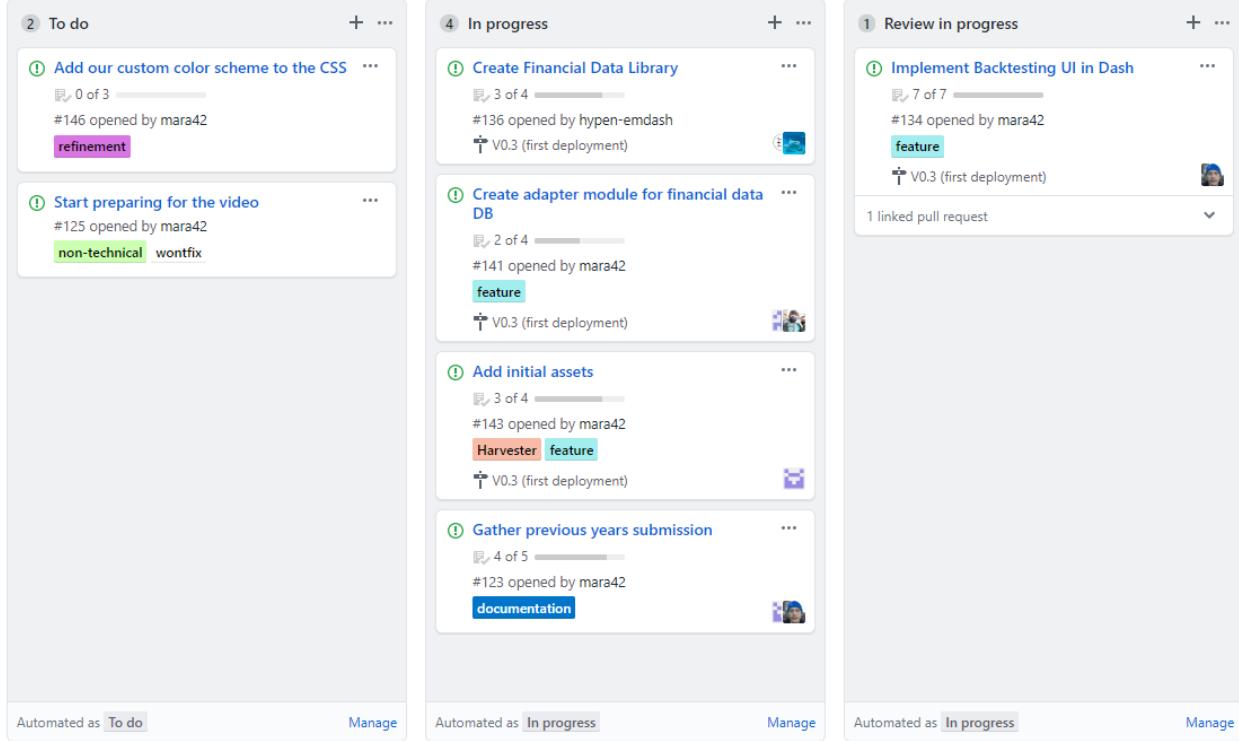


Figure 20: Scrum board - truncated

The workflow then largely depended on the task at hand and the person working on it. It was up to them if they met for pair-programming, or chose to work individually. In the beginning, we all agreed on the developing environment and coding standards (for a detailed analysis, see 6.4). Since this report was a relevant portion of the workload, we decided to treat it as code. The report was written in L^AT_EX with an overall structure of the document in a main.tex file which linked to individual sections. This let us to work on different sections separately, just like features in our software.

Regardless of the nature of the ticket, the week, or in case of some larger tasks, two weeks, ended by the creator indicating that the changes were ready to be integrated. This was done by publishing the changes and creating a new pull request, which - upon approval - pushed the changes onto the master branch. To ensure code quality, we set up a continuous integration (CI) environment (more on that in 6.4). After all checks have passed, the changes were successfully integrated into the code base.

6.3 Key Implementation Decisions

Let us now discuss the main technologies used in our project. These can be categorised as follows:

6.3.1 Web Framework

Choosing the right web framework was more controversial than originally planned. The two major options for Python are Django and Flask. Although we decided to use Django for our MVP in the first semester, we had to spend a significant portion of the time available learning the framework, so we had to decide whether we were going to stick with it or learn Flask. In the end, we opted to go with Flask for the following reasons:

- Django has one architecture that all projects must share, and we have designed the architecture for our project ourselves. While neither architecture is wrong, the two are not compatible. Flask, on the other hand, is structure-agnostic, so we can lay out the code as we see fit.
- Flask comes with the bare minimum for web-development, which means that we don't need to manage the complexity of any feature we're not using. Django has a more complete feature-set from the beginning. This would be desirable in a large web application, but introduces significant overhead in our case, where the website has only a handful of pages.
- Django all but insists on using its ORM for all database interaction, while we plan to have a more manual approach through the use of Finda.

Thus, we chose Flask, a Python-based micro web framework. Flask provides functionality for building web applications, such as managing HTTP requests, using the ‘requests’ library, and rendering templates [55]. Flask leverages ‘Jinja2’ as a template engine [56]. We used the ‘render-template’ library to process and parse the HTML templates for all the pages, except for the dashboard, which is described in more detail below. We used some of Flask’s extensions, such as Flask-Login, which provides user session management for Flask [57]. To handle web forms we used the ‘flask-wtf’ extension, which is a wrapper around the ‘WTForms’ package [58].

6.3.2 GUI

Dash is great for applications that require data visualisation, modelling, or analysis [31], which is precisely what is needed for our portfolio backtesting software. It allows us to create reactive single-page apps, which means that with the use of tables, drop-down menus, and other forms of input, the app is updated instantly. This is done with the use of ‘callbacks’ [59]. Callbacks allow us to exchange data with our web server asynchronously through AJAX requests whenever a user input changes. Upon receiving a response to such a request, the corresponding output is updated immediately without refreshing the page.

A typical callback using an example from our application can be seen below:

```

1 def register_add_portfolio_button(dashapp):
2     dashapp.callback(
3         Output("portfolios-container", "children"),
4         [Input("add-portfolio-btn", "n_clicks")],
5         [State("portfolios-container", "children")],
6     )(add_portfolio)

```

Listing 4: setup.py - Development environment

In this code snippet, we register a listener to the button with the unique id: ‘add-portfolio-btn’, and its property ‘n_clicks’, or number of clicks. Whenever the input property changes, the corresponding function gets called automatically. This also happens upon launching the Dash page. Therefore, it is common to start the called function by catering to this case and raising a ‘PreventUpdate’ exception if the input parameter is ‘None’. An example can be seen in [60].

The called function then receives the input properties as parameters. Furthermore, we can add any number of parameters as ‘State’ objects, as has been done above. When returning from a function call, the same number of values as is specified in the callback must be returned. Mismatching the number of return values results in a runtime error. Mismatching the input parameters, on the other hand, yields a different function signature, and so the callback is not registered with the desired function. Doing this does not raise any errors.

Fortunately, Dash offers great developer tools [61], such as the ‘Callback Graph’ visible on Figure21, which can be accessed by running the application in debug mode and activating the Dev tools. This graph came in handy often, especially due to the arguments presented in 6.3.2.

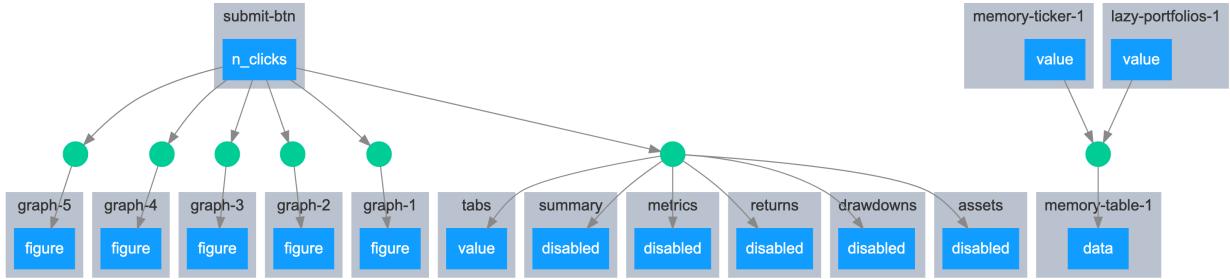


Figure 21: Callback Graph on Thalia

Limitations

Although Dash is fairly intuitive and great for one-page applications, it has its limitations when it comes to dynamically creating components. To be more precise, all callbacks need to be registered before runtime. In our application the user is able to enter multiple portfolios via the ‘Add Portfolio’ button. Doing so creates new input fields for the portfolio, which all need to be registered.

As this problem persisted with other components as well, we needed to find a general solution. Many approaches have been suggested by the Dash Community, some even by the developer team [62]. In the end, we have decided to use the most stable approach and generate all the components beforehand and hide them until activated. This means that the relevant function now also has the responsibility of returning the visibility of the required div element. Note that this is a highly requested feature in the Dash community, and is likely to be added soon, but for now, we needed to resort to workarounds.

Furthermore, the data we show in the dashboard is greatly centered around the strategy object discussed in 5.4. At the beginning of every backtesting process, a new strategy object is initialised, which is a time-costly task. As this object holds most of the information our result pages require, many functions would need to access the same object.

One solution would be to have a callback for initialising the object, and then share it with all others. Although there exist ways to do so [63], we were not satisfied with any of the approaches listed. Consequently, we decided to go forward with the original approach that Dash favours, and combine more outputs in one function. Even though this introduced more complexity to our code, it made the application faster overall.

Layout

The layout of the application, as discussed in 6.3.2, has been created with the use of the ‘Dash HTML components’ library. Consequently, there was no need to write HTML for the dashboard, since the layout was composed using Python, which then generates HTML content. Dash uses the Flask web framework, which makes it fairly simple to embed a Dash app at a specific route of an existing Flask app [64]. The framework also handles most of the communications between the front-end and back-end.

The User-Interface was made with the use of ‘Bulma’, which is an open-source CSS framework [30]. Bulma is ideal for creating responsive elements, as it automatically adapts the website to different screen sizes [65]. It is also compatible with various web browsers. Another advantage of Bulma is that we could create a very modern-looking website quickly since the framework is doing a lot of the heavy lifting for us.

It also provides pre-styled elements, components (such as the navigation bar) and alerts for the users. Additionally, it is intuitive to learn when compared to ‘Bootstrap’, as Bulma uses memorable class names, such as ‘.button’, and has a very straightforward modifier system. It consists of nothing but CSS stylesheets without any use of JavaScript, which makes it very compatible with a framework like Dash.

6.3.3 Business Logic

The business logic module sits at the very core of our application. We require it for producing a time series of a portfolio’s performance as well as selected key risk metrics of a given asset allocation. This output is consumed by our web application and presented in plots and tables as described in 6.3.2.

Research into developing this module was initiated by listing our requirements for its desired behaviour. We identified that it should support:

- Specification of a portfolio as a set of pandas dataframes, the common data exchange format in our application.
- Calculation of a portfolio's return over time.
- Regular contributions to mimic saving.
- Rebalancing strategies to reestablish the desired weighting of a portfolio.
- Calculation of key metrics, including the Sharpe and Sortino Ratio, Max Drawdown, Best and Worst Year.
- Collecting and reinvesting dividends for equities.

Using these requirements, we struggled to find an open-source library that would handle these tasks for us. While backtesting libraries written in Python are available in abundance (e.g. PyAlgoTrade [66] and bt [67]), each of these was lacking in at least one critical aspect. None of them support specifying a portfolio using absolute or relative weights and instead seem to focus on backtesting trading strategies involving just a single asset while using technical indicators. Thus, we made the decision to develop our own library for handling the aforementioned tasks.

The result of this effort is ‘Anda’. For each of its functions, Anda takes as input a ‘Strategy’ object that specifies the entire list of parameters for a backtest, including contribution dates, a list of assets with associated price data, dividends for equities, etc. Calculations are performed on a per-metric basis by separating them into individual functions. This approach has allowed us to tailor the entire business logic module exactly to our needs without having to produce complicated wrapper code for existing backtesting libraries.

6.3.4 Database and Finda

Another major technology decision was the choice of appropriate database management system (DBMS) for storing historical price data collected by the data harvester. Before committing to a specific technology, we identified the following requirements a suitable DBMS should fulfil:

- **Schema:** The structure of our data is relatively simple. Consequently, Thalia does not require support for sophisticated features and data types. A suitable DBMS should be able to accommodate the database schema designed last term, with the addition of simple integrity constraints and cascade operations.
- **Support:** Ideally, the DBMS should be cross platform, as this would allow us to defer commitment to a specific deployment platform until we are ready to start the CD process.
- **License and pricing:** The DBMS should be free to use and have a non-restrictive license.
- **Performance:** The DBMS should be able to handle a high volume of concurrent reads to fulfill user requests. The data will be updated daily, meaning efficient write operations are a lower priority.
- **Usability:** As our team lacks experience in this field, a suitable DBMS should be relatively simple to learn. Ideally, team members should be able to learn the basics in a single week long sprint.
- **Security:** The DBMS should have a mature code base and be relatively secure, as access to financial data is a key component of our business model. Later, it will likely also store data that is not available through public APIs, meaning potential data breaches could expose us to legal liability. [68]
- **Type:** Since the project constraints specify we use SQL queries, only relational DBMSs supporting a version of SQL are appropriate.

MySQL, PostgreSQL, SQLite and MariaDB were subject to in-depth comparison based on fulfilment of the above requirements and industry adoption [69] [70]. Our final decision was to use SQLite for the following reasons:

- It is user-friendly and easy to deploy, allowing us to start continuous deployment faster.
- It has a small footprint and offers good performance [69].
- Portable serverless design aids with development and testing.
- All team members have experience working with SQLite from previous term. This helps to reduce the overhead of knowledge transfer.

The main drawbacks of using SQLite, namely scalability and performance, are not a concern at this stage, as the current version of Thalia is meant to be a high-quality industrial prototype. As such, it will not contain the full range of financial data needed for marketability. Should SQLite prove to be inadequate in the future, we would be able to switch to a different DBMS with relatively little trouble, as the process of database migration is exceedingly well-documented [71] [72]. To pre-empt any difficulties that might arise, the decision was made to design the data layer to easily accommodate such a migration.

6.3.5 Data Harvester

The role of the Data Harvester is the collection of live data from multiple source, parsing it to a format that makes it usable by Anda and easily storables by our in house DBMS system Finda. In addition to this, a mechanism for changing the data collected and adapting to the market demands has to exist. Another requirement relates to the limitations that are imposed by the third-party APIs we use. The first design decision made in order to accomplish the above was the creation of the updating mechanism.

All financial backtesting applications rely on having accurate data to work with. We required a system that can aggregate the data from multiple APIs and that would keep our database updated with live data. The requirements for such a system were:

- **Standard data format:** To be able to store the data into a database with Finda and have the financial analysis done by Anda, the Harvester needs to parse all the data received through the APIs to a standard format without losing numerical accuracy or granularity.
- **Update the database:** The database has to be updated regularly to contain live data.
- **Data Interpolation:** Some financial assets require interpolation to account for days without any transactions (e.g. weekends and public holidays).
- **Redundancy:** We have no power over the APIs we are using. The required system needs to have redundancy measures so that we can retrieve live data even if an API fails.
- **Maintenance:** The financial world is in a continuous change. Because of this, it is important to make sure that we can add and delete assets as time goes on.

6.3.6 Continuous Updating Mechanism

The Updating consists of three components:

- **The update procedure** made out of a series of methods. The procedure starts by looking at what APIs are available to be called. Each API has an update list that contains the tickers of the financial assets that we wish to get data on and the last and earliest record we have on that asset. If it is the first time when asking for data on a specific asset then we ask for all data available between 1970/1/1 up until the last trading day available. If the number of allowed calls for an API is greater than the number of assets to be updated, then all assets can be updated in a daily run. However, if it is not possible to update all assets in one update run, then we will update the rest of the assets in the following run so that we can get as close as we can to live data.

- A `run_updates` script that starts an update round. It contains the configuration with the API names, the location of the access key (if it exists) and the number of calls to be done for each API.
- An **initialiser** module that is based on the persistent data directory which contains the asset classes and the assets in the form of CSV files. The initialiser gets a configuration similar to that given to the ‘`run_updates`‘ script. Based on the asset classes, the assets and the APIs given, the initialiser creates an ‘`update_list`‘ for each API.

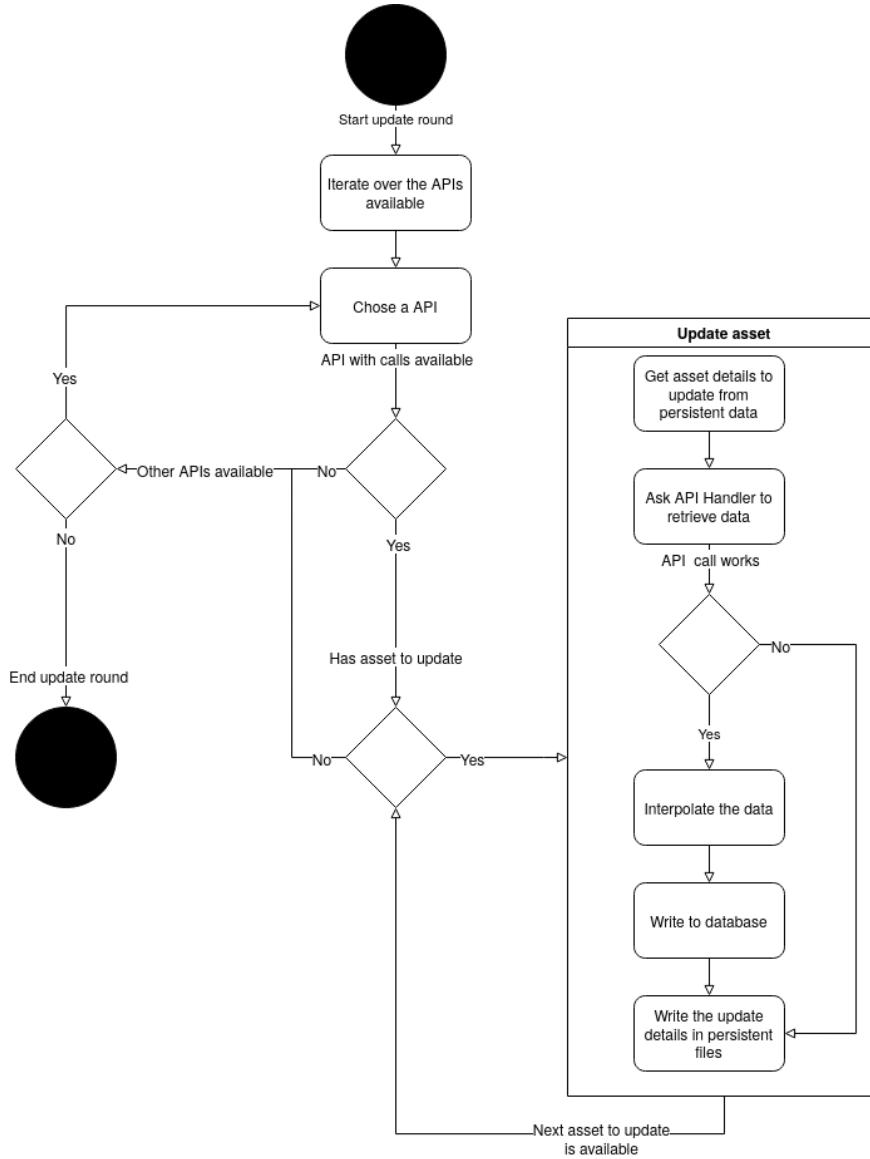


Figure 22: Data Harvester Update Round [2]

To know more about how the updating mechanism works, please refer to 12.0.1.

6.3.7 API Handler

For each API, the API Handler has a ‘call‘ method and a method that formats the data so that it fits our data model. Additionally, there is an ‘`api_selector`‘ method that makes the right API call, based on what

asset class is requested and a data frame format checker that does a final check before sending the data to the updating mechanism.

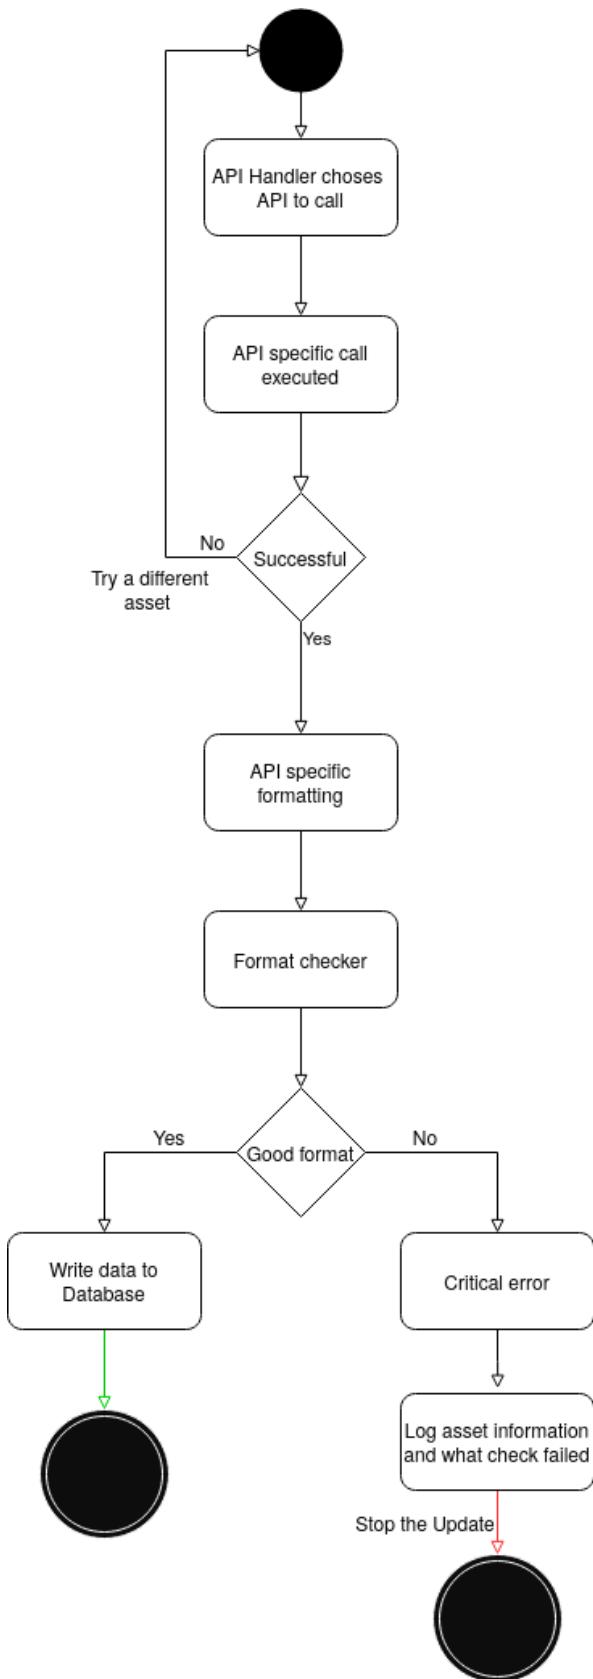


Figure 23: API Handler Flow Chart [2]

6.3.8 Standardization of Data

All data written to the database has to be in the following format:

```
1 Columns: [AOpen<Decimal.Decimal>, AClose<Decimal.Decimal>,
2           AHigh<Decimal.Decimal>, ALow<Decimal.Decimal>,
3           IsInterpolated<Integer>]
4
5 Index: [AssetTicker<String>, ADate <datetime.date>]
```

Listing 5: Pandas Data Frame Format

The data standardization is conducted as soon as the data frame has been received from the API. All APIs need to have a method that takes the data frame received and modifies it so that it abides by our standard. To further enforce our standard, we have also added a format checker in the API handler. It can be used to verify any newly added API or to catch any data frame modifications done by the API owners before they get written in our database.

6.3.9 Interpolation

Our business logic library requires continuous data in order to be able to perform its numerical analysis. In order to transform real-life data that includes weekends and bank holidays into a continuous data set, we have used interpolation. The type of interpolation used is nearest neighbour. To be more precise, the algorithm designed goes over the data in pairs of adjacent dates. If there are any missing days, then the older nearest neighbour is duplicated over the missing days. In addition to this, the duplicated rows are flagged as interpolated in the database. Here is a example of how it may look like:

Pre-Interpolation:

Index	ADate	AHigh	ALow	AOpen	A Close	AssetTicker	IsInterpolated
3	2000-11-16	126.86	126.89	126.86	87.36	VFIAX	0
4	2000-11-17	126.44	126.44	126.44	87.07	VFIAX	0
5	2000-11-20	124.12	124.15	124.15	85.47	VFIAX	0
6	2000-11-21	124.55	124.45	124.55	85.78	VFIAX	0

Post-Interpolation:

Index	ADate	AHigh	ALow	AOpen	A Close	AssetTicker	IsInterpolated
3	2000-11-16	126.86	126.89	126.86	87.36	VFIAX	0
4	2000-11-17	126.44	126.44	126.44	87.07	VFIAX	0
5	2000-11-18	126.44	126.44	126.44	87.07	VFIAX	1
6	2000-11-19	126.44	126.44	126.44	87.07	VFIAX	1
7	2000-11-20	124.12	124.15	124.15	85.47	VFIAX	0
8	2000-11-21	124.55	124.45	124.55	85.78	VFIAX	0

6.3.10 Security

To identify possible vulnerabilities, we have created a scenario where an attacker is trying to disrupt our flow of data - or worse, gain control of our updating process. During this phase, we have identified three security concerns that will be listed below.

Security Concerns

1. **API calls:** The data received from an API can be intercepted and modified.
2. **API keys:** The keys used by each API have to be stored safely.
3. **Rogue API:** An API can send a malicious package instead of the data we are expecting.

Solutions

1. **API calls:** Depending on the API settings, we will use either an IPSec policy or a VPN connection. The details of this will be discussed with the financial data provider.
2. **API keys:** We will salt and then hash the API keys with SHA512 before storing it. The number of financial data APIs on the market is low enough to allow us to use SHA512 without having to worry about computation time.
3. **Rogue API:** We have added several checks that make sure that the data we retrieve from the APIs are what we expect it to be. If an API sends malformed data, then a critical error will be displayed and logs of what did the API sent will be saved for inspection.

6.4 Integration and Deployment

Writing well documented and good quality code is one thing, but making sure it all works together as a whole is a completely different story. In the first term, many hours have been wasted on trying to integrate different components of the system which did not want to fit together. Even then, we had some DevOps tools in place [73], but considering that we had to produce significantly less code and more documentation last term, this was not a priority.

Starting afresh the coming term, we have decided to set up the development environment again. Upon opening the ‘setup.py’ file, we see a list of required libraries, and the following lines of code:

```

1 """A setuptools based setup module."""
2 from os import path
3
4 from setuptools import find_packages, setup
5
6 here = path.abspath(path.dirname(__file__))
7
8 install_requires = [
9     "flask",
10    "flask-login >= 0.5",
11    "flask-migrate",
12    "flask-wtf",
13    "pandas",
14    "dash",
15 ]
16
17 tests_require = ["pytest", "coverage"]
18
19 extras_require = {"dev": ["black", "flake8", "pre-commit"], "test": tests_require}
20
21 setup(
22     name="Thalia",
23     version="0.2.0",
24     packages=find_packages(),
25     install_requires=install_requires,
26     extras_require=extras_require,
27 )
28

```

Listing 6: setup.py - Development environment

One of the first decisions we had to make was to choose a standard coding style. This is exactly what ‘flake8’ is for, which we can see amongst the extras in the code snippet above. The original documentation of flake8 defines it as “[...] a command-line utility for enforcing style consistency across Python projects. By default it includes lint checks provided by the PyFlakes project, PEP-0008 inspired style checks provided by the PyCodeStyle project, and McCabe complexity checking provided by the McCabe project“ [74]. However, as many other developers, we also decided to redefine the maximum line-length from 79 to 88 as we found this convention to be a hindrance when writing code.

Another tool used for enforcing uniform style was the ‘black’ auto-formatter for Python [75], which formatted the code for us upon every save (if enabled) and when committing code to GitHub. This has been achieved by the use of githooks [76], which are programs that are triggered upon certain git actions. Setting

up these hooks required pre-commit package. This ensured that both flake8 and black were guaranteed to run before publishing changes.

6.4.1 Continuous Integration

Many studies have investigated the positive effects of developing in a continuous integration (CI) environment (see, for example, [77] and [78]). Regardless of the exact implementation, its obvious benefit is that it provides security and uniformity for projects. We already made some steps to achieve a uniform style, but had no means to know whether the code published has been also passed its tests. Note that, in this section, we will only discuss testing as a part of CI and not the testing strategy (for this, see 7).

Another important part of the DevOps toolchain is the use of containers, which is what Docker helps to achieve [79]. Docker helps developers focus on writing code rather than worrying about the system the application will be running on and also helps to reveal dependency and library issues. As Docker is open source, there are many free to use docker images available online [80]. When choosing the CI environment, Docker support was one of the main requirements.

The most promising candidate for this was CircleCI [81], which is a cloud-based system with first-class Docker support and a free trial. After connecting our GitHub repository to CircleCI and setting up a configuration, CircleCI now does the following on every pull-request:

1. Sets up a clean Docker container or virtual machine for the code.
2. Checks previously cached requirements (for more detail see Figure24).
3. Installs the requirements from ‘requirements.txt’
4. Caches the requirements for faster future performance.
5. Clones the branch needed to be merged.
6. Runs flake8 one last time and saves results.
7. Runs tests and saves results.
8. Deploys the master branch to AWS, see 6.4.2

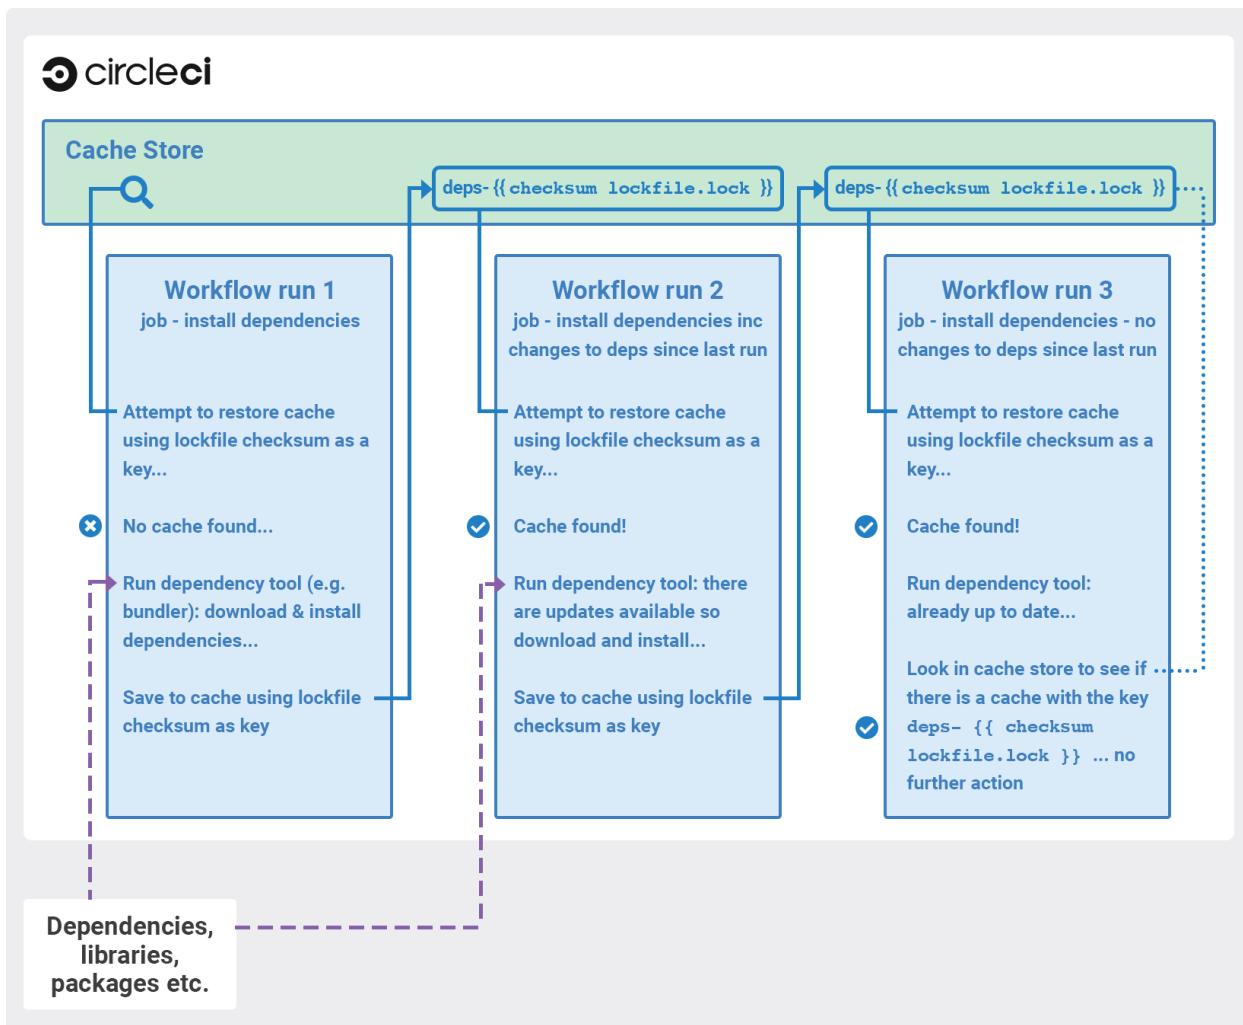


Figure 24: CircleCI on Caching (source: <https://circleci.com/docs/2.0/caching/>)

The outcome of these steps is visible on the CircleCI web platform, but - more importantly - also on GitHub. CircleCI will prevent a merge with master if failing tests (or no tests) have been found.

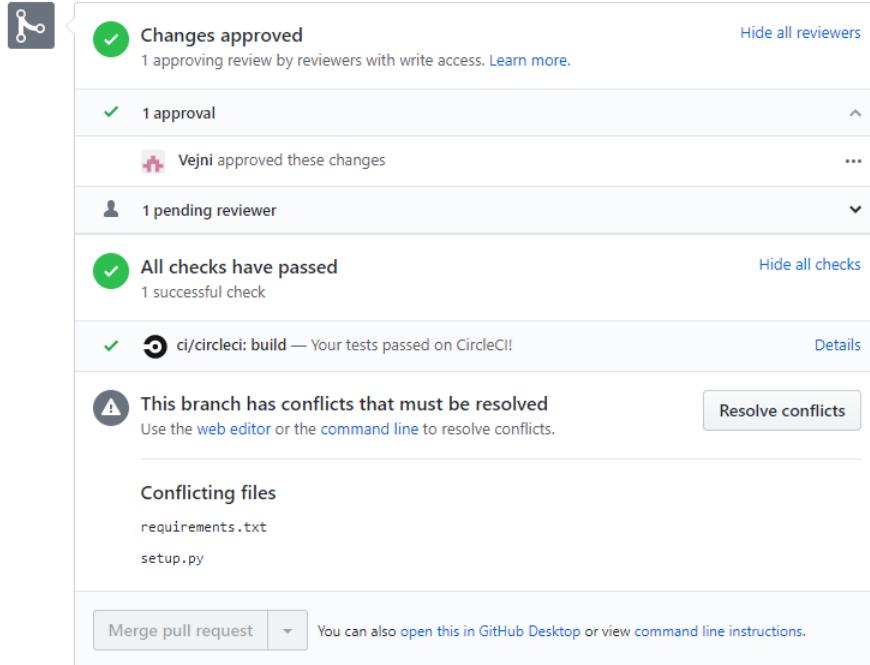


Figure 25: CircleCI on GitHub (source: <https://github.com/>)

With these steps, we managed to ensure that the code written is uniform and of good quality. It significantly reduced the time needed for integrating and code reviewing. The last step was now to deploy the system.

6.4.2 Hosting and Continuous Deployment

An overview over the literature covering deployment reveals a plethora of different strategies for deploying and hosting web applications [82]. Our decision of how to choose among them involved the following considerations:

- Price - since we have severe budget constraints, we were looking for a cheap hosting solution
- CircleCI support - The target host should be supported by CircleCI natively to ease development of a continuous deployment (CD) pipeline

The upfront cost of buying physical servers ruled it out as an option for us. Thus, we turned our attention to using a solution that involved deployment to a virtual machine in the Cloud.

Due to native CircleCI support and a free-tier service, we initially chose Heroku [83] as our initial hosting provider. This allowed us to host our application for free in the initial stages of development while providing ample opportunity for horizontal and vertical scaling later on, if required.

However, it quickly became apparent that Heroku was not appropriate for our purposes, due to its lack of support for sqlite databases [84]. Hence, we moved to AWS [85], which is also supported by CircleCI [86].

The benefits of using continuous deployment have been well established for multiple years and involve “the ability to get faster feedback, the ability to deploy more often to keep customers satisfied, and improved quality and productivity” [87]. Using AWS in combination with CircleCI, our CD pipeline involves the following simple steps:

1. Upon commits to the master branch on GitHub, CircleCI triggers a workflow.
2. The workflow first executes the steps listed in 6.4.1 to ensure the validity of the current codebase state.

3. If this step is successful, the master branch is pushed to a remote repository recognized by AWS via git.
4. AWS executes the ‘Procfile’ script stored in the root of our project to start the application using a ‘gunicorn’ web server [88].

Our deployment process is thus fully automated and immune to failing tests, as it will only complete successfully if the application is in a correct state. The full CI/CD workflow is captured by the flowchart on Figure26.

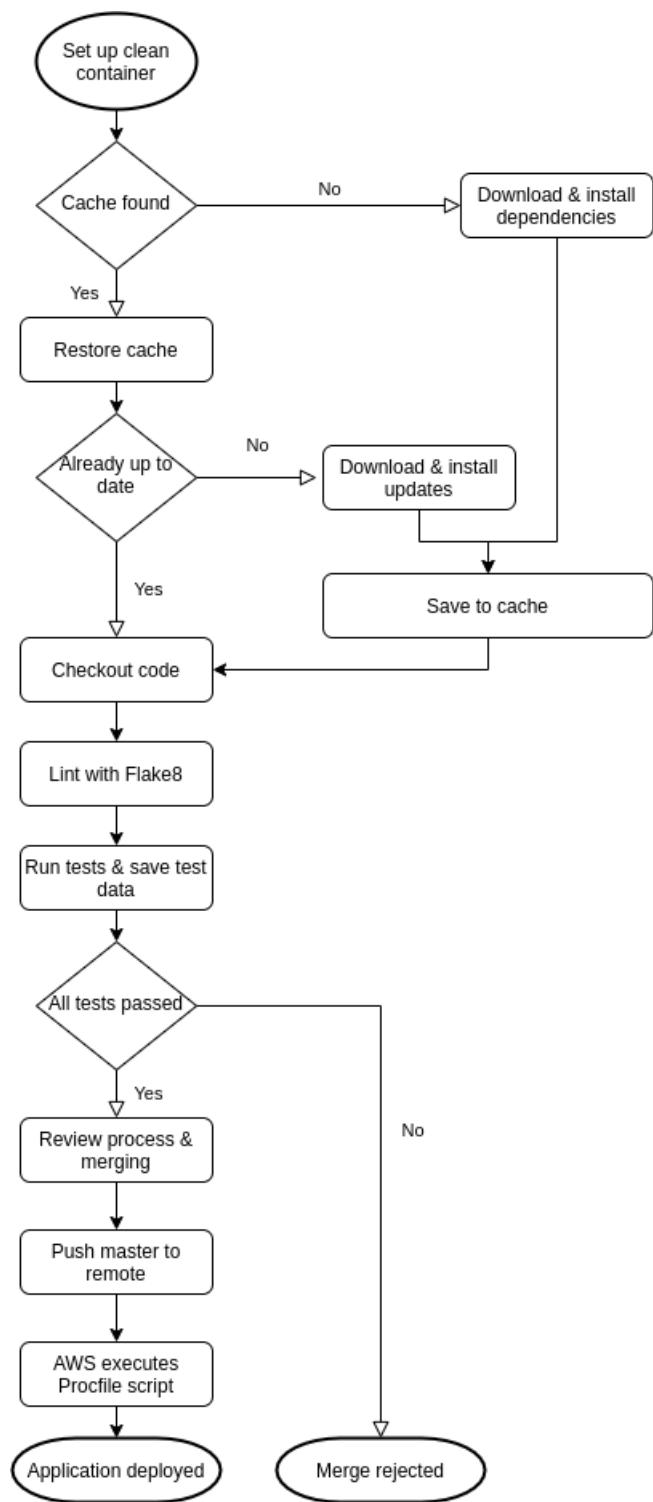


Figure 26: CI/CD workflow

7 Testing

7.1 Testing Strategy

Our initial testing strategy devised during the first semester remained mostly unchanged throughout development. Since we opted to redesign most of our first semester's prototype from scratch, we were also forced to recreate our testing suite and test data. This ultimately proved to be beneficial, as we able to reassess and improve the design of our tests. We also opted to spend more developer time on testing so that we could achieve high overall test coverage to help increase code quality, reliability and maintainability [89]. Additionally, we used this opportunity to integrate new tools into our testing suite to aid with comprehensive testing (last semester we only used our testing harness pytest [90]).

Our testing strategy so far consisted of three main types of tests: unit, integration and end-to-end tests. We opted to use a mix of white and black box testing techniques for the first two and conducted our end-to-end tests purely as black-box tests. In addition to this, the Data Harvester was designed to perform tests on the data gathered from APIs at runtime and record their results in a series of log files. The majority of our tests were automated and included in our pytest testing suite, although a minority was also performed by hand or with the use of auxiliary code. Some examples of this would be hand testing of API constraints, as this is a task that only needed to be performed once, and timing of different SQL queries to determine which is faster.

```
Tests/
|-- Finda
|   |-- conftest.py
|   |-- finData.db
|   |-- helpers.py
|   |-- test_fd_manager.py
|   |-- test_fd_read.py
|   |-- test_fd_remove.py
|   '-- test_fd_write.py
|-- Harvester
|-- Thalia
|   |-- conftest.py
|   |-- test_callbacks.py
|   |-- test_factory.py
|   |-- test_user.py
|   '-- test_views.py
|-- analyse_data
    |-- test_anda.py
    '-- test_data
        |-- AAPL.csv
        |-- AAPL_dividends.csv
        |-- BRK-A.csv
        |-- MSFT.csv
        '-- MSFT_dividends.csv
            '-- risk_free_rate.csv
```

Figure 27: Tree of files in testing suite

7.1.1 Unit Testing

All team members were required to write automated unit tests for any code they wished to merge into the project's master branch on GitHub. These were designed to test individual component modules and work independently from each other. The running of automated unit tests was performed by our CircleCI workflow. Although we are currently constrained by our budget, in the future, our team plans on using the advanced metrics provided as part of CircleCI's paid plan to gain insight into how testing was performed throughout development. As part of the code review process, team members reviewed each other's tests as well as production code. This helped us guarantee that unit tests were exhaustive and tested the full range of input classes and boundary cases for each software component. When designing these tests, team members also used white box testing techniques to maximize test coverage and make sure all possible paths of execution were accounted for.

In addition to uncovering bugs in our code, unit tests provided us with several other benefits. Firstly, we found that this strategy meshed well with our Agile approach to software development, as units of work loosely corresponding to functional requirements could be added to the live version of Thalia without the

risk of breaking it. This, in turn, meant that after the continuous integration process was set up, we could continually test to see that all of our project's modules worked properly with each other. Since each functional requirement had an associated suite of unit tests, we could be reasonably certain that it was implemented correctly and would not block our progress down the line. Moreover, team members could work on Thalia's components interchangeably, as per our management strategy, since any undesired side effects of one person's changes would cause the tests to fail. Finally, unit tests helped us to identify bugs resulting from merge conflicts and in doing so helped keep the live version of Thalia bug-free. This was especially important as we were aiming to merge as quickly and as often as possible, in line with our Agile approach.

7.1.2 Integration Testing

Equally important to the unit tests were our integration tests. As with unit testing, we used a mixture of white and black box testing techniques. The majority of integration testing was done upon first setting up the CI process, as this was when the individual components of Thalia first had to interact with one another. Subsequently, as we added new features, additional testing was performed to make sure new dependencies between components were integrated properly. Integration tests also helped us refactor existing code to work better with other modules, since designing integration tests before writing code helped highlight areas where components' interfaces differed.

7.1.3 End-to-end Testing

Finally, upon completion of the prototype, we conducted extensive end-to-end testing. We used these tests to help identify system dependencies and to check whether data integrity was maintained throughout execution. They were also key in conclusively evaluating the functionality of our system, as they tested it holistically with real-world data as input, which is as close to a real world environment as we could get before the prototype release.

7.1.4 Data Harvester Logs

Since our data is collected from 3rd party APIs, we opted to implement a series of logs to track their behaviour at runtime. These aim to record issues with received data that stop short of causing a catastrophic failure in the Data Harvester's execution. Such a situation might occur if an API were to change its policy or experience downtime, or if there is a problem with the network on Thalia's side. In this case, we are able to look through the logs and determine what caused the error and, consequently, what parts of the Data Harvester may need modification.

The logs record where and when data retrieval failed, what data was successfully retrieved and whether the data was successfully written to the database. In addition, they perform checks at runtime to test whether the API's data is of the expected type and format.

```

1 INFO:root:
2
3 Started update at: 2020-03-09 19:49:13.691662
4 INFO:root:Started updating yfinance at 2020-03-09 19:49:13.691850
5 INFO:root:Updating for yfinance doing 1 calls
6 INFO:root:calling for:
7 asset_class: index_funds ticker: VFIAX start_date : 1970-1-1 end_date: 2020-03-08
8 INFO:root:api returned <class 'pandas.core.frame.DataFrame'>
9 INFO:root:Received data frame with data between 2000-11-13 - 2020-03-08
10 INFO:root:Start interpolationdataframe shape: (4857, 7)
11 INFO:root:Data Frame Interpolateddataframe shape: (7054, 7)
12 INFO:root:Writing interpolated dataframe to DB
13 INFO:root:Finished updating yfinance at 2020-03-09 19:49:33.868327
14 INFO:root:Started updating nomics at 2020-03-09 19:49:33.868475
15 INFO:root:Updating for nomics doing 1 calls
16 INFO:root:calling for:
17 asset_class: currency ticker: ALL start_date : 1970-1-1 end_date: 2020-03-08
18 INFO:root:api returned <class 'pandas.core.frame.DataFrame'>
19 INFO:root:Received data frame with data between 2018-05-31 - 2020-03-08
20 INFO:root:Start interpolationdataframe shape: (648, 7)
21 INFO:root:Data Frame Interpolateddataframe shape: (648, 7)
22 INFO:root:Writing interpolated dataframe to DB
23 INFO:root:Finished updating nomics at 2020-03-09 19:49:36.457801
24 INFO:root:Finished all updates at: 2020-03-09 19:49:36.457934

```

Figure 28: Example of data logged by harvester

7.2 Testing Tools

The following is the list of 3rd party tools used by our testing suite to aid with the testing of Thalia and its components:

1. **pytest:** For our testing harness, we opted to use the standard python testing framework pytest. With it, we could easily create a testing suite independent of production code. By using simple assert statements, pytest remains easy to use, helping to reduce overhead for our team, most of whom had little to no experience writing automated tests for their code. Although simple, pytest offers several powerful features that helped us during testing:

- Detailed introspection of failing tests helped decrease time spent debugging
- Fixtures allowed us to automate test setup and teardown
- Test discovery helped to keep the structure of the testing suite simple and physically separated from production code
- Testing for expected exceptions allowed us to properly test for unexpected flows of events

```

/usr/lib/python3.8/site-packages/werkzeug/local.py:378: in <lambda>
    __getitem__ = lambda x, i: x._get_current_object()[i]
-----
self = <SecureCookieSession {'_fresh': True, '_id': '3da5fe10fd627b080e6c2d3525f339fe2ca24ae875ef9b44ed1e2c3aefd26cf5daf4
eb44a705e2fa7723ae6a7cec477948bad9cbd569357bfb219b430b3b733', 'user_id': '1'}>
key = '_user_id'

def __getitem__(self, key):
    self.accessed = True
>   return super(SecureCookieSession, self).__getitem__(key)
E   KeyError: '_user_id'

/usr/lib/python3.8/site-packages/flask/sessions.py:84: KeyError
===== 2 failed, 46 passed in 30.72s =====

```

Figure 29: Example of pytest assertion introspection

2. **Mock:** The Mock [91] package allowed us to replace parts of the system with mock objects. This was useful when testing API's as it allowed us to control the input to the Data Harvester and allowed for repeatable execution and predictable output.
3. **Selenium:** Selenium [92] is a testing tool we used to create unit and integration tests for Thalia's website. Using Firefox or Chromium in headless or full browser mode, it allowed us to automate user tasks such as clicking links and text input. With it, we were able to test logging in, running a simulation, registering a new account, and accessing the various pages of our website.
4. **Coverage:** Coverage [93] is a tool to measure code coverage in Python. It creates detailed reports on what lines of code are passed through during program execution. We used this in conjunction with pytest to monitor what code was covered by tests to ensure our testing suite covered all possible branches of execution. In addition to validating test coverage, Coverage aided us in designing white-box tests by showing what parts of the codebase were in need of additional testing. Although not related to testing, the tool proved useful when refactoring code, as it allowed us to identify dead code.

7.3 Test Data

We selected test data for our testing suite with the following considerations in mind:

- For Black-box tests, the test data should contain representatives of all major equivalence classes of possible inputs.
- For Black-Box tests, the test data should encompass all major boundaries of equivalence classes in the accepted inputs.
- For White-Box testing, the data should be selected so tested code is executed exhaustively.

Broadly speaking, the data we used for testing can be separated into two categories, based on how it was collected. Either we wrote code to procedurally generate the data, or we used a subsection of the real world data included in the final prototype and collected from financial data APIs. Each of these had its own benefits and drawbacks when being used to design tests.

7.3.1 Mock Data

In the real world, financial data, i.e. the prices of assets, tend to experience frequent and unpredictable fluctuations. Part of our testing strategy was the generation of mock sets of historical price data that we designed to be as clean and easy to understand as possible. This allowed us to work with neat, easily understandable datasets. For example, we created a fictitious asset whose price increased linearly over time, an asset whose price decreased monotonically and an asset whose price remained fixed. All mock data was either generated procedurally by auxiliary code we wrote ourselves or hand-coded in cases where large amounts of data were not necessary (for example when testing the functionality of the database adapter). Being able to design these custom data sets helped us to overcome the following difficulties when designing and writing tests:

- Predictability: Since complex financial equations are a key part of our service, a core requirement for our testing strategy was to be able to independently verify the results generated by the Anda library. Having simple data showing, for example, a linear or quadratic increase in the daily price of an asset meant we could predict the expected results and compare them to Anda's output. With real-world data, calculating expected values by hand would have been effectively impossible.
- Interpretation of results: Components of both Thalia Web and Finda modify financial data as it passes through them. It proved difficult to see the effect processing had on real-world data. Using clean datasets meant that the effect methods had on data could be examined by a human. This was a significant aid to development and saved us considerable time in the long run.

- Designing tests: When designing black-box tests for Thalia’s components, it was helpful to be able to create datasets that had specific properties for use as edge cases. A good example of this would be creating a series of prices that never decreased, as this represents an edge case for the calculation of an assets maximum draw-down.

7.3.2 Real-world testing data

In addition to generating our own assets for testing, we used a limited subset of real-world data for testing across components. We chose to include assets from across all supported asset classes over a significant period of time (several years as the simulation will at most span several decades). We found that this approach was easier than trying to emulate the complex fluctuations of real-world asset prices ourselves. While predicting the expected output of methods when handling real-world data was more difficult, we still found it proved useful in the following circumstances:

- Validation: Correct behaviour of system can be better confirmed after running them on real data [94]. In our case, this is especially true since the data we work with is so complex.
- Performance: Performance of some of Thalia’s features might vary based on the complexity of the input. As a result of this, real data was needed to properly assess performance.
- Design: An important design consideration was how clear the UI elements, especially the ones displaying data visually (for example the price graph) looked on screens of varying form factors. Having realistic test data helped us to better assess the quality of the user experience.

7.4 Testing Results

With all unit, integration and end-to-end tests passing, it is safe to assume that the core functionality of Thalia works as expected (short of a manageable risk of further bugs appearing, as no testing suite of this scale is perfect). Our unit tests are exhaustive, testing for the majority of possible edge cases and achieving a high level of coverage for each module tested. The functionality of the final prototype was additionally tested by hand and found to be working for the included data (a range of assets from all major asset classes with over 10 years of historical data). As such, it is safe to assume adding additional structurally similar data will not break the product. At the time of writing, we deem the testing of Thalia to be successful and to have demonstrated the suitability of our product for further development and, eventually, bringing it to market. Going forward, we would like to implement a test driven development approach with the hopes of further reducing software defects and improving quality of design[95].

8 Evaluation

8.1 Approach

Evaluation of our product began after the team finished development and testing. As mentioned previously, code was only merged after successfully passing automated tests and review by at least one other team member. We conclude that implemented features work correctly and therefore implement their corresponding functional requirements correctly. Additional evaluation was required to assess whether or not identified non-functional requirements were fulfilled. This evaluation was performed on a requirement-by-requirement basis, and in some cases required the use of additional external tools (for example the installation of assistive technologies when evaluating accessibility requirements). In some cases, proper evaluation would have required access to resources not available to our team due to budget constraints. In these cases, the team attempted to perform evaluation as best possible and outline the difficulties experienced in this report. The team’s main focus during evaluation was objectivity. To achieve this goal, we relied on external tools and metrics for the evaluation of requirements where possible.

Although our evaluation process is extensive, we recognize that no amount of heuristic analysis is a perfect replacement for proper user testing, as these are two different approaches that address separate aspects of software usability and adequacy [96]. Additionally, some of the requirements outlined in the first

semester, although important, are based on a user's subjective experience and are therefore impossible to evaluate without extensive user testing. Upon consulting with our guide and course coordinator, we were advised that user testing would be outside the scope of this course, and as such, the team decided not to conduct user testing at this stage. Extensive user testing will be performed during Thalia's Alpha and Beta releases to fully evaluate our product's usability and user experience. While conducting user testing, there is a possibility the team identifies new user personas, and consequently new requirements. As we are following an Agile approach to software development, these features can easily be added to subsequent releases of Thalia by implementing a Release on Demand process [97]. In the following sections, we will first discuss the results of the evaluation conducted for functional and non-functional requirements, followed by a discussion of the overall success of the project.

8.2 Functional Requirements

As mentioned in last year's Technical Report, we introduced some milestones or releases for our product based on incremental sets of features. This was done to help us stay on track with development. The roadmap based on this schedule can be seen on Figure30.

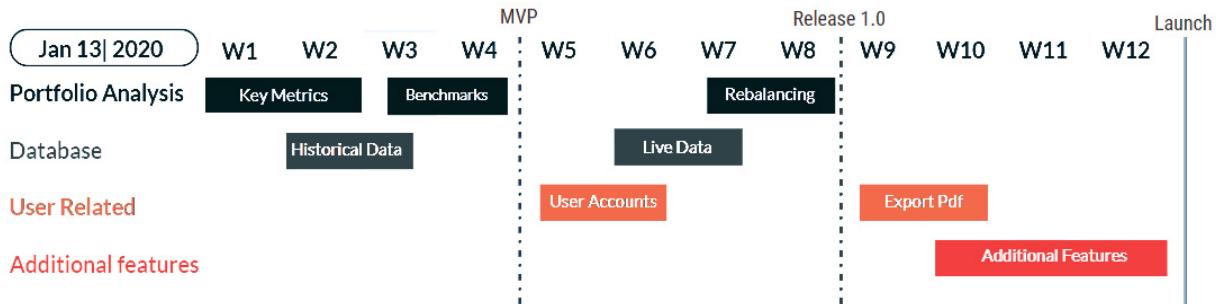


Figure 30: Technical Report: Roadmap - revisited

Due to our team's focus on correctness and code quality, the implementation of many features ended up taking longer than we initially predicted. Although the implementation of a single feature would generally be assigned as a single ticket, the writing of tests and refactoring of code to adjust for feedback would often mean tickets would take more than one week (our team's sprint length) to complete. Additionally, our team spent a lot of time implementing 'under the hood' features to improve extensibility and maintainability. Most notably, the Anda module was designed as a separate library and Finda was designed to include advanced functionality seen in other ORMs (permission management and support for managing of several databases). Although these features contribute to the quality of our product, they ended up taking developer time away from the development of functional requirements in the early stages of our project. Throughout our development process, the team would hold weekly retrospective meetings to try and better estimate velocity. These allowed us to react to this problem halfway through our allotted time. We came to the consensus decision to lock features that would be implemented and increase target weekly hours team members were expected to put into the project from 10 to 20. Additionally, we loosened some of the constraints necessary for deliverables to pass the review process and specialised the roles of team members to reduce the overhead of knowledge transfer. The result of this was that at the end of the project, all functional requirements except one were properly implemented.

The functional requirements not implemented at the time of writing are the ability for a user to search for assets by category when creating a portfolio and front end support for currency conversion. Their omission is a result of a lack of resources, architectural limitations and our internal assessment of their importance. Halfway through development, the team assessed each functional requirement not yet implemented and determined these two to be of low importance. This decision was based on the fact that searching assets by category only provides users with a slight improvement to quality of life without giving any additional insight into the viability of their portfolios. Additionally, our team determined that implementation of this feature would require architectural changes to the Thalia Web and Finda modules, meaning that its implementation

would represent a bad cost to value ratio when compared with other remaining features. Although back-end support for different currencies is implemented in the Anda module, front-end support would require sourcing and collecting a large amount of new data. The team decided that front-end support for this feature is not critical for a prototype and would prefer implementing it later on when adding support for languages and other elements of localization. Even without these features, we believe this prototype to be fully featured and usable and therefore ready for the next stages of development.

8.3 Non-Functional Requirements

The following are the results of our evaluation of non-functional requirements identified by our team - listed by category - along with some general notes on our approach to evaluating each.

8.3.1 Usability

Usability requirements were particularly hard to assess. This was due to their subjective nature, relying on the experiences of users. Hence, we will only conclusively be able to say that these requirements are fulfilled after performing user testing.

- The product must be easily usable for users who already have some financial investment experience.
Since Thalia's design and features are based on feedback from finance students and enthusiasts interviewed in the first semester, we can infer that someone of similar experience would likely be familiar with most of Thalia's terminology, features and charts.
- The basic backtesting interface needs to look familiar to people already experienced with it.
The design of our interface closely follows that of our major competitor (Portfolio Visualizer), so it is safe to say that a user with moderate experience using it or other backtesting tools would have no trouble finding their way around.
- The product must have detailed instructions on how to use its advertised functions.
A link to an 'About' section featuring basic instructions on how to use Thalia is featured in the navigation bar on Thalia's landing page. Additional material is also available in the user manual 11.
- All major functions must be visible from the initial landing page.
All pages are accessible through the navigation bar at the top of Thalia's landing page.
- Must work in both desktop and mobile browsers.
Since Thalia's dashboard is built using the Dash platform, it automatically scales to the size of the browser window, even when accessed from a mobile browser. All other pages have also been designed to work on mobile platforms.
- The results page should scale with mobile
Same as for the previous requirement.

8.3.2 Reliability

- The product must have a greater than 99% uptime.
Thalia is hosted by Amazon AWS, a reputable hosting provider who's service level agreement (a legally binding contract) promises a monthly uptime of over 99.9%[98].
- All our assets need to have up-to-date daily data where the asset is still publicly tradeable.
All assets included in Thalia are gathered from APIs by the Data Harvester. All API usage constraints were thoroughly researched and tested to make sure we are able to access data daily and that the data provided was updated on all trading days. This practice will continue for any API that we wish to add to Thalia.
- All assets supported by the system must provide all publicly available historical data.
Both sources of data currently used by the Data Harvester (Yahoo Finance and Nomics) are publicly available and do not impose any restrictions on the use of said data.

8.3.3 Performance

Performance and load times of Thalia's web page hosted on our cloud hosting provider were tested with the use of Google Lighthouse [99], a tool for evaluating web page quality.

- The website should load within 3 seconds on mobile.

Although Google Lighthouse estimated Thalia's 'time to interactive' (time needed for all interactive elements of the website to become accessible to the user) to be above the 3-second threshold, all other features of the page were loaded in this time. Additionally, Thalia scored well overall in performance.

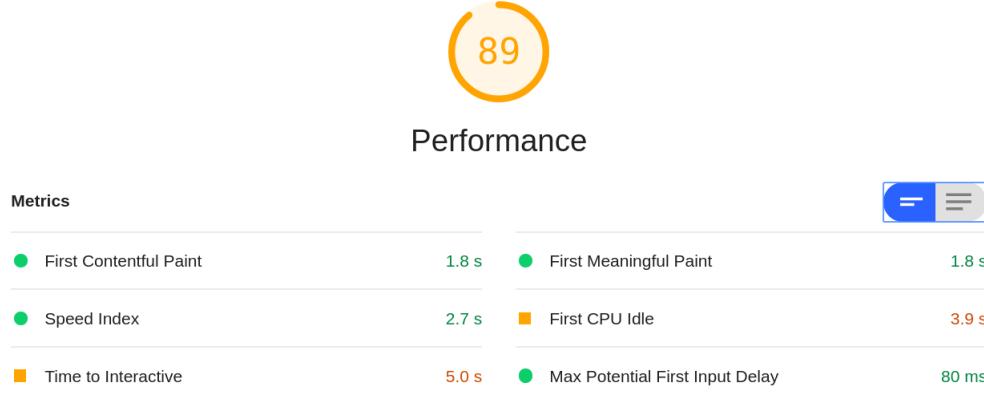


Figure 31: Results of Lighthouse Performance Tests

- Large portfolios must be supported - up to 300 different assets.

There is no reasonable limit on the size of the portfolio supported, and our tool has been successfully tested on portfolios of over 300 assets. That said, the tool does become unwieldy when entering portfolios of that scale since the table assets are entered into becomes too large to fit on a screen without scrolling. This problem is especially pronounced on mobile.

8.3.4 Supportability

- The system should be well documented and easy to maintain. *The architecture and design of the system are documented in this report. Additional documentation on the individual modules is also available on the project wiki. Additionally, the team has created a maintenance manual 12 explaining in detail how to maintain and extend Thalia.*
- The system should be fault-tolerant and be designed to support graceful degradation. *The system has been designed so that faults resulting from bad input and API problems do not crash the website. This means that users not directly affected can continue to run simulations.*
- Installation and migration between hosting providers should be simple. *In an appropriate environment, both Thalia and its dependencies can be installed with a single command. Installation is extensively documented in the user manual 11.*
- The system should be internationalized and be easy to extend to support new languages and currencies. *The current version of the Anda library supports currency conversion, hence adding new currencies is supported by our architecture. Versions of Thalia's website and dashboard in different languages can easily be implemented by hiring a translator. This is, however, outside the scope of our project which is merely a high quality industrial prototype.*

8.3.5 Implementation

- The system needs to work on a cloud hosting provider.
The system is successfully deployed on Amazon AWS cloud hosting services.

8.3.6 Interfacing

- The Data Gathering Module must never use APIs stated to-be-deprecated within a month.
Neither Yahoo Finance nor Nomics has made any statement to this effect. As such, we can assume both APIs will remain unchanged and available in the near future. In addition to this, Nomics API's SLA (service level agreement) guarantees a high average uptime [100].
- The Data Gathering Module must not exceed its contractual usage limits.
Fulfilled as previously stated.

8.3.7 Operations

- An administrator on-call will be necessary for unexpected issues.
This requirement would require the hiring of additional staff and is therefore outside the scope of this project. We have, however, included a form for submitting complaints on the Thalia website.

8.4 Packaging

- The product needs to work inside a Linux container (e.g. Docker).
As previously mentioned, our CircleCI process' tests are all run in a Docker container. This means that throughout development all working versions of Thalia have been tested and are therefore proved to be working inside one.
- All dependencies need to be installable with a single command.
Dependencies necessary for deployment are installed automatically when Thalia is installed with pip. Testing and development dependencies can be installed additionally, but are not necessary for running Thalia.

8.4.1 Legal

Proper evaluation of the fulfilment of legal requirements would require consulting a legal professional, as none of the team have an in-depth understanding of British law. This is, again, outside the scope of this project. Nevertheless, the team has attempted to comply with the identified legal requirements to the best of our abilities so as to try and minimise potential future changes. The steps we have taken to fulfil these requirements are outlined below.

- All user testing must be done with ethical approval from the University.
We were advised by the course coordinator that user testing is not necessary for this stage of our project. Additionally, we were advised that as long as no personal information is published, user testing could proceed without such approval. Therefore, this is a false requirement.
- UI must display a clear legal disclaimer about the service not providing financial advice.
A disclaimer is clearly displayed on Thalia's website. Since our tool doesn't propose a specific course of action to our users, it is likely it would constitute at most financial guidance and not financial advice [101]. In addition, our disclaimer is based on disclaimers found in the terms of service of competing software [102] and should, therefore, require little modification after review by a legal professional.
- All third-party code should allow for commercial use without requiring source disclosure (e.g. no GPL-3).
With the use of the pip package manager, we were able to confirm that the majority of third-party code does not require us to disclose our source code. The remaining (unknown) package was identified to be nomics-python, which is used to access the Nomics API and is distributed under the MIT license, meaning it also allows commercial use without source code disclosure.

```

python -m piplicenses --summary
Count License
1 Apache
2 Apache 2.0
6 BSD
1 BSD License
6 BSD-3-Clause
1 BSD-like
1 Dual License
1 LGPL
17 MIT
1 MPL-2.0
1 UNKNOWN

```

Figure 32: Licensing of 3rd party packages

- User data handling should comply with GDPR.
The team has done its best to familiarize itself and comply with GDPR and related legislation[103].
- Provided services should not constitute financial advice under UK law to avoid being subject to financial advice legislation and potential liability.
Thalia is merely a tool for running simulations, and as such we at no point advise users on how to invest their real-life assets. This is also stated explicitly in our legal disclaimer.

8.4.2 Accessibility

- Display items should be clearly labelled.
All UI elements are clearly labelled based on their purpose. All elements of the dashboard have been given descriptive names.
- UI should scale to accommodate different screen sizes and aspect ratios.
All pages on Thalia's website have been designed to scale to different-sized desktop and mobile displays. Dash apps automatically scale and rearrange elements to fit the display, meaning the main dashboard scales as well.
- UI elements and text superimposed over one another should have high contrast in their colours.
When selecting Thalia's colour scheme, care was taken to allow for such contrast. On the dashboard, foreground elements and text are rendered in dark blue and black, while background elements are white or bone. A similar colour scheme is used for the other pages on Thalia's website, with foreground and background colours inverted.
- UI should allow for the use of assistive technologies to accommodate for individuals with accessibility issues.
Selection of assistive technologies to test our website was made based on guidelines set out by the UK government identifying the 5 most commonly used assistive technologies [104]. As these work by magnifying parts of the screen and reading out website text, they did not encounter any issues on our website. This means that Thalia in its current form reaches the high standard set for UK government institutions and associated businesses.

8.5 Evaluation Results

The prototype of Thalia we have delivered along with this report fulfills the vast majority of identified functional requirements. We have also provided adequate reasoning for why we decided not to implement the remaining functional requirements. Where possible, evaluation was also performed for each non-functional requirement. We believe that we have shown why these, too, are adequately fulfilled. In addition, our team has focused on the quality of our process throughout Thalia’s development and, as a result of this, has been able to deliver high-quality code that has been extensively reviewed and tested.

9 Conclusion and Further Work

9.1 Conclusion

In this report, we have presented Thalia, a portfolio backtesting web application. After engaging in discussion of the product requirements and our approach to development, we have provided justifications for any design and implementation decisions as well as technology choices. In the following sections, we then went on to specifying our testing strategy and the results of our evaluation process.

As a consequence of the preceding evaluation section, we perceive our project to have been successful. We have delivered a high quality, functioning piece of software that fits our identified business needs and market niche. Since we have already provided extensive discussion of whether our requirements have been met, we may now provide an outline of further work that may be done in the future.

9.2 Further Work

9.2.1 Payment Service Integration

Charging for the service we provide requires processing of customer payments. Given the sensitive nature of payment data (e.g. credit card numbers), our preferred approach is to outsource processing of payments to a third-party service, such as Stripe [105]. Our responsibility would then be reduced to integrating the vendor payment system into our codebase.

9.2.2 Implementation of a Scripting Language

One of our optional features from last semester’s technical report was the development of a scripting language that would allow investors to specify dynamic trading strategies. As an example, an investor may want to sell their equities for bonds whenever there is a drawdown greater than 10% in equity price indices. While support for such a scripting language was outside of the scope of this project, it is technically feasible, as has been demonstrated by other backtesting tools, e.g. Stockbacktest [16].

9.2.3 Additional Risk Metrics

Our aim for this project was to provide values for the most commonly encountered risk metrics when backtesting a portfolio. However, many other, more exotic risk metrics exist. Over time, we may want to incorporate additional metrics into our dashboard to enable an even more nuanced analysis of an investment strategy. To prevent more novice users from being overwhelmed, these could be deactivated by default and enabled on demand.

9.2.4 Expanding Asset Selection

Our team is proud of the breadth of assets available for analysis in our tool. As discussed previously, the inclusion of European market assets sets us apart from many of the commercially available tools. Nevertheless, we were unable to support every asset class and even every asset within the asset classes we support. In particular, we would be interested in supporting standard exchange traded derivatives, such as Futures and Options. In order to appeal to other investor demographics, it may also be worthwhile integrating price data for African, Asian and Latin American assets.

9.2.5 Equity Dividends

Investors receiving dividends for the equities they own face a choice: Should dividends be reinvested in the equities themselves or spread over the whole portfolio? Supporting calculations that involve dividends would enable us to provide empiric answers to these questions. Our business logic library already provides support for dividend calculations, hence this would be a simple matter of making appropriate changes to the UI.

10 Appendices

11 Appendix A - User Manual

11.1 Launching Thalia

As Thalia is a web application, it is accessible in a web browser via the url: <http://thaliabacktest.xyz/>.

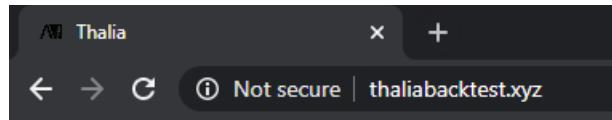


Figure 33: Thalia Web

11.2 Browsing Thalia

Navigation within the application can be achieved mainly through use of the navigation bar, visible on Figure34.



Figure 34: Thalia Navigation Bar

In case Thalia is launched on a different device, such as mobile or in a smaller window, the layout changes to fit the screen. In this case, the navigation bar becomes a so-called "hamburger button" and dropdown menu visible on Figure35.

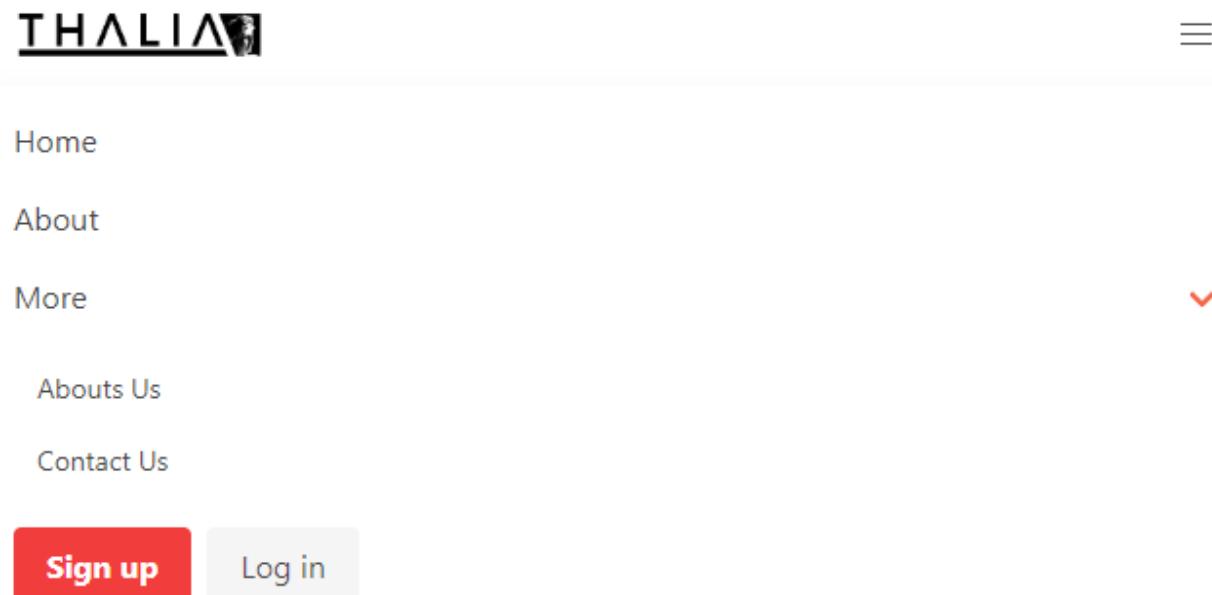


Figure 35: Thalia Navigation Bar - Hamburger

For the rest of this manual, we shall assume that the application was launched on a desktop, although the layout is identical and intuitive in both the desktop and mobile case.

11.2.1 Homepage

By default, the user ends up on the homepage, although some other pages are accessible as well, given the correct URL. The Homepage is visible on

Figure36.

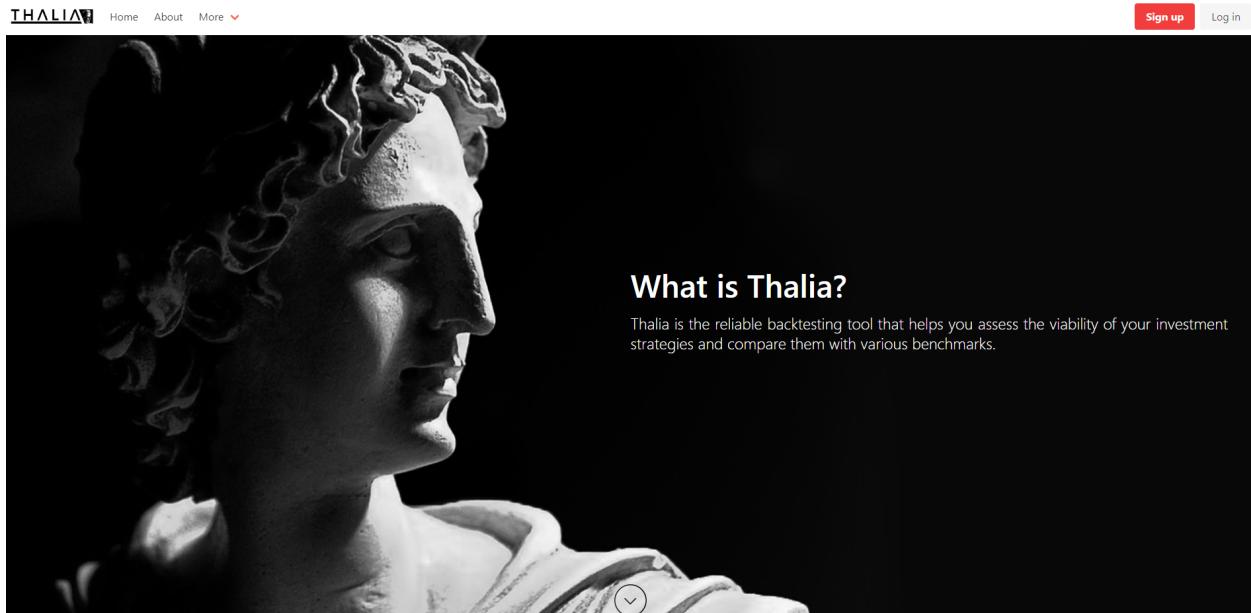


Figure 36: Thalia Homepage (source: <http://thaliabacktest.xyz/>)

The purpose of the homepage is to give a cover to our application. When users arrive at the homepage, they are greeted with a short description of what our software entails.

As we do not wish to have users jumping blindly into the process of portfolio analysis, we also provide some basic information on the backtesting process. A link at the end of the description then leads to the 'About' page, where users can read more on the topic.

Scrolling down further, the user may encounter a small register form, or in case the user is logged in, our tweets and social media links.

Still with us?

Go ahead and create an account, or log in if you already have one. Once registered, you can backtest to your heart's content.

Register

[Register](#)
Already have an account? [Login](#)

Already signed in?

Why not jump right in and create your dream portfolio? If you would like to stick around, make sure to follow us on our social media pages to see our latest.

Tweets by @Thalia99627941

Thalia
@Thalia99627941

Help pass the time during lockdown by perfecting your investment strategy with the **#Thalia** backtesting tool!

April 5, 2020

[Load more Tweets](#)

Figure 37: Top: User not logged in; Bottom: User recognised

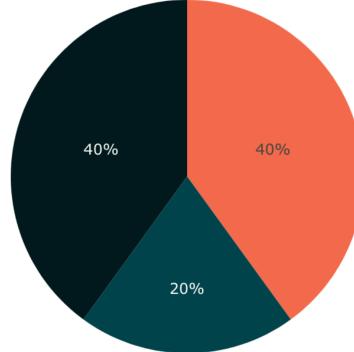
11.2.2 About Page

The about page of our application provides a more detailed description of the problem domain and is available at <http://thaliabacktest.xyz/about/>.

The user can arrive at this page either by directly typing in the URL, clicking on the learn more option on the homepage, or by navigating here using the navigation bar.

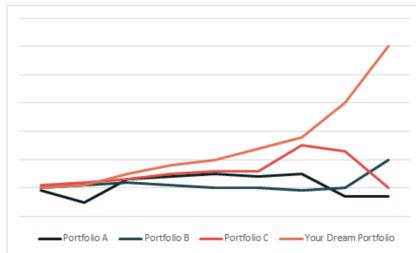
Our About Page looks as follows:

- Index Funds
- Stocks
- Currency



But how does it work?

Imagine you're nearing retirement and can't afford to sit patiently through another financial crisis waiting for your investments to regain the value that they have lost. You're also passionate about finance, and devised an investment strategy, which couldn't possibly fail, you say. Imagine this dream portfolio consists of the following:



That's great, but wouldn't it be useful to consider similar portfolios, comparing how much they've lost in the worst year out of the last 50? Just to make sure that the dream portfolio indeed works. This process is known as backtesting and is usually done with the assistance of a computer.

That is where Thalia can help you, with Thalia, you can:

- Specify an asset allocation using hundreds of assets across major asset classes
- Plot the performance of strategies against each other and an index benchmark
- Gain insight into a strategy's key risk and performance metrics
- Model regular contributions and re-balancing strategies
- Save and share a portfolio to compare and discuss with peers



Figure 38: Thalia About Page (source: <http://thaliabacktest.xyz/about/>)

11.2.3 Log In and Sign Up Pages

In case the User is not yet logged in, links for the 'Log In' and 'Sign Up' pages are visible on the navigation bar as seen on Figure34 or Figure35. In addition, they are directly available at <http://thaliabacktest.xyz/login/> and <http://thaliabacktest.xyz/register/>.

Both of these forms are quite common, with the login requiring:

- Username
- Password
- Remember me (optional)

And for signing up, the fields are:

- Username
- Password (minimum eight characters, at least one letter, one number and one special character)
- Confirm Password

By convention, the registration fails when the user enters different values to the Password and Confirm Password fields. In this case, the user is prompted to try again.

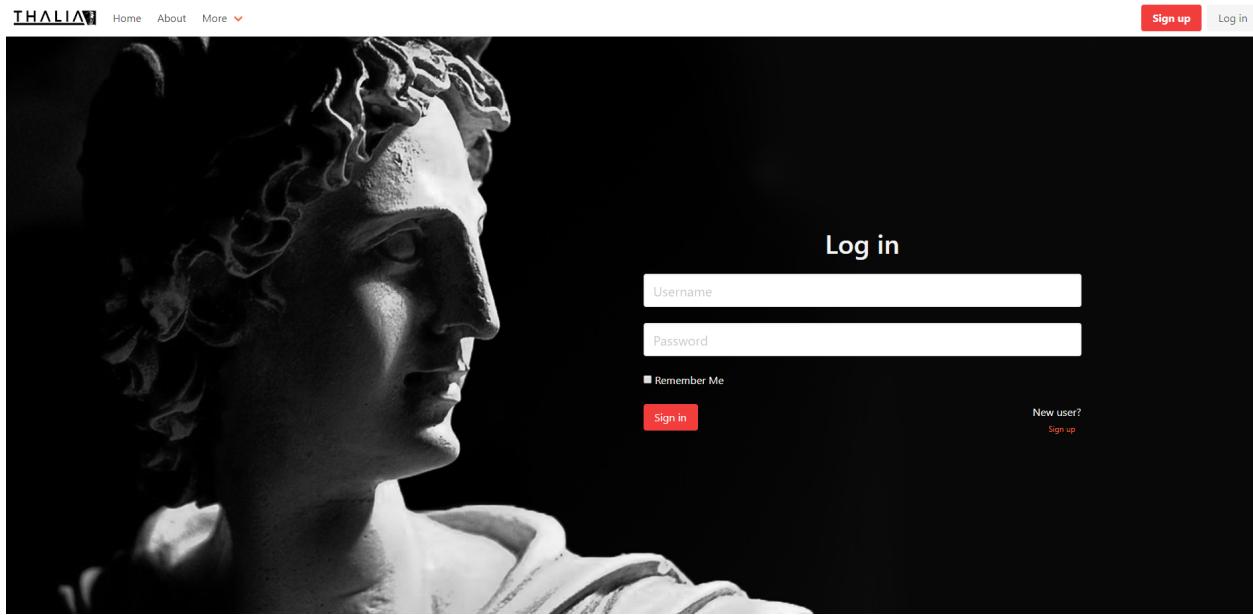


Figure 39: Thalia Log In Page (source: <http://thaliabacktest.xyz/login/>)

In case the user is already logged in and attempts to access these pages, they will be redirected to the homepage. In addition, the navigation bar changes, allowing one to log out as visible on Figure40.

Opting to log out, the user finds themselves at the homepage.

11.2.4 Contact Us Page

The 'Contact Us' page is available at <http://thaliabacktest.xyz/contact/> or via opening the dropdown menu on the navigation bar and clicking the link. This short form is for users to provide feedback, report issues or request features.

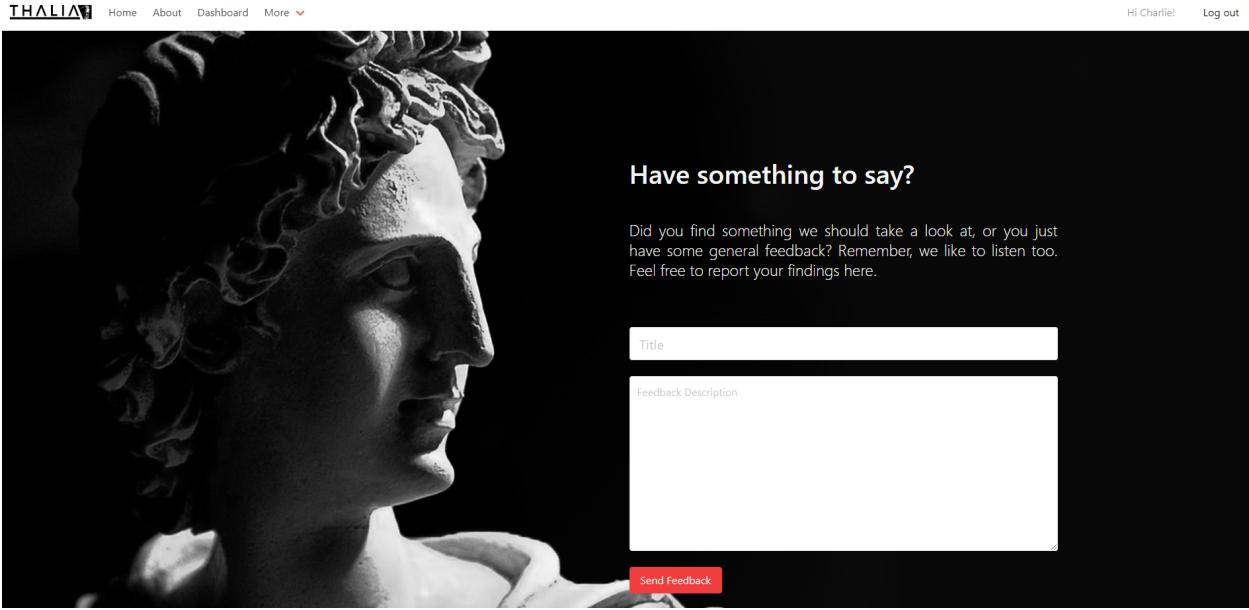


Figure 40: Thalia Report Issues Page (source: <http://thaliabacktest.xyz/contact/>)

11.2.5 About Us Page

We also wanted to take pride in our project and sign our work, following the advice in [49]. This simple page is dedicated to the team behind Thalia.

11.2.6 My Portfolios Page

The ‘My Portfolios’ Page is only accessible for authenticated users. Here, the user may view their saved portfolios with their asset allocations. If you are running Thalia for the first time, this page will be empty, encouraging you to go and create a portfolio. To create a portfolio, simply head to the dashboard and allocate your assets following the instructions in 11.3.2. Before submitting your investment strategy, make sure to save it for later. You can achieve this by clicking on the default portfolio name and changing it to something more fitting. Then, use the ‘Save Portfolio’ button. When you next visit the ‘My Portfolios’ Page, you can see the portfolio and its associated asset weights. Clicking on the box will take you back to the dashboard. The corresponding allocations will already be loaded to the table. Alternatively, you can select any of your previously saved portfolios directly from the dashboard by using the dropdown menu and selecting your chosen portfolio.

The following image shows the aforementioned page with some example portfolios saved.

Dream Portfolio A		Dream Portfolio B		Dream Portfolio C	
JEP	20.00%	JPY	10.00%	ALL	50.00%
KRW	40.00%	SB	70.00%	NGN	50.00%
ALL	20.00%	GTQ	10.00%		
JPY	20.00%	O	10.00%		

Figure 41: My Portfolios - Example (source: <http://thaliabacktest.xyz/gallery/>)

11.3 Running a Backtest

To run a backtest, the user must be authenticated. After a successful registration and/or login process, the user can head to the Dashboard.

11.3.1 Dashboard

The Dashboard, the heart of our application, is available for registered users at <http://thaliabacktest.xyz/dashboard/> or via the link on the navigation bar.

This web page is divided into tabs, which are:

- Ticker Selector
- Summary
- Metrics
- Returns
- Drawdowns
- Assets
- Overfitting

At first, only the Ticker Selector Page is accessible to the user.

This is because the following tabs are used for displaying the results of backtesting one or multiple portfolios and, as a consequence, are initially empty.



Figure 42: Thalia Dashboard - Disabled Tabs

Let us now consider each tab individually.

11.3.2 Ticker Selector

Here, the user is required to input their backtesting strategy. As Thalia supports testing multiple portfolios at once (currently up to 5), there are input fields which are portfolio specific and others that are not. The latter consists only of:

- Start Date: The day the investment is made.
- End Date: The last day of the investment, which defaults to the present date.
- Initial Amount: Initial balance. The default currency is dollars, as is convention for financial applications.

The first two of these are standard date-selectors, but also allow for typing the date directly. Next, the user enters the data specific to the current strategy.

- Portfolio Name: Defaults to Portfolio 1, Portfolio 2, etc.
- Contribution Amount: The amount that should be added to the total investments on a regular basis (e.g. after receiving a monthly paycheck).
- Contribution Frequency: Specify how regular these contributions should occur.

- Rebalancing Frequency: Specify how often rebalancing (reestablishing the initial weights of the portfolio by buying and selling assets) should happen.
- Asset Allocations: Discussed below.

The last one is the most crucial information in a strategy. First, the user selects the desired asset from the dropdown menu. The menu item also allows for typing in order to find an asset faster. The selected item is then added into the table below.

Ticker:

A screenshot of the Thalia Dashboard. At the top, there is a search bar labeled "Ticker:" containing the text "JPY". Below the search bar is a table titled "Asset Ticker" with columns "Asset Ticker", "Name", and "Allocation". The table contains three rows: "SPY" with allocation 1, "VTSAX" with allocation 2, and "JPY" with allocation 1. The "Allocation" column for "JPY" is highlighted with a pink background.

	Asset Ticker	Name	Allocation
x	SPY		1
x	VTSAX		2
x	JPY		1

Figure 43: Thalia Dashboard - Asset Table (source: <http://thaliabacktest.xyz/dashboard/>)

The user then specifies the weight of the asset as an element of his portfolio. This is required to be a numeric value, representing the *relative* weight of the asset. For example, given asset A with relative weight 1, and asset B with relative weight 2, then 33% of the investment is allocated to asset A and 66% to asset B.

In case the user is content with the portfolio, they can either add another strategy via the “Add Portfolio“ button or click “Submit“ to see the results. In case the user has already entered the maximum number of portfolios, i.e. 5, the button is disabled and the only possibility is to submit.

If the user has not entered some required data, such as an initial amount, or selected at least one asset, a warning message is shown, telling them to complete the form correctly. Note that this also happens when the user selects a time frame shorter than a year, as many of our metrics do not make sense for investments for such short periods of time.

In many cases, the user may want to compare their portfolio to a benchmark, i.e. to a small set of standard portfolios. These can be selected from the “Lazy“ dropdown menu at each portfolio, which then populates the table with the desired proportions.

Alternatively, the user may want to test an allocation on price data only they possess. If the user has the data in the format of a CSV file, this can be done by simple dragging and dropping the file to the corresponding field.

Upload your own ticker: [?](#)

A screenshot of a user uploaded data field. It features a dashed rectangular border and a central text area with the placeholder "Drag and Drop or Select Files".

Figure 44: User Uploaded Data Field

11.4 Viewing the Results

The result of the backtesting process is divided between the remaining tabs of the dashboard. Let us now discuss what each of them tell the user.

11.4.1 Summary

If all required fields are populated, the user is taken to the ‘Summary’ tab. This, as well as all other tabs, are now unlocked. At this point, the user is shown one of the key components of our application, i.e. the portfolio growth graph. All graphs are fully interactive. The user may zoom in on selected areas, hover over desired data points, save the plot as an image, etc.



Figure 45: Dash Functionalities

The plot seen at the top of Figure 46 shows the total return of each portfolio. Users will typically use this graph for comparison. In addition, for each portfolio the user has specified, the following is shown:

- Selected Key Metrics: Final Balance, Best Year, Worst Year, etc.
- Pie Chart: Proportion of each asset.
- Bar Chart: Shows how the total return changed over the years relative to the last.

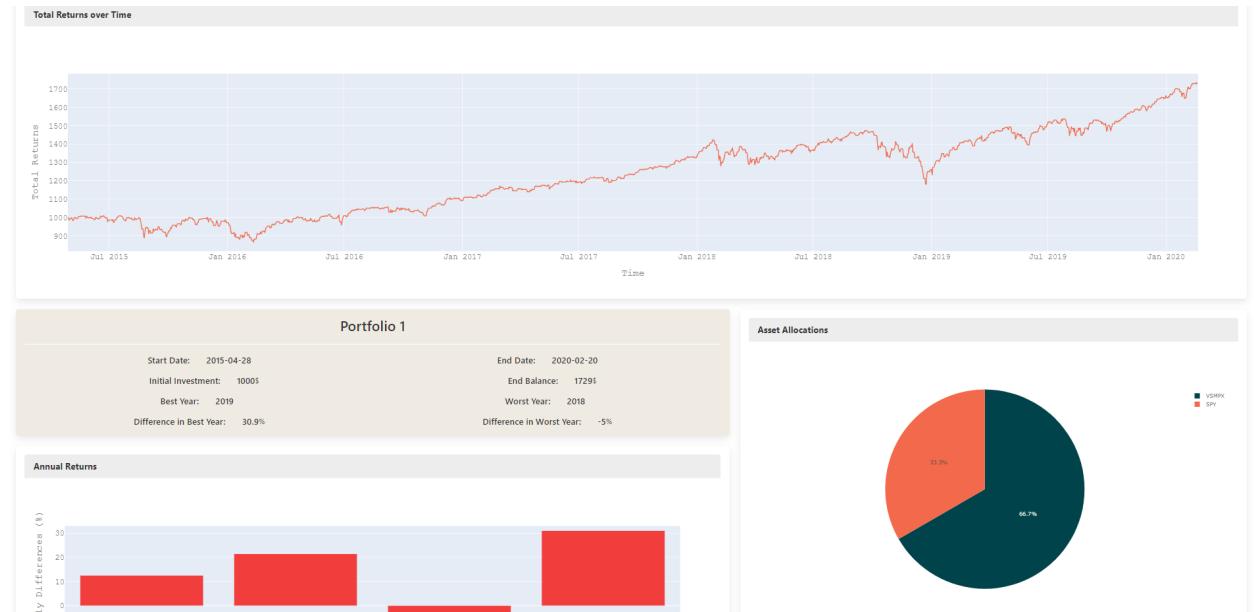


Figure 46: Summary Tab (source: <http://thaliabacktest.xyz/dashboard/>)

It is also important to note that in some cases, there is no data for a given asset. For example, there is no price data for cryptocurrencies before the 2000s as they did not yet exist. In the edge cases, where the portfolio is made up entirely of these assets, the strategy may not be valid over a given investment period.

If there is only one portfolio given as input and affected by the above, the simulation will be run for a time frame over which data for all assets exists. In the case where the user provides more than one investment strategies with different time frames, we provide output where price data is available.

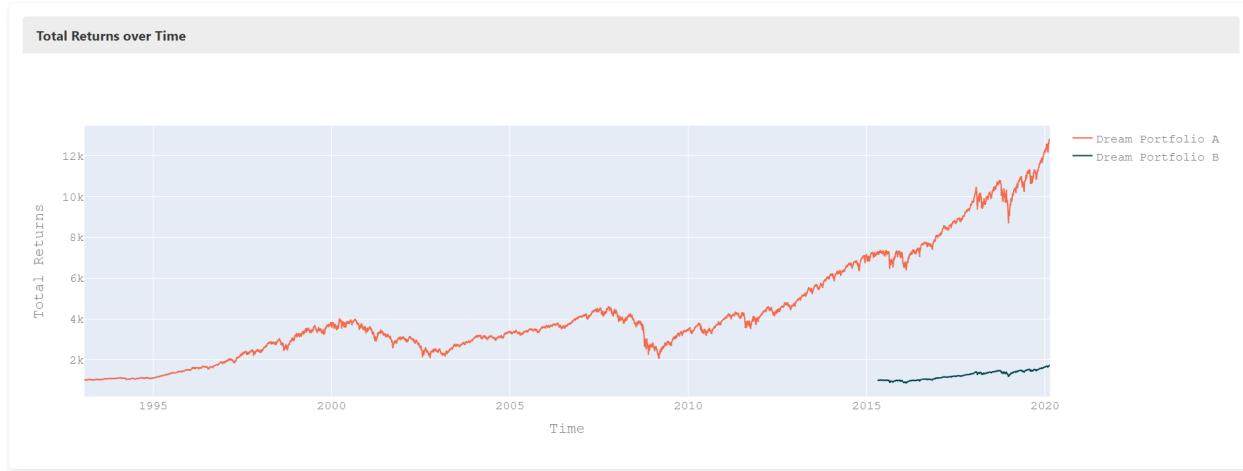


Figure 47: Different Time Frames per Portfolio (source: <http://thaliabacktest.xyz/dashboard/>)

11.4.2 Metrics

The ‘Metrics’ tab shows a simple table of key metrics for each portfolio. This allows for a quick comparison of strategies. Although this page does not show much more data than the previous one, it could easily be expanded later. In case the table gets bigger than the browser window, a horizontal scrolling can be used for viewing the rest.

Key Metrics

Metric	Dream Portfolio A	Dream Portfolio B
Initial Balance	1000	1000
End Balance	12737.24	1718.19
Best Year	38.05	30.79
Worst Year	-36.8	-5.16
Max Drawdown	55.19	20.16
Sortino Ratio	0.94	1.34
Sharpe Ratio	0.61	0.9

Figure 48: Metrics Tab (source: <http://thaliabacktest.xyz/dashboard/>)

11.4.3 Returns

The ‘Returns’ tab allows the user to compare the annual returns of each portfolio. It consists of a bar chart similar to the one in the Summary tab and a table similar to the one in the Metrics tab. The graph can be seen on Figure49.

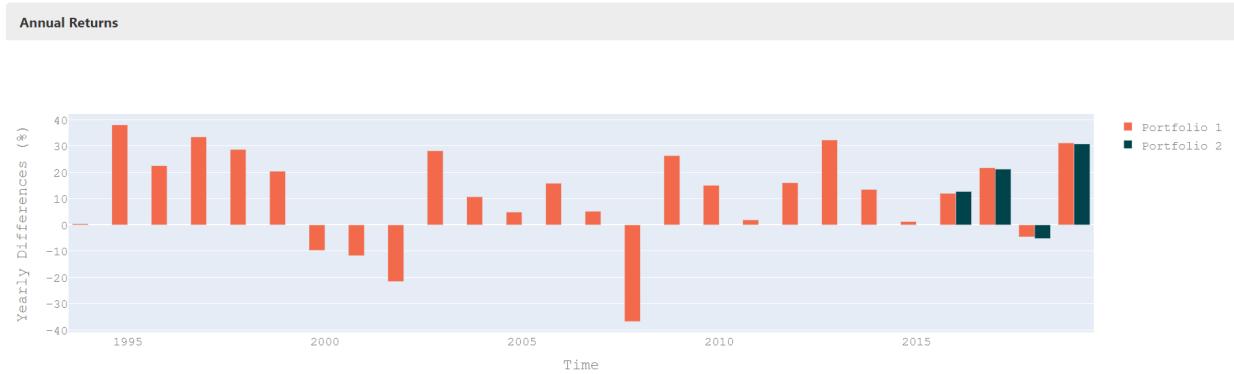


Figure 49: Metrics Tab (source: <http://thaliabacktest.xyz/dashboard/>)

It may happen that one of the portfolios was not available at certain times and there is no data for their assets. In these cases, we do not show any values.

11.4.4 Drawdowns

The ‘Drawdowns’ tab shows declines during the investment periods. This is shown both on a graph comparing each portfolio visible on Figure50.

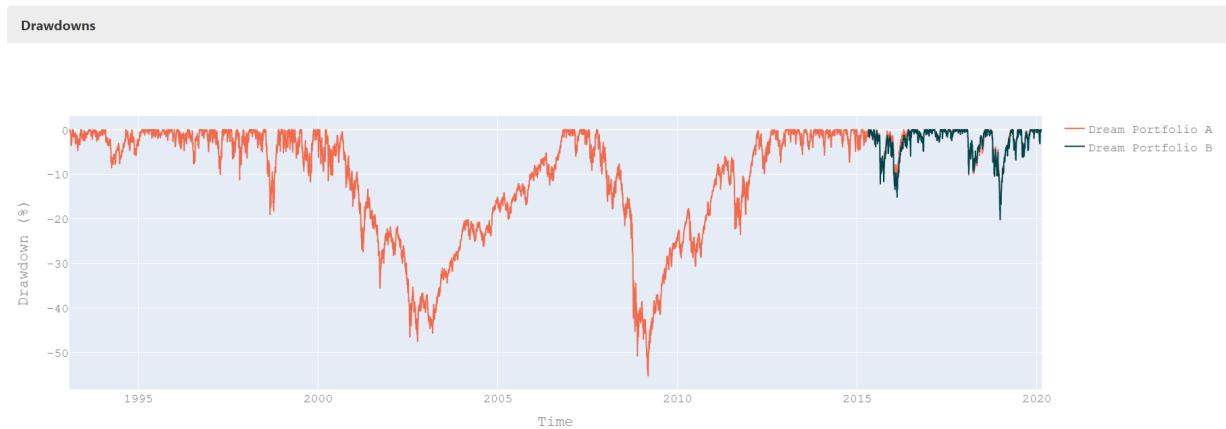


Figure 50: Drawdowns Tab - Graph (source: <http://thaliabacktest.xyz/dashboard/>)

For deeper insight, a table of the top drawdowns is generated for each portfolio. By default, this will be the top 10 or as many as exist in case the portfolio does not have 10 drawdown periods.

Dream Portfolio A Top Drawdowns						
Drawdown (%)	Start	End	Recovery	Length	Recovery Time	Underwater Period
-55.19	10/10/2007	09/03/2009	16/08/2012	1 year, 4 months	3 years	4 years
-47.52	27/03/2000	09/10/2002	26/10/2006	2 years	4 years	6 years
-19.35	21/09/2018	24/12/2018	12/04/2019	3 months	3 months	6 months
-19.03	21/07/1998	31/08/1998	23/11/1998	a month	2 months	4 months
-13.02	21/07/2015	11/02/2016	18/04/2016	6 months	2 months	8 months
-11.7	19/07/1999	15/10/1999	17/11/1999	2 months	a month	3 months
-11.2	08/10/1997	27/10/1997	05/12/1997	19 days	a month	a month
-10.1	29/01/2018	08/02/2018	06/08/2018	10 days	5 months	6 months
-10.04	19/02/1997	11/04/1997	05/05/1997	a month	24 days	2 months
-9.3	20/01/2000	25/02/2000	17/03/2000	a month	21 days	a month

Dream Portfolio B Top Drawdowns						
Drawdown (%)	Start	End	Recovery	Length	Recovery Time	Underwater Period
-20.16	21/09/2018	24/12/2018	23/04/2019	3 months	3 months	7 months
-15.11	24/06/2015	11/02/2016	08/06/2016	7 months	3 months	11 months
-9.9	29/01/2018	08/02/2018	25/07/2018	10 days	5 months	5 months
-6.77	06/05/2019	03/06/2019	20/06/2019	28 days	17 days	a month
-6.18	29/07/2019	14/08/2019	28/10/2019	16 days	2 months	2 months
-5.9	09/06/2016	27/06/2016	08/07/2016	18 days	11 days	29 days
-4.8	08/09/2016	04/11/2016	15/11/2016	a month	11 days	2 months
-3.11	21/01/2020	31/01/2020	05/02/2020	10 days	5 days	15 days
-2.77	02/03/2017	13/04/2017	27/04/2017	a month	14 days	a month
-2.43	26/07/2017	18/08/2017	11/09/2017	23 days	24 days	a month

Figure 51: Drawdowns Tab - Tables (source: <http://thaliabacktest.xyz/dashboard/>)

11.4.5 Assets

Here, the user may encounter a pie chart similar to the one on the summary tab. However, the pie chart shows the proportions of the asset classes and not the assets allocations. This is crucial information for the user as it can show how diverse their portfolio is.

Assets Breakdown

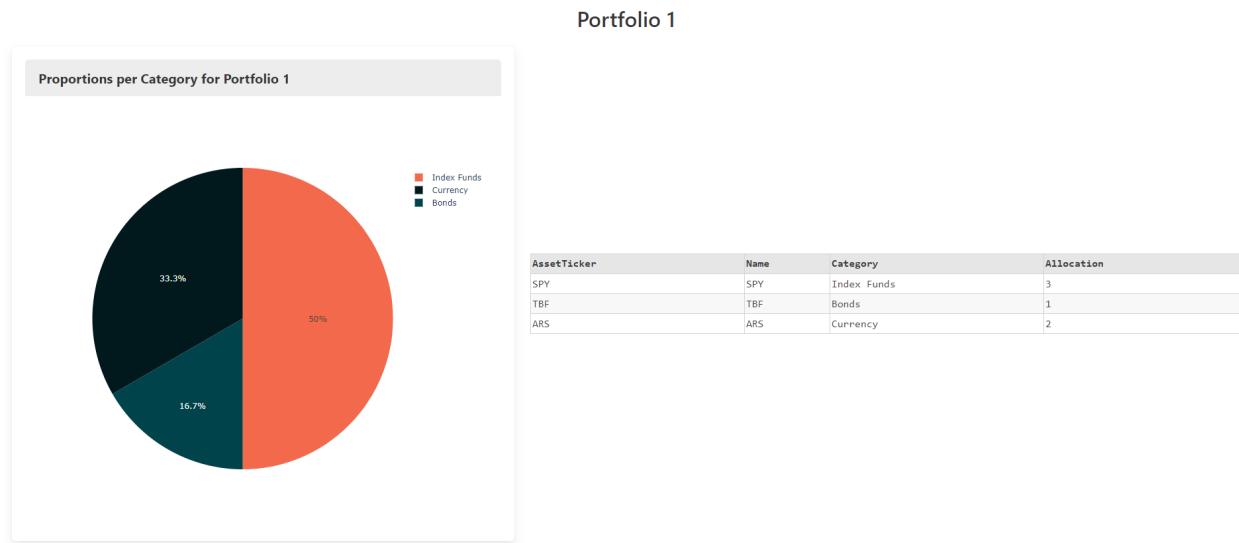


Figure 52: Assets Tab (source: <http://thaliabacktest.xyz/dashboard/>)

11.4.6 Overfitting

Our last results tab encompasses a relatively complex feature, which is overfitting. Here, the user may test if any of their portfolios exhibits bias to particular market conditions, for example by choosing a specific timeframe. Overfitting is dangerous as it may skew the investor's perception of the risk-reward characteristics

of a portfolio. In case the user decides to run the overfitting test and it fails to detect an overfitting tendency, the following message is displayed:

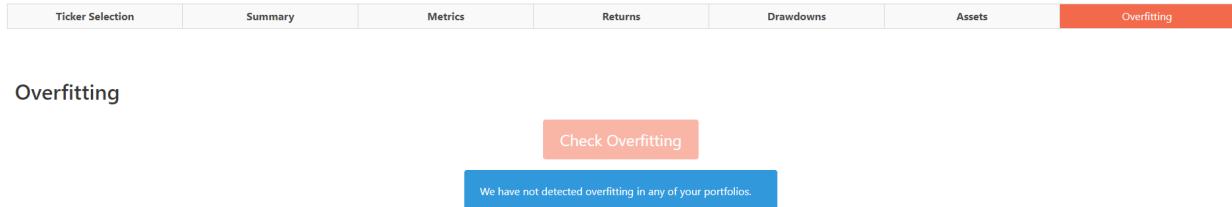


Figure 53: Overfitting Tab (source: <http://thaliabacktest.xyz/dashboard/>)

Alternatively, if the test detects overfitting, it will warn the user, highlighting the corresponding portfolio(s).

12 Appendix B - Maintenance Manual

Please keep in mind that we wrote sections such as:

↳ Early in the inception phase of development, we decided that our goal was not to have fixed responsibilities in our team, allowing everybody to work on each component of the system. This decision has also eased the code review process, as we had no status differences to distort the error-correcting mechanism [51].

So maybe not go into much detail about our specializations...

12.0.1 Maintenance Design

All the financial assets that the Data Harvester will retrieve data for are stored in a folder that contains a CSV file for each financial asset class, where the asset class name is the name of that CSV file. Each row of a CSV file will contain the ticker, the last and earliest record and the full name of an asset. If there is no data on the asset in the database, then the last and earliest records are left empty so that the Data Harvester will pull all available information on the first update of that asset.

Logs

An extensive logging mechanism that logs information from all running parts of the Data Harvester has been implemented. Logs should be consulted before and after making any modifications. All successful updating sessions have the following information:

1. Date and time of when an update has started.
2. The name of the API currently being updated and how many calls are going to be made.
3. For each API call:
 - The asset class name, asset ticker and time range for the API call.
 - The type of data frame retrieved. It is <class ‘pandas.core.frame.DataFrame’> if the API returned that the call worked.
 - The time range in the data frame retrieved. The newest date always is the date of the last trading day for that asset.
 - The shape of the data frame before the interpolation.
 - The shape after the interpolation. For those assets that are traded every day, no new rows appear after the interpolation.
 - A message that we started writing the data frame to the database.
4. A message that the requested number of API calls has been done.
5. Repeated logs from point 2 if there are other APIs to be used or a log when the update has finished (i.e. if all APIs have been used).

Maintenance of the Assets:

- **Add/remove an asset class:** Create a new CSV file with the name of the asset class as a filename within the asset classes folder and add the columns in the following form:

```
1 Ticker,Last_Update,Earliest_Record,Name  
2
```

It is important to remember to add the new asset class in the list of supported asset classes of the API that supports it.

To remove an asset class, delete the respective CSV file and re-run the initialization script so that the asset class will be removed from the future updates.

- **Add/remove an asset:** Go into the asset class CSV file that contains the asset and add the required details under the correct columns. At the end of the CSV asset class file, add a line containing the ticker and name of the asset.

Example:

```

1 Ticker ,Last_Update ,Earliest_Record ,Name
2 SPY ,,,S&P 500
3

```

To delete an asset, you simply have to remove the row that contains it.

- **Re-add an asset:** First check in the database for the last and earliest records. Then introduce a new row that contains them.

Maintenance of the APIs:

- **Adding a new API:**

The difficulty of adding a new API greatly varies. It depends on ease of API use for data access. As a general rule, the organisations that have financial incentives to provide functional APIs often have an easier integration process. On the other end, the governmental organisation or the NGOs that provide free data can have a lengthy implementation process.

Here are the steps required for adding a new API.

1. Clearly outline the reason why you are adding the API. Is it for adding new assets that were not available before? Is it to replace an older API? Are you adding redundancy for a more reliable service? In any of those cases, the appropriate adding and removing of assets and asset classes is required.
 2. Add a wrapper to the API call so that it works with the updating mechanism. The updating mechanism passes an asset ticker, an asset class name and a time range for the dates it is requesting data for. In return, it expects a data frame with data from 1970-1-1 (in the case of assets that have been traded for a long time) or the first available date for newer ones. In the case that the API has no data on the requested ticker, the update mechanism expects single value 1. The call wrapper should have a name of the form '<api_name>.call'.
 3. Create a method that formats the data frame retrieved by the API. Use the format checker method from the API Handler class to verify whether your formatting is correct.
- **The API changes over time:** While it is hard to predict what will change, we know that if any changes happen with the APIs, it will be seen from the logs.

- **Format changes:** If the format of the data retrieved changes, then the format checker combined with the API logs will be able to tell what is the exact new format and what has to be changed and what the maintainer has to do to fix the problem.
- **Network Errors:** Depending on the error number, appropriate action should be taken. The network errors are match HTTP status codes.
- **Empty data frame:** It is easy to see if the data frame retrieved is of the right size by looking at the logs. In this case, there is likely a problem with the API and extra work is required to find a solution.
- **Malformed data:** This is hard to verify and unlikely to happen. If an API is returning malformed data for all assets, then it should not be used for any further updates. If an API is working correctly but it provides bad data for one asset, then that asset should be removed. Depending on the combination of assets and API, one can check for malformed data by implementing cross-checks with multiple APIs for the same assets.
- **APIs Internal Errors:** Verify the latest documentation and perform the required modifications of the calling methods. There is a good chance that our code is using an outdated version of the API.

13 Appendix C - Problem Domain Glossary

- Portfolio - A portfolio is a grouping of financial assets such as stocks, bonds, commodities, currencies and cash equivalents, as well as their fund counterparts, including mutual, exchange-traded and closed funds. An investor's portfolio is the group of assets they have currently invested in.
[<https://www.investopedia.com/terms/p/portfolio.asp>]
- Asset - Generally an asset that gets its value from being owned; can be traded on financial markets. Stocks, bonds, commodities, (crypto-)currencies are all types of financial assets.
[<https://www.investopedia.com/terms/a/asset.asp>]
- Asset Class - An asset class is a grouping of investments that exhibit similar characteristics and are subject to the same laws and regulations. Asset classes are made up of instruments which often behave similarly to one another in the marketplace.
[<https://www.investopedia.com/terms/a/assetclasses.asp>]
- Backtesting - Backtesting is the general method for seeing how well a strategy or model would have done ex-post. Backtesting assesses the viability of a trading strategy by discovering how it would play out using historical data.
[<https://www.investopedia.com/terms/b/backtesting.asp>]
- Standard Strategy / Lazy Portfolios - A lazy portfolio is a collection of investments that require very little maintenance.
[<https://www.thebalance.com/how-to-build-the-best-lazy-portfolio-2466533>]
- Rebalancing - Rebalancing is the process of realigning the weightings of a portfolio of assets. Rebalancing involves periodically buying or selling assets in a portfolio to maintain an original or desired level of asset allocation or risk.
[<https://www.investopedia.com/terms/r/rebalancing.asp>]
- Key metrics - Performance measures of a portfolio that are of high interest to the majority of investors.
- Standard Deviation - The standard deviation is a statistic that measures the dispersion of a dataset relative to its mean.
[<https://www.investopedia.com/terms/s/standarddeviation.asp>]
- Worst Year - The worst performance over any given 365 day period starting from January 1st of some year.
- Sharpe Ratio - The Sharpe ratio was developed by Nobel laureate William F. Sharpe and is used to help investors understand the return of an investment compared to its risk.
[<https://www.investopedia.com/terms/s/sharperatio.asp>]
- Sortino Ratio - The Sortino ratio is a variation of the Sharpe ratio that differentiates harmful volatility from total overall volatility by using the asset's standard deviation of negative portfolio returns, called downside deviation, instead of the total standard deviation of portfolio returns.
[<https://www.investopedia.com/terms/s/sortinoratio.asp>]
- Inflation - Inflation is a quantitative measure of the rate at which the average price level of a basket of selected goods and services in an economy increases over a period of time.
[<https://www.investopedia.com/terms/i/inflation.asp>]
- Nominal Values - A value that is unadjusted for inflation.
- Real Values - A value that is adjusted for inflation.
- Equity - Equity is typically referred to as shareholder equity (also known as shareholders' equity) which represents the amount of money that would be returned to a company's shareholders if all of the assets were liquidated and all of the company's debt was paid off.

- Fixed Income - Fixed income is a type of investment security that pays investors fixed interest payments until its maturity date.
[\[https://www.investopedia.com/terms/f/fixedincome.asp\]](https://www.investopedia.com/terms/f/fixedincome.asp)
- Commodity - A commodity is a basic good used in commerce that is interchangeable with other commodities of the same type. Commodities are most often used as inputs in the production of other goods or services. The quality of a given commodity may differ slightly, but it is essentially uniform across producers.
[\[https://www.investopedia.com/terms/c/commodity.asp\]](https://www.investopedia.com/terms/c/commodity.asp)
- FOREX / FX - Forex (FX) is the marketplace where various national currencies are traded. The forex market is the largest, most liquid market in the world, with trillions of dollars changing hands every day.
[\[https://www.investopedia.com/terms/f/forex.asp\]](https://www.investopedia.com/terms/f/forex.asp)
- Overfitting - Overfitting is a modeling error that occurs when a function is too closely fit for a limited set of data points.
[\[https://www.investopedia.com/terms/o/overfitting.asp\]](https://www.investopedia.com/terms/o/overfitting.asp)
- Leverage - Leverage results from using borrowed capital as a funding source when investing to expand the firm's asset base and generate returns on risk capital.
[\[https://www.investopedia.com/terms/l/leverage.asp\]](https://www.investopedia.com/terms/l/leverage.asp)
- Technical Analysis - Technical analysis is a trading discipline employed to evaluate investments and identify trading opportunities by analyzing statistical trends gathered from trading activity, such as price movement and volume.
[\[https://www.investopedia.com/terms/t/technicalanalysis.asp\]](https://www.investopedia.com/terms/t/technicalanalysis.asp)

14 Appendix D - Project Terminology

- Anda - A library that calculates a portfolio's historical performance and associated risk metrics. The name has been chosen as the abbreviation of 'analyse data'.
- Data Harvester - Component of our system that fetches data from third party financial data APIs and writes it to the Thalia database. Retrieves both historical and live price data as well as dividends data for equities.
- Finda - An interface for our databases. Provides functionality for registering, creating and connecting to databases as well as reading from and writing data to them. The name has been chosen as the abbreviation of 'find data'.
- Thalia Web - Shorthand for the Thalia Web Server providing the service our customers interact with.

15 Appendix E - UML Diagrams

The UML Diagrams for the business logic library 'Anda' are the only ones that show the parameter type for the methods as well as the return types. This is because we believe that the discussion of Anda requires more detail, since it deals with numerical values such as 'Decimal' or 'float'.

15.1 Anda

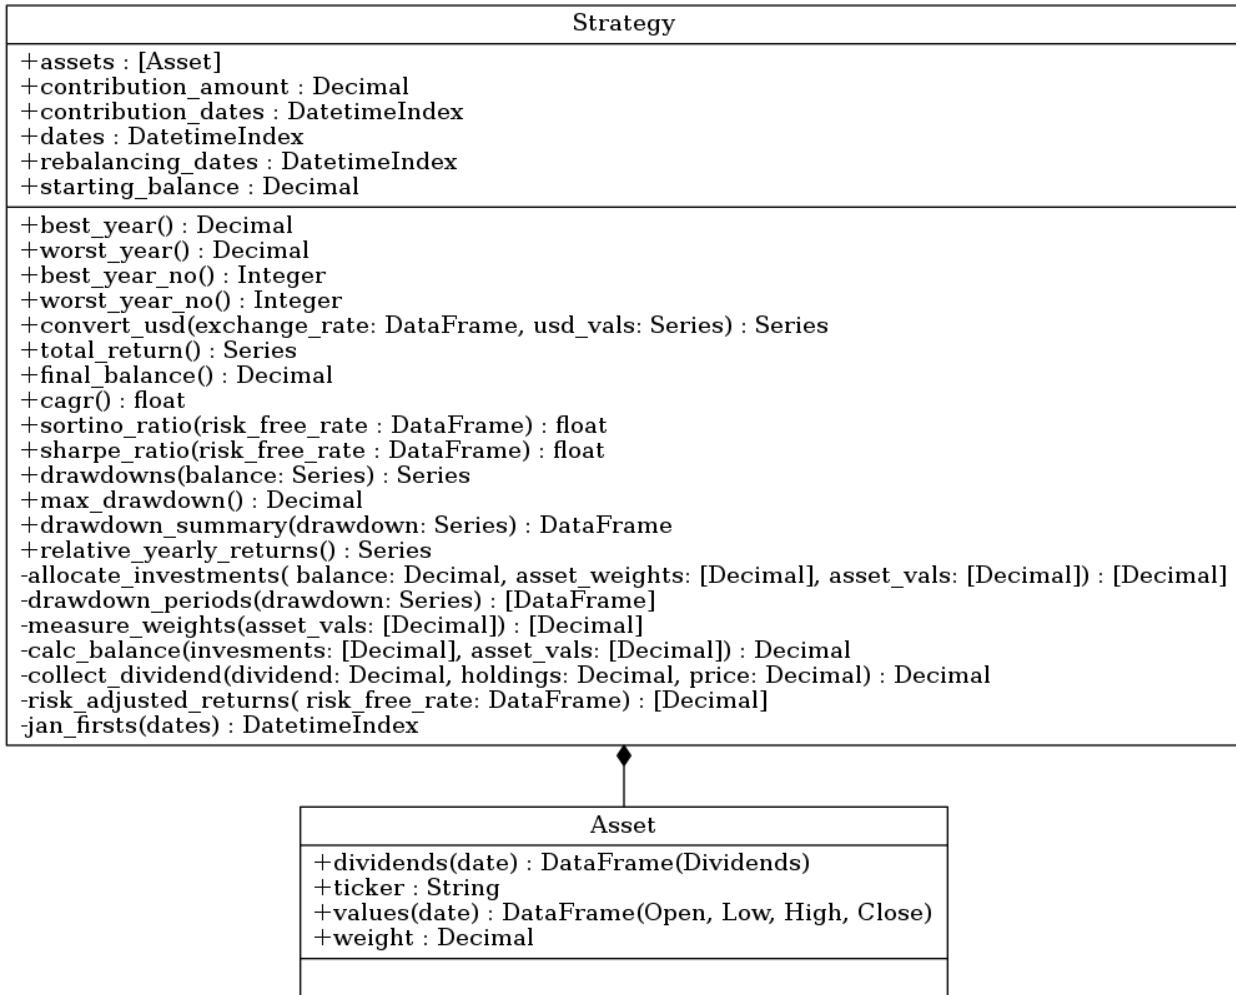


Figure 54: UML Diagrams - Anda

15.2 Data Harvester

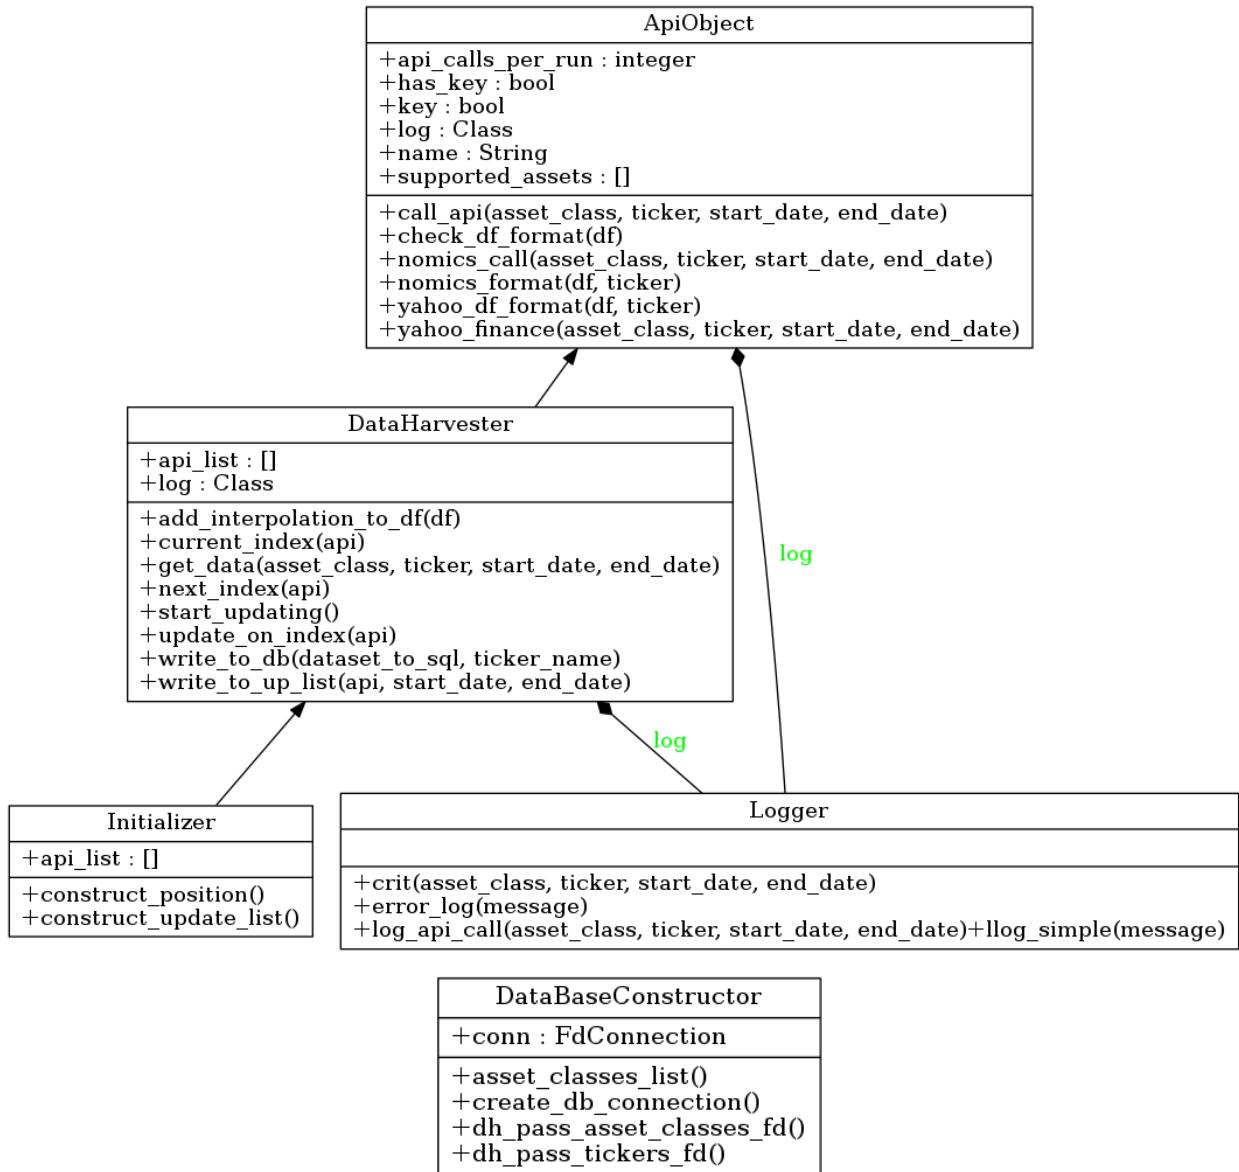


Figure 55: UML Diagrams - Data Harvester

15.3 Finda

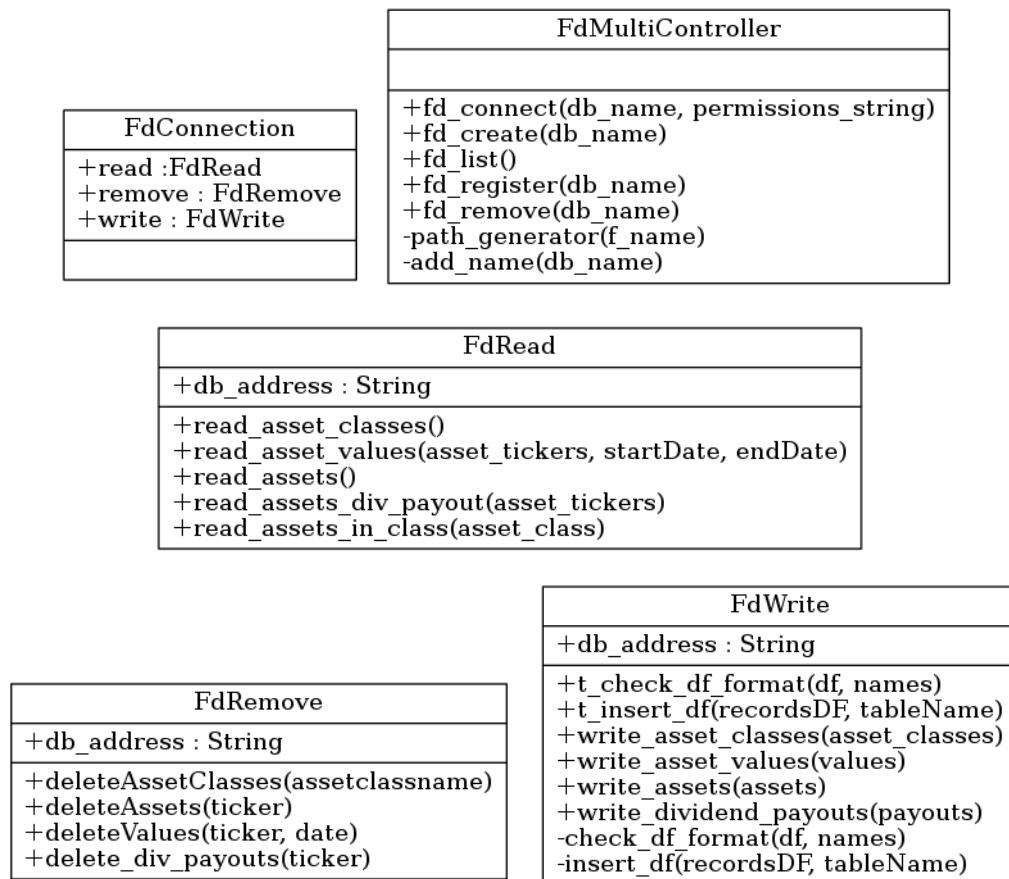


Figure 56: UML Diagrams - Finda

15.4 Thalia Web

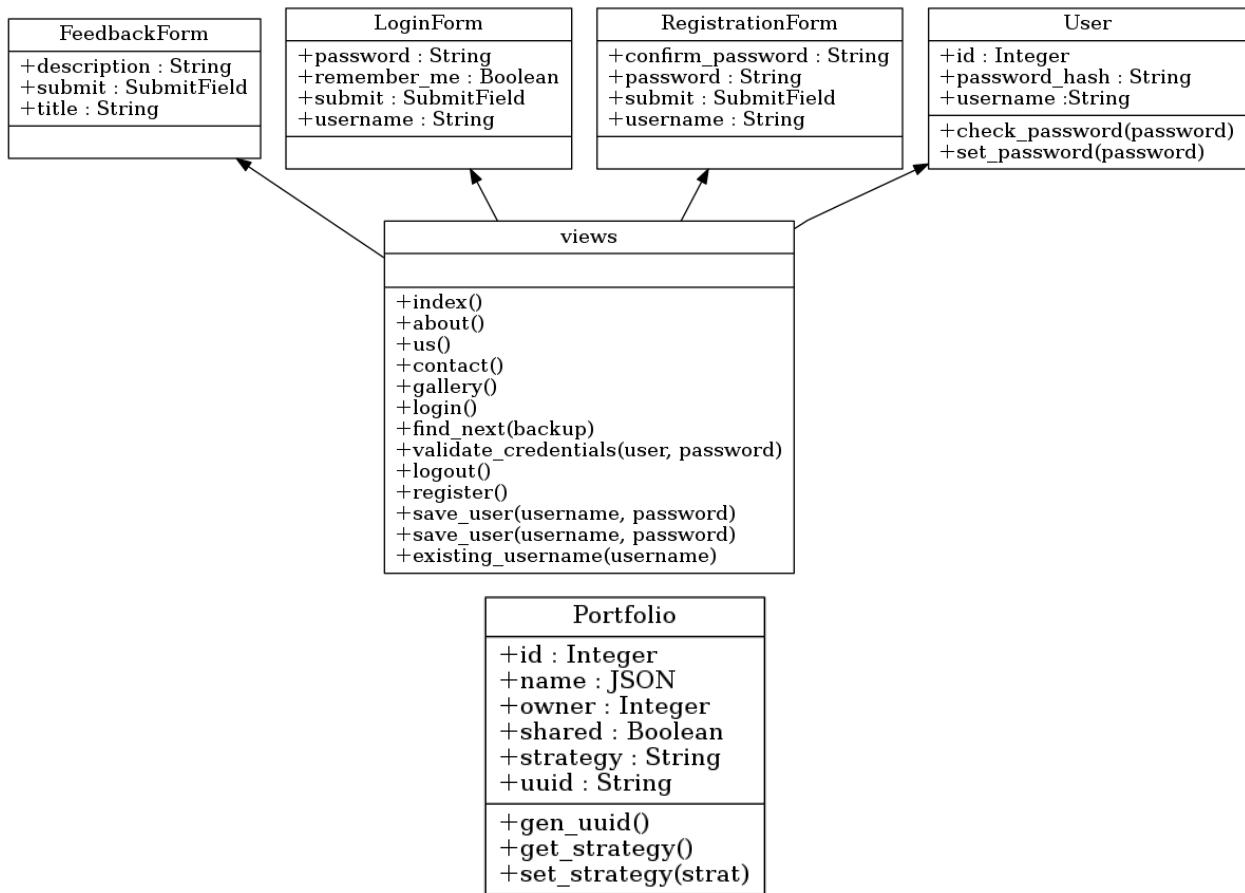


Figure 57: UML Diagrams - Thalia Web

16 Appendix F - User Types Diagrams

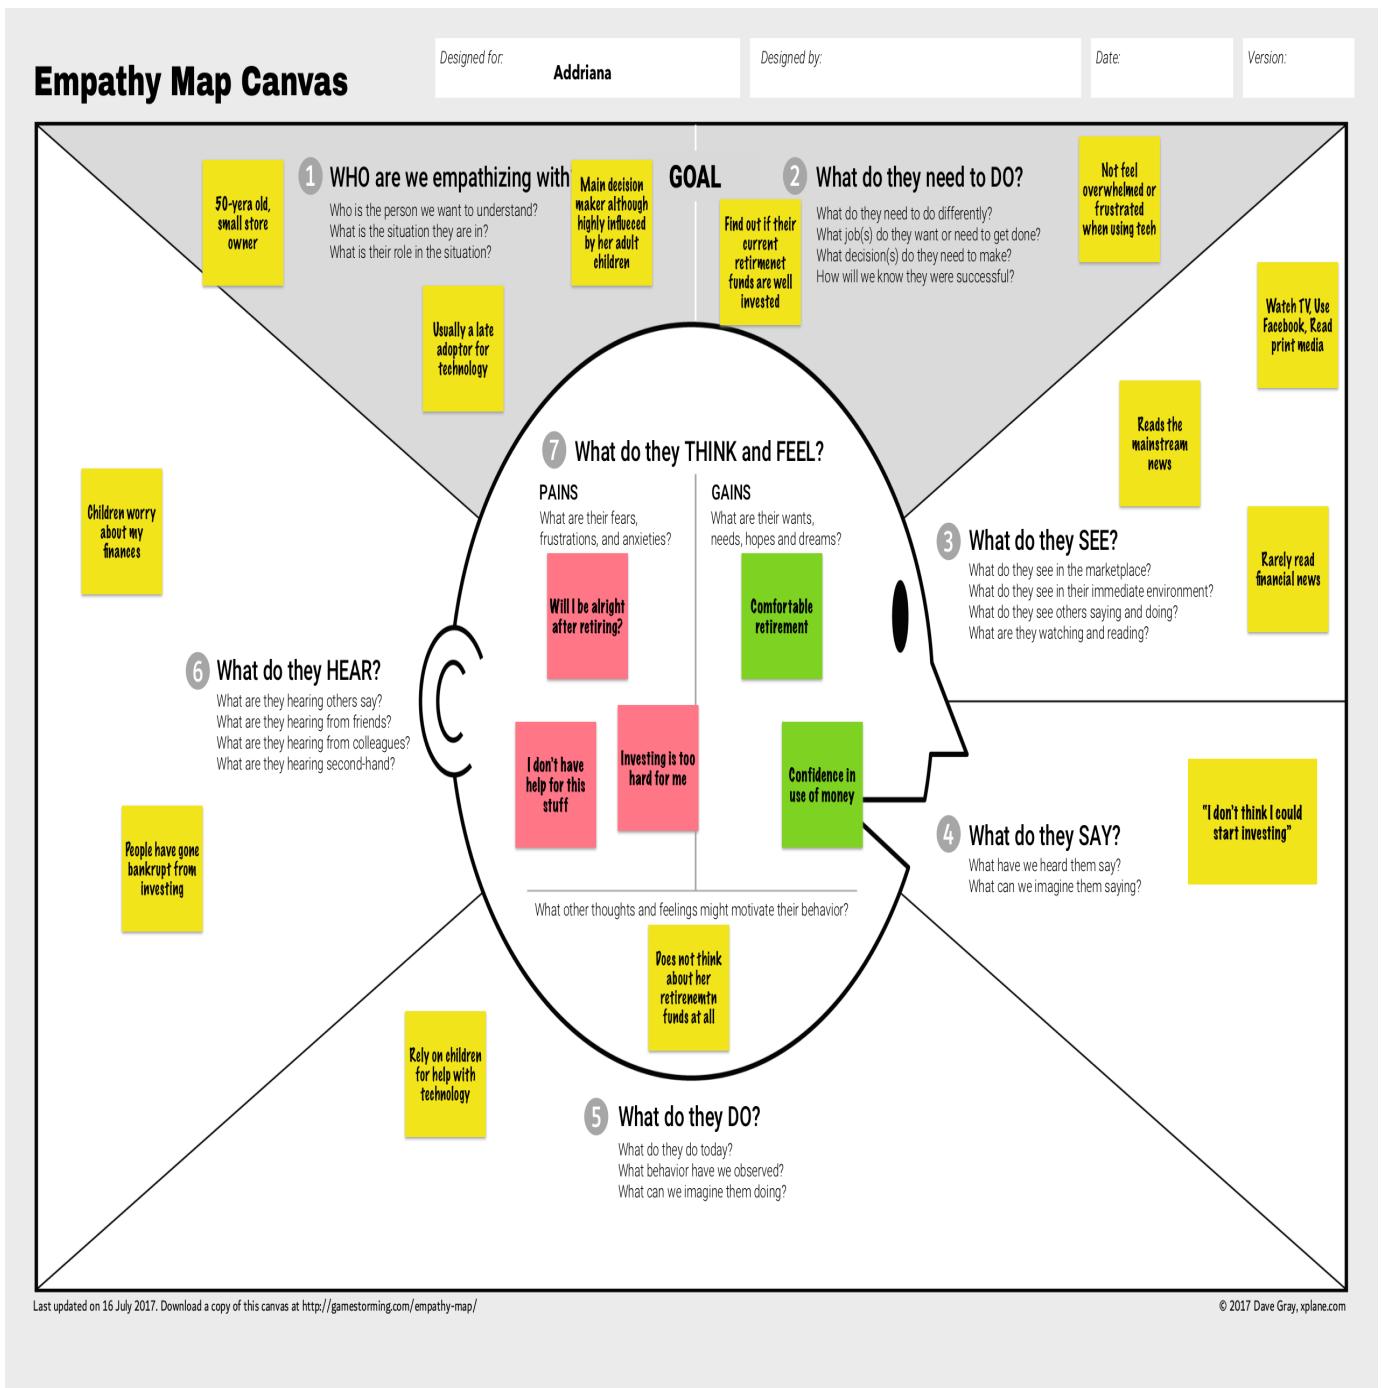


Figure 58: Addriana Empathy Map

Empathy Map Canvas

Designed for: Allison

Designed by:

Date:

Version:

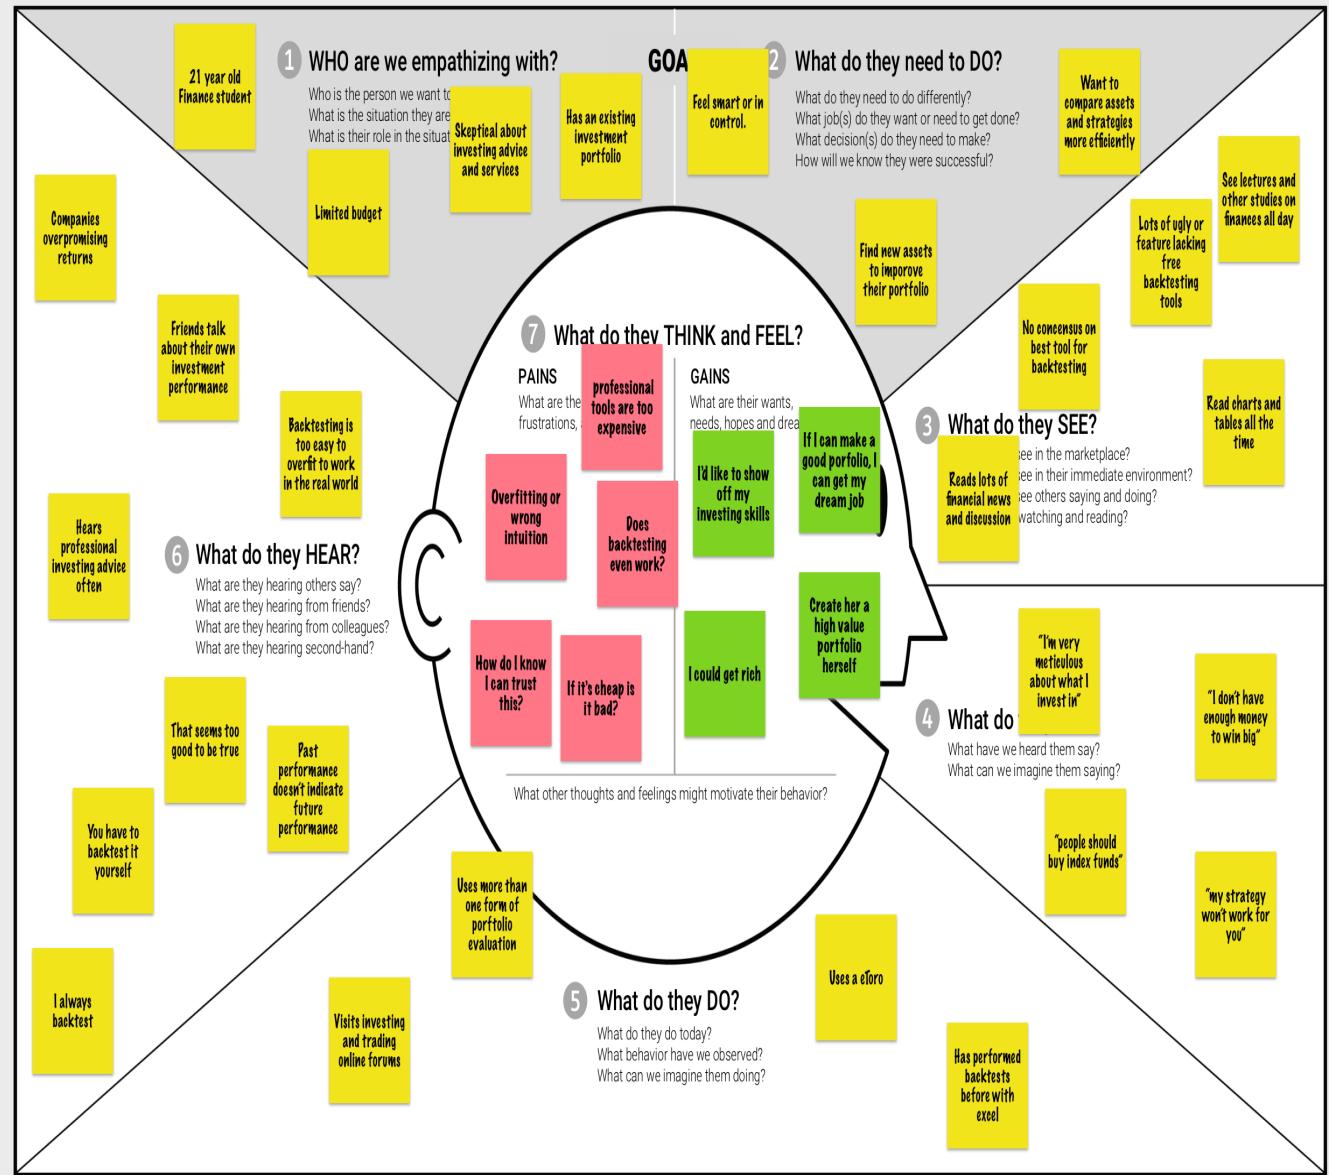


Figure 59: Allison Empathy Map

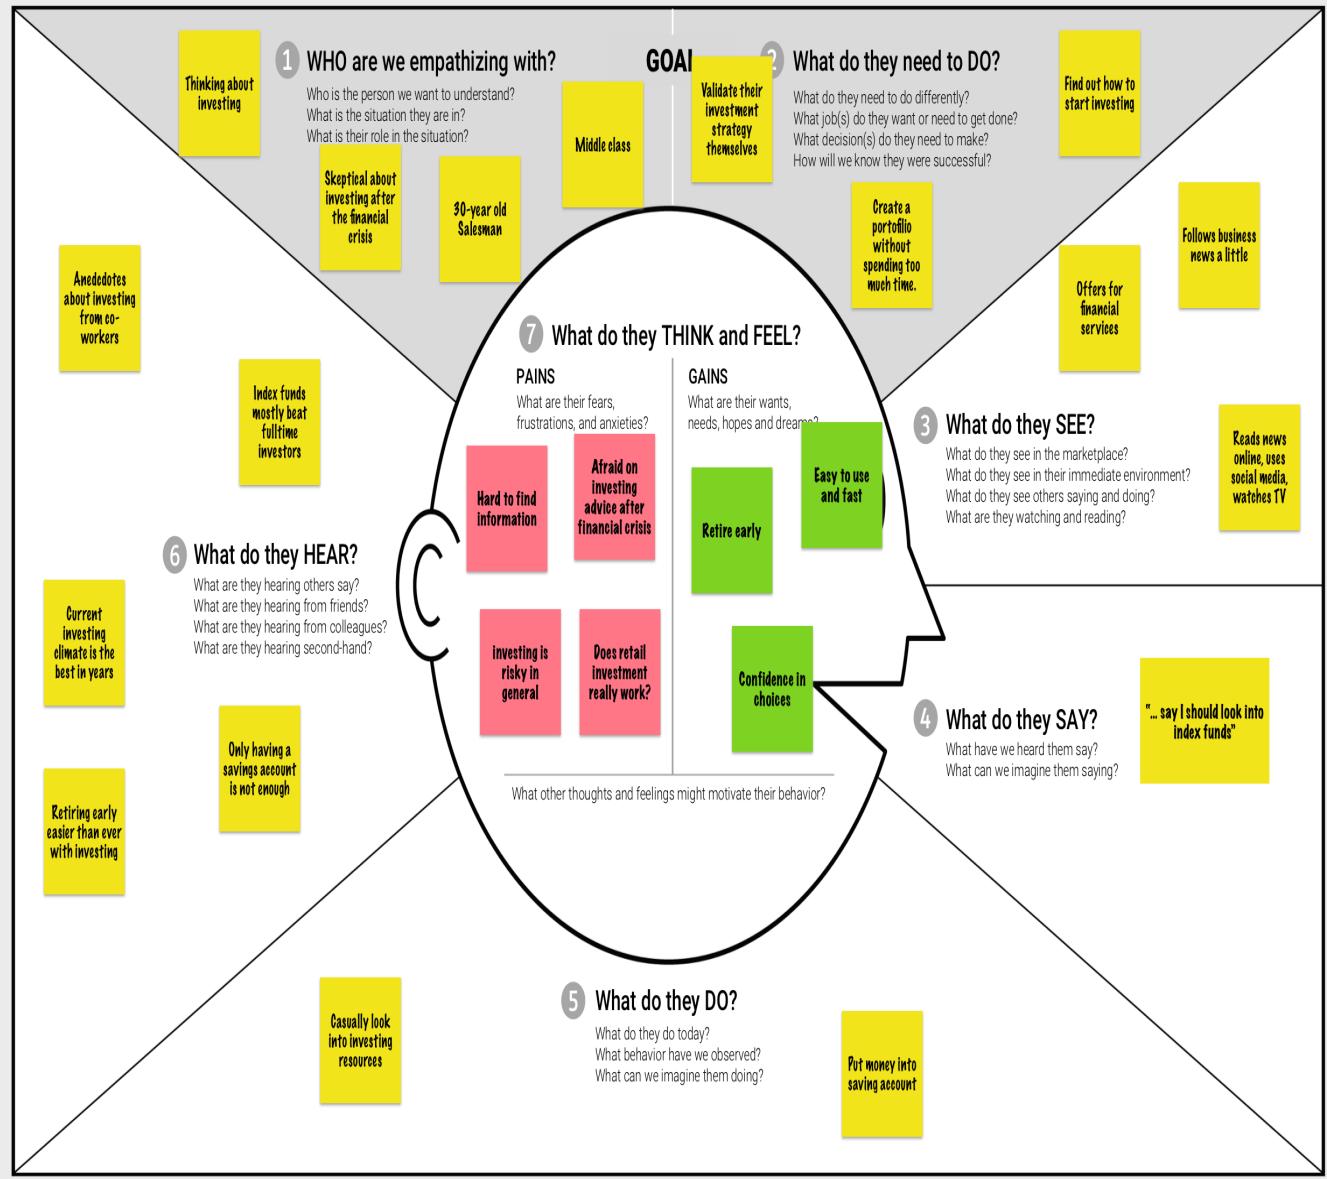
Empathy Map Canvas

Designed for: Sam

Designed by:

Date:

Version:



Last updated on 16 July 2017. Download a copy of this canvas at <http://gagelab.com/empathy-map/>

© 2017 Dave Gray, xplane.com

Figure 60: Sam Empathy Map

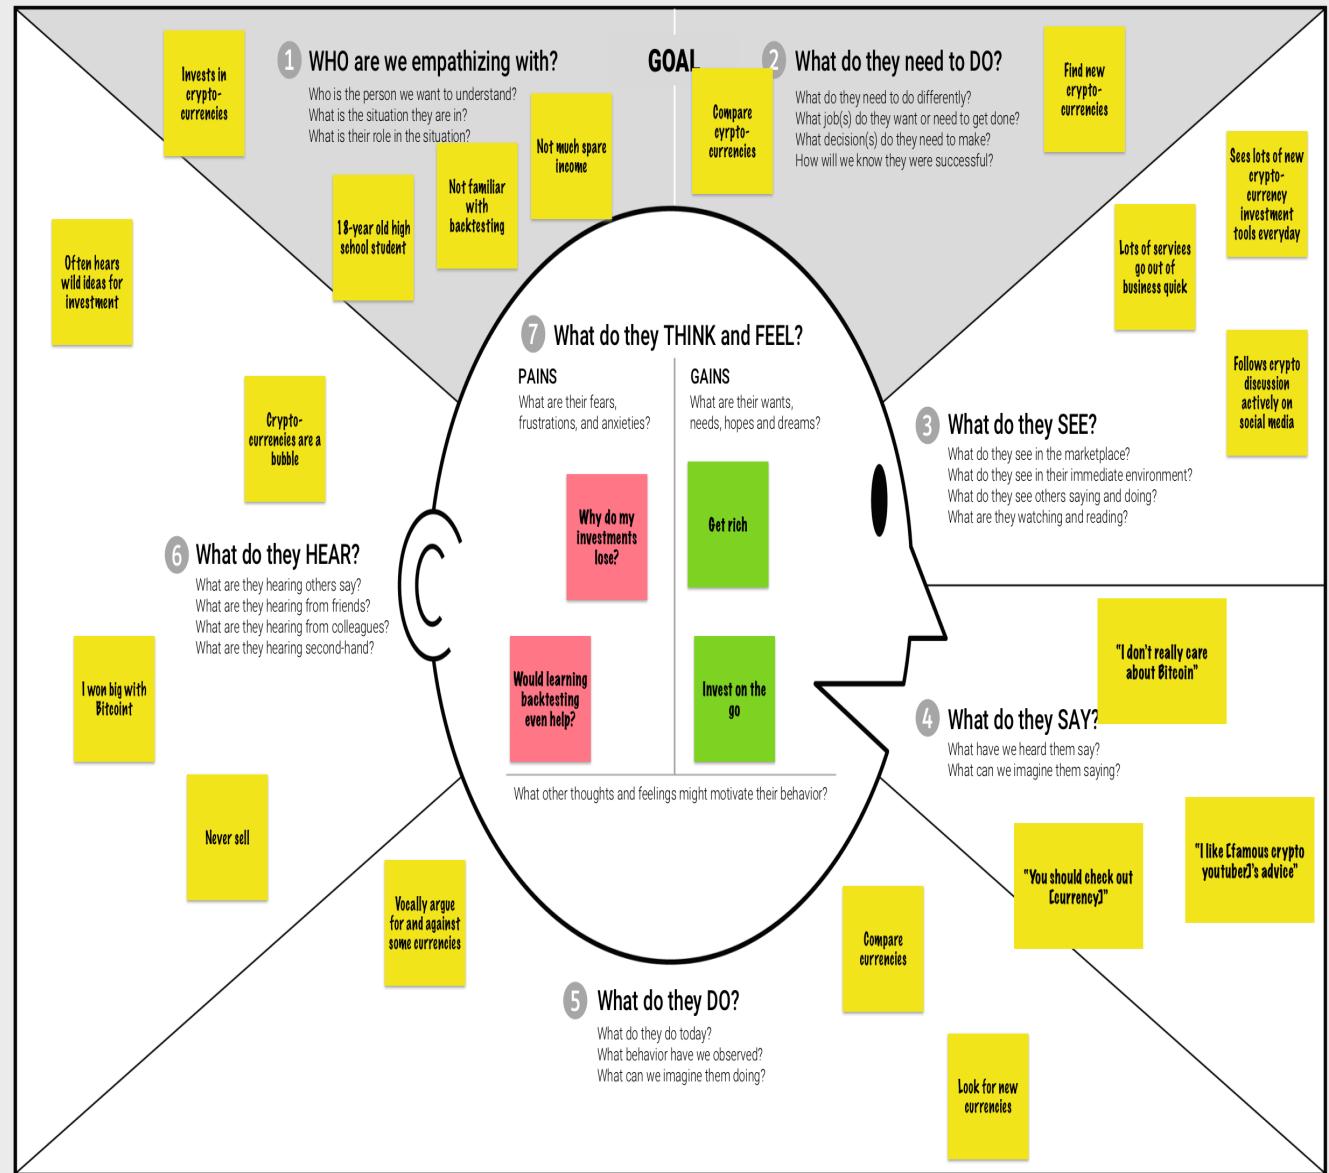
Empathy Map Canvas

Designed for: Timmy

Designed by:

Date:

Version:



Last updated on 16 July 2017. Download a copy of this canvas at <http://gagelab.com/empathy-map/>

© 2017 Dave Gray, xplane.com

Figure 61: Timmy Empathy Map

PERSONA

Personas are used to capture the type of people involved in and using your service. They are not necessarily a representation of an individual, but show the characters, opinions and qualities of a group of users. Try to create 5 personas associated with your service.

What do they look like? Draw/create an image here.



Photo by Artem Beliaikin on Unsplash

WHAT ARE THEY MOST LIKELY TO SAY ABOUT THE SERVICE?

“

I don't understand

”

Try out investing

WHAT WOULD THEY USE THIS SERVICE FOR?

NAME:

Addriana

AGE:

50

OCCUPATION:

Small store manager

LIKES:

Simple things

DISLIKES:

Technology

FAMILY AND FRIENDS:

Single



Figure 62: Addriana Persona

PERSONA

Personas are used to capture the type of people involved in and using your service. They are not necessarily a representation of an individual, but show the characters, opinions and qualities of a group of users. Try to create 5 personas associated with your service.

What do they look like? Draw/create an image here.



Photo by Max Ilienerwise on Unsplash

WHAT ARE THEY MOST LIKELY TO SAY ABOUT THE SERVICE?

“

easy and cheap option

”,

NAME:

Allison

AGE:

21

OCCUPATION:

Finance Student

LIKES:

Finance, mobile UI

DISLIKES:

Lack of feedback

FAMILY AND FRIENDS:

Finance students

WHAT WOULD THEY USE THIS SERVICE FOR?

Help validate strategy



Figure 63: Allison Persona

PERSONA

Personas are used to capture the type of people involved in and using your service. They are not necessarily a representation of an individual, but show the characters, opinions and qualities of a group of users. Try to create 5 personas associated with your service.

What do they look like? Draw/create an image here.



Photo by Andrew Robinson on Unsplash

NAME:

Sam

AGE:

30

OCCUPATION:

B2B Salesman

LIKES:

Trying new things

DISLIKES:

Feeling left out

FAMILY AND FRIENDS:

Wife

WHAT WOULD THEY USE THIS SERVICE FOR?

Come up with an investment
strategy



WHAT ARE THEY MOST LIKELY TO SAY ABOUT THE SERVICE?

“

Looks promising

”

Figure 64: Sam Persona

PERSONA

Personas are used to capture the type of people involved in and using your service. They are not necessarily a representation of an individual, but show the characters, opinions and qualities of a group of users. Try to create 5 personas associated with your service.

What do they look like? Draw/create an image here.

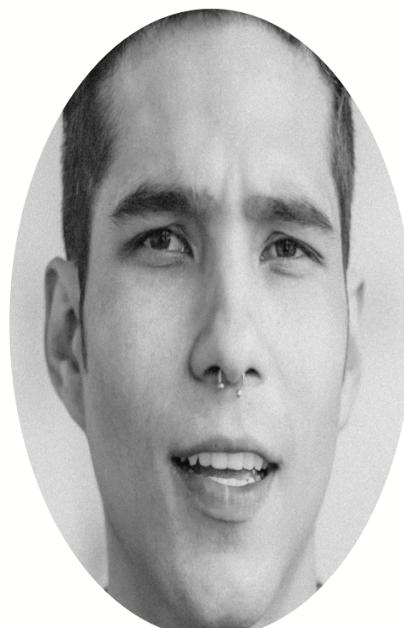


Photo by Roberto Delgado Webb on Unsplash

WHAT ARE THEY MOST LIKELY TO SAY ABOUT THE SERVICE?

“

they've got all the cryptos
”,

NAME:

Timmy

AGE:

18

OCCUPATION:

High school student

LIKES:

Cryptocurrencies

DISLIKES:

Studying

FAMILY AND FRIENDS:

Lives with parents

WHAT WOULD THEY USE THIS SERVICE FOR?

Compare cryptocurrencies



Figure 65: Timmy Persona

References

- [1] M. Aukia, A. Boehm, A.L. Heath, D. Joffe, G. Stoian, and M. Veiner. Thalia White Paper, 2019.
- [2] M. Aukia, A. Boehm, A.L. Heath, D. Joffe, G. Stoian, and M. Veiner. Thalia Technical Report, 2019.
- [3] PayScale. Average Software Developer Salary in United Kingdom. Available at: https://www.payscale.com/research/UK/Job=Software_Developer/Salary.
- [4] B. M. Ferreira, S. D. J. Barbosa, and T. Conte. PATHY: Using Empathy with Personas to Design Applications that Meet the Users' Needs, 2016.
- [5] Interaction Design Foundation. Ux Design Courses for Beginners and Professionals. Available at: <https://www.interaction-design.org/>.
- [6] R. F. Dam and Y. S. Teo. Personas – A Simple Introduction. Available at: <https://www.interaction-design.org/literature/article/personas-why-and-how-you-should-use-them>.
- [7] L. H. Shawn. Accessibility in User-Centered Design: Personas. Available at: <http://uiaccess.com/accessucd/personas.html>.
- [8] D. Bland. Agile Coaching Tip: What Is an Empathy Map? Available at: <https://www.solutionsiq.com/resource/blog-post/what-is-an-empathy-map/>.
- [9] Quantpedia. A comprehensive list of tools for quantitative traders. Available at: <https://quantpedia.com/links-tools/?category=historical-data>.
- [10] Fidelity. Fidelity Active Trader Pro. Available at: <https://www.fidelity.com/trading/advanced-trading-tools/active-trader-pro/overview-c>.
- [11] MetaTrader. MetaTrader4 Trading Platform. Available at: <https://www.metatrader4.com/en>.
- [12] NinjaTrader. NinjaTrader Platform. Available at: <https://ninjaTrader.com/>.
- [13] Portfolio Visualizer. Portfolio Visualizer Backtesting Tool. Available at: <https://www.portfoliovisualizer.com/backtest-portfolio>.
- [14] ETFReplay. Etfreplay Backtesting Tool. Available at: <https://www.etfreplay.com/combine.aspx>.
- [15] Backtest Curvo. Backest Curvo Backtesting Tool. Available at: <https://backtest.curvo.eu/>.
- [16] Stockbacktest. Stockbacktest Backtesting Tool. Available at: <http://www.stockbacktest.com/>.
- [17] J. Williams. The Basics of Branding. Available at: <https://www.entrepreneur.com/article/77408>.
- [18] R. Clifton. *Brands and Branding*. Economist Books, 2009.
- [19] Pantone. Pantone Color of the Year 2020. Available at: <https://www.pantone.com/color-intelligence/color-of-the-year/color-of-the-year-2020>.
- [20] Sessions College. Triadic Color Scheme Wheel. Available at: <https://www.sessions.edu/color-calculator/>.
- [21] San Jose State University. FURPS+ Model. Available at: <http://www.cs.sjsu.edu/faculty/pearce/modules/lectures/ooa/requirements/IdentifyingURPS.htm>.
- [22] A. B. Al-Badareen, M. H. Selamat, M. A. Jabar, J. Din, and S. Turaev. Software Quality Models: A Comparative Study. *Communications in Computer and Information Science*, 2011.
- [23] R. Dhall. Designing Graceful Degradation in Software Systems. *Proceedings of the Second International Conference on Research in Intelligent and Computing in Engineering*, 10(-):171–179, 2017.

- [24] D. P. Drljaca and B. Latinovic. Using TELOS for the planning of the information system audit. *IOP Conference Series: Materials Science and Engineering*, 294:012022, 2018.
- [25] B. Olsen and L. Kolvereid. Development of new ventures over time: strategy, profitability and growth in new Scandinavian firms. *Entrepreneurship & Regional Development*, 6(4):357–370, 1994.
- [26] FCA. Understanding ‘advice’ and ‘guidance’ on investments, 2018. Available at: <https://www.fca.org.uk/consumers/understanding-advice-guidance-investments>.
- [27] NumFOCUS. pandas - Python Data Analysis Library. Available at: <https://pandas.pydata.org/>.
- [28] W. McKinney et al. pandas: a foundational Python library for data analysis and statistics. *Python for High Performance and Scientific Computing*, 14(9), 2011.
- [29] MDN Web Docs. HTML5 Developer Guides. Available at: <https://developer.mozilla.org/en-US/docs/Web/Guide/HTML/HTML5>.
- [30] Bulma. Bulma: Free, Open Source, and Modern CSS Framework based on Flexbox. Available at: <https://bulma.io/>.
- [31] plotly. Introducing Dash. Available at: <https://medium.com/plotly/introducing-dash-5ecf7191b503>.
- [32] MDN Web Docs. Web Technology for Developers - JavaScript. Available at: <https://developer.mozilla.org/en-US/docs/Web/JavaScript>.
- [33] plotly. Dash Layout. Available at: <https://dash.plotly.com/layout>.
- [34] plotly. Dash DataTable. Available at: <https://dash.plotly.com/datatable>.
- [35] R. Brath and M. Peters. Dashboard Design: Why Design is important. *DM Review Online*, 2004.
- [36] A. Janes, a. Sillitti, and G. Succi. Effective Dashboard Design. *CUTTER IT JOURNAL*, 26, 2004.
- [37] Python Software Foundation. Decimal Fixed Point and Floating Point Arithmetic, 2020. Available at: <https://docs.python.org/3/library/decimal.html>.
- [38] S. v.d. Walt, S. C. Colbert, and G. Varoquaux. The NumPy array: a structure for efficient numerical computation. *Computing in Science & Engineering*, 13(2):22–30, 2011.
- [39] T. Facen. Nomics API Documentation. Available at: <https://pypi.org/project/nomics-python/>.
- [40] apilayer. Fixer API Documentation. Available at: <https://fixer.io/documentation>.
- [41] R. Aroussi. Yahoo Finance API Wrapper Documentation. Available at: <https://pypi.org/project/yfinance/>.
- [42] Alpha Vantage. Alpha Vantage API Documentation. Available at: <https://www.alphavantage.co/documentation/>.
- [43] Quandl Inc. Quandl API Documentation. Available at: <https://www.quandl.com/>.
- [44] J. v. d. Bossche and K. Sheppard. Pandas Datareader Documentation. Available at: <https://pandas-datareader.readthedocs.io/en/latest/>.
- [45] The Internet Society. Common Format and MIME Type for Comma-Separated Values (CSV) Files. Available at: <https://tools.ietf.org/html/rfc4180>.
- [46] Cisco Security. A Framework to Protect Data Through Segmentation. Available at: https://tools.cisco.com/security/center/resources/framework_segmentation.
- [47] C. Dye. Multiple Schema Versus Multiple Databases. Available at: <https://www.oreilly.com/library/view/oracle-distributed-systems/1565924320/ch01s04.html>.

- [48] E. Gamma, R. Helm, R. Johnson, and J. Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1994.
- [49] D. Thomas and A. Hunt. *The Pragmatic Programmer - 20th Anniversary Edition*. The Pragmatic Bookshelf, 2019.
- [50] S. Nanz and C. A. Furia. A Comparative Study of Programming Languages in Rosetta Code. *IEEE/ACM 37th IEEE International Conference on Software Engineering*, 2015.
- [51] A. G. Edge and W. Remus. The Impact of Hierarchical and Egalitarian Organization Structure on Group Decision Making and Attitudes. *Developments in Business Simulation & Experiential Learning*, 11, 1984.
- [52] H. Sousa. How to code review in a Pull Request. Available at: <https://blog.codacy.com/how-to-code-review-in-a-pull-request/>.
- [53] A. Cockburn and L. Williams. The Costs and Benefits of Pair Programming. Available at: <https://collaboration.csc.ncsu.edu/laurie/Papers/XPSardinia.PDF>, 2000.
- [54] J. Bancroft-Connors. Why physical task boards still matter. Available at: <https://www.leadingagile.com/2015/05/why-physical-task-boards-still-matter/>.
- [55] P. Smyth. Creating Web APIs with Python and Flask, 2018. Available at: <https://programminghistorian.org/en/lessons/creating-apis-with-python-and-flask>.
- [56] Pallets. Flask Template Documentation, 2010. Available at: <https://flask.palletsprojects.com/en/1.1.x/templates/>.
- [57] Flask-Login Team. Available at: <https://flask-login.readthedocs.io/en/latest/>.
- [58] WTForms Team. WTForms Documentation. Available at: <https://wtforms.readthedocs.io/en/stable/forms.html>.
- [59] plotly. Sharing data between Dash Callbacks. Available at: <https://dash.plot.ly/sharing-data-between-callbacks>.
- [60] plotly. Advanced Callbacks. Available at: <https://dash.plotly.com/advanced-callbacks>.
- [61] plotly. Dash Dev Tools. Available at: <https://dash.plotly.com/devtools>.
- [62] plotly Dash Community. Dash Dynamic Callbacks Possible Solution. Available at: <https://community.plotly.com/t/dynamic-controls-and-dynamic-output-components/5519/2>.
- [63] plotly. Sharing Data Between Callbacks. Available at: <https://dash.plotly.com/sharing-data-between-callbacks>.
- [64] plotly. Integrating Dash with Existing Web Apps. Available at: <https://dash.plot.ly/integrating-dash>.
- [65] J. Thomas. Bulma Documentation, 2020. Available at: <https://bulma.io/documentation/overview/responsiveness/>.
- [66] G. Becedillas. PyAlgoTrade Documentation. Available at: <https://github.com/gbeced/pyalgotrade>.
- [67] P. Morissette. bt Documentation. Available at: <https://pmorissette.github.io/bt/>.
- [68] UK Government. The Data Protection Act. Available at: <https://www.gov.uk/data-protection>.
- [69] Db-Engines. Database Engine Ranking. Available at: <https://db-engines.com/en/ranking>.
- [70] Datanyze. Databases Market Share Report. Available at: <https://www.datanyze.com/market-share/databases-272>.
- [71] A. M. Al Balushi. A Framework for Data Migration between Various Types of Relational Database Management Systems. *International Journal of Computer Applications*, 123(8):38–43, 2015.

- [72] S. Pearlman. Understanding Data Migration: Strategy and Best Practices. <https://www.talend.com/resources/understanding-data-migration-strategies-best-practices/>.
- [73] E. Mueller. What is DevOps? Available at: <https://theagileadmin.com/what-is-devops/>.
- [74] I. Stapleton. Flake8 Official Documentation. Available at: <https://flake8.pycqa.org/en/latest/manpage.html>.
- [75] L. Langa. Black Documentation. Available at: <https://black.readthedocs.io/en/stable/>.
- [76] Software Freedom Conservancy. githooks - Hooks used by Git. Available at: <https://git-scm.com/docs/githooks>.
- [77] Y. Zhao, A. Serebrenik, Y. Zhou, V. Filkov, and B. Vasilescu. The impact of continuous integration on other software development practices: A large-scale empirical study. *32nd IEEE/ACM International Conference on Automated Software Engineering (ASE)*, 2017.
- [78] J. Holck and N. Jørgensen. Continuous Integration and Quality Assurance: A Case Study of Two Open Source Projects. *Australasian Journal of Information Systems*, Special Issue 2003/2004.
- [79] Docker Inc. What is a container? Available at: <https://www.docker.com/resources/what-container>.
- [80] Docker Inc. Docker Documentation. Available at: <https://docs.docker.com/>.
- [81] CircleCI. CircleCI Continuous Integration and Delivery. Available at: <https://circleci.com/>.
- [82] R. Connolly. *Fundamentals of Web development*. Pearson Education, 2015.
- [83] Heroku. Heroku Cloud Application Platform, 2020. Available at: <https://www.heroku.com/>.
- [84] Heroku. SQLite on Heroku, 2017. Available at: <https://devcenter.heroku.com/articles/sqlite3>.
- [85] Amazon Web Services. Aws, 2020. Available at: <https://aws.amazon.com/>.
- [86] CircleCI. CircleCI and AWS. Available at <https://circleci.com/integrations/aws/>.
- [87] M. Leppanen, S. Makinen, M. Pagels, V.-P. Eloranta, J. Itkonen, M. V. Mantyla, and T. Mannisto. The highways and country roads to continuous deployment. *IEEE Software*, 32(2):64–72, 2015.
- [88] C. Benoit. Gunicorn Documentation, 2020. Available at: <https://docs.gunicorn.org/en/stable/>.
- [89] Intland Software. Quality Assurance Using Code Coverage Analysis. Available at: <https://content.intland.com/blog/sdlc/quality-assurance-using-code-coverage-analysis>.
- [90] H. Krekel. pytest Documentation. Available at: <https://docs.pytest.org/>.
- [91] M. Foord. Mock Documentation. Available at: <https://mock.readthedocs.io/en/latest/>.
- [92] Software Freedom Conservancy. The Selenium Browser Automation Project. Available at: <https://www.selenium.dev/documentation/en/>.
- [93] N. Batchelder. Coverage Documentation. Available at: <https://coverage.readthedocs.io/>.
- [94] Software Testing Help. What is Test Data? Test Data Preparation Techniques with Example. Available at: <https://www.softwaretestinghelp.com/tips-to-design-test-data-before-executing-your-test-cases/>.
- [95] S W.. Ambler. Introduction to Test Driven Development. Available at: <http://agiledata.org/essays/tdd.html>.
- [96] W.-S. Tan, D. Liu, and R. Bishu. Web evaluation: Heuristic evaluation vs. user testing. *International Journal of Industrial Ergonomics*, 39(4):621–627, 2009.
- [97] SAFe. Release on Demand. Available at: <https://www.scaledagileframework.com/release-on-demand/>.

- [98] Amazon Web Services. Amazon Compute Service Level Agreement. Available at: <https://aws.amazon.com/compute/sla/>.
- [99] Google Developers. Tools for Web Developers - Lighthouse. Available at: <https://developers.google.com/web/tools/lighthouse>.
- [100] Nomics. Nomics Cryptocurrency Bitcoin API. Available at: <https://p.nomics.com/cryptocurrency-bitcoin-api>.
- [101] Financial Conduct Authority. Understanding ‘advice’ and ‘guidance’ on investments. Available at: <https://www.fca.org.uk/consumers/understanding-advice-guidance-investments>.
- [102] Portfolio Visualizer. Portfolio visualizer - Terms of Service. Available at: <https://www.portfoliovisualizer.com/terms-of-service>.
- [103] Trust Arc. *Essential Guide to the GDPR*. Trust Arc, 2018.
- [104] UK Government. Assistive technology tools you can test with at no cost. Available at: <https://accessibility.blog.gov.uk/2018/09/27/assistive-technology-tools-you-can-use-at-no-cost/>.
- [105] Inc. Stripe. Stripe Payment Platform. Available at: <https://stripe.com/>.